

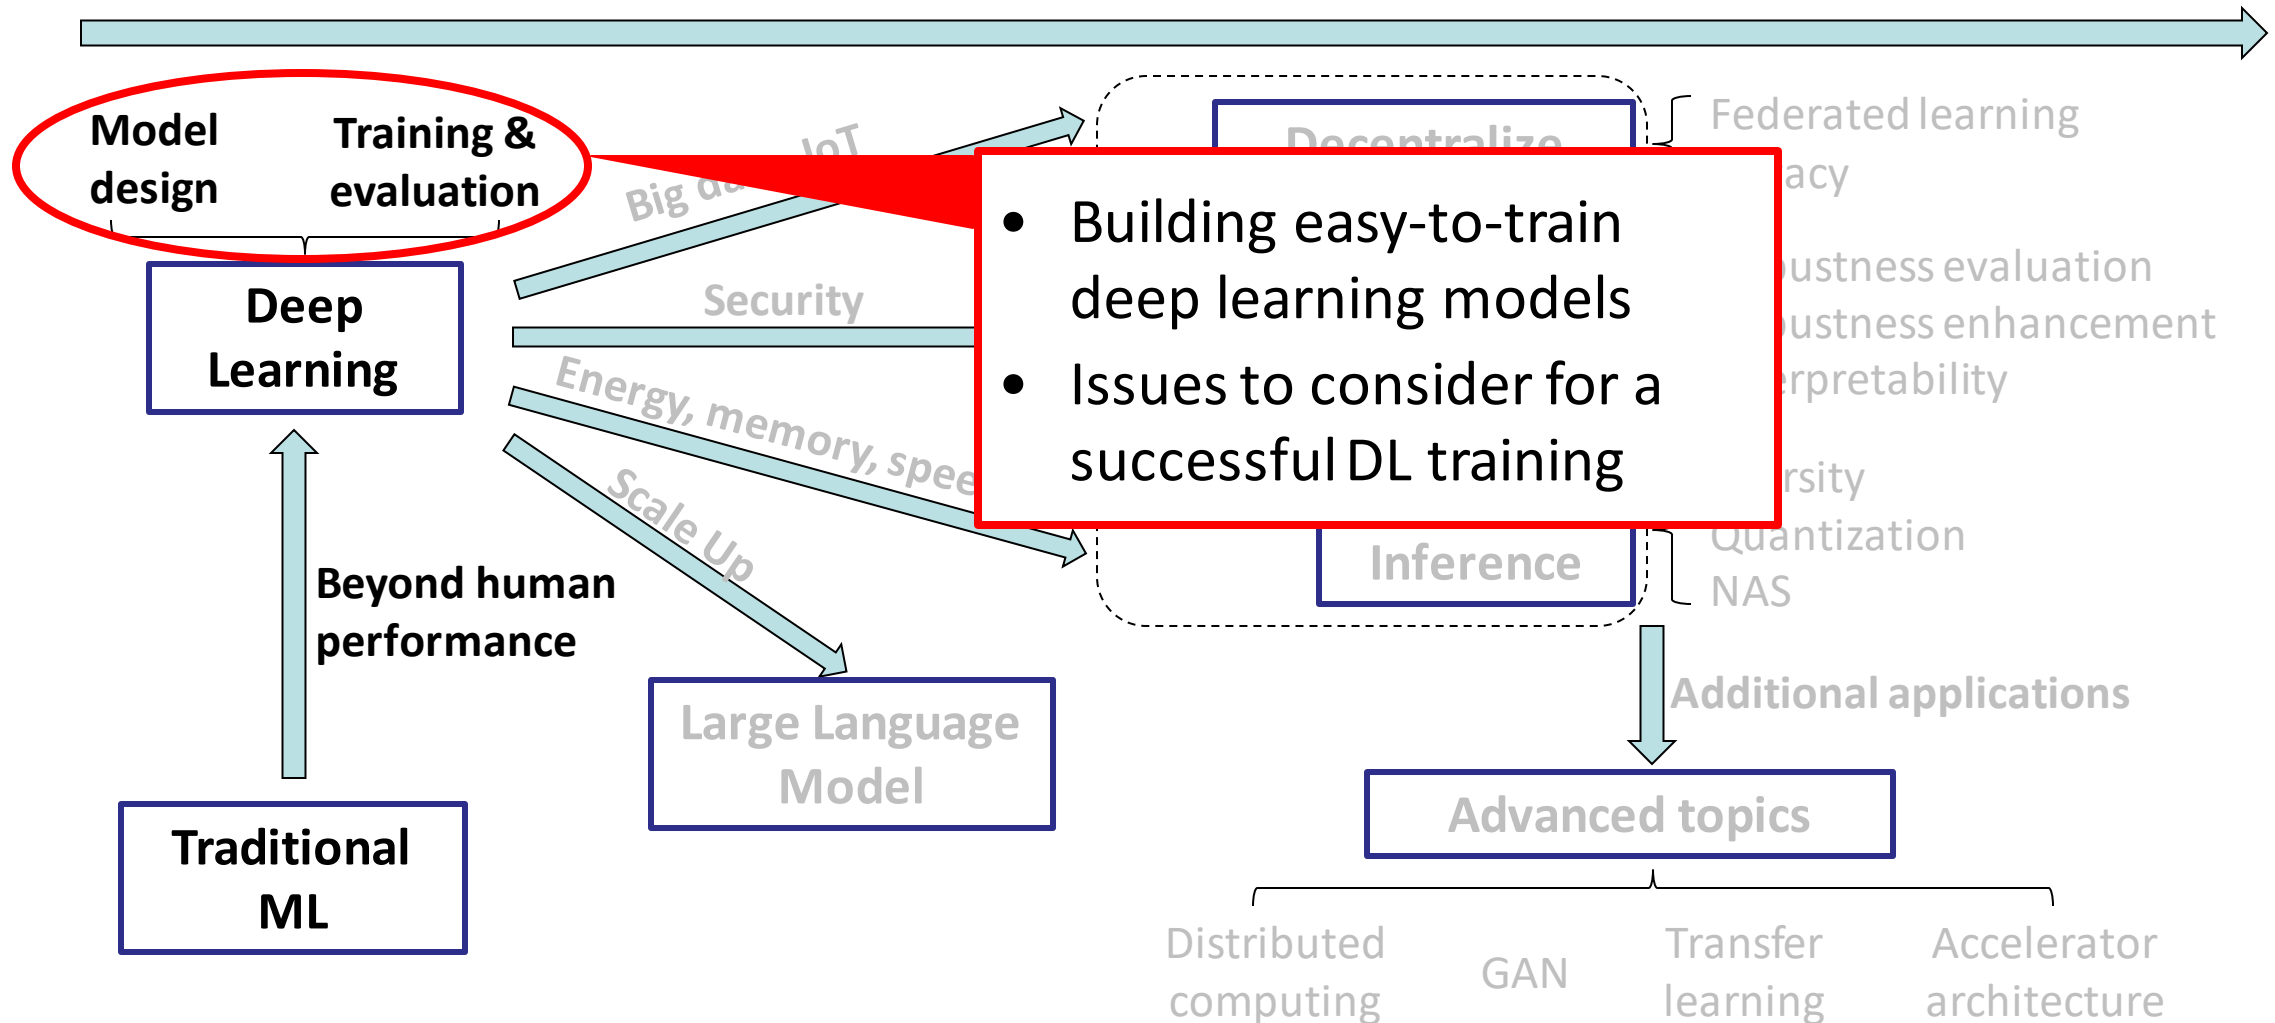


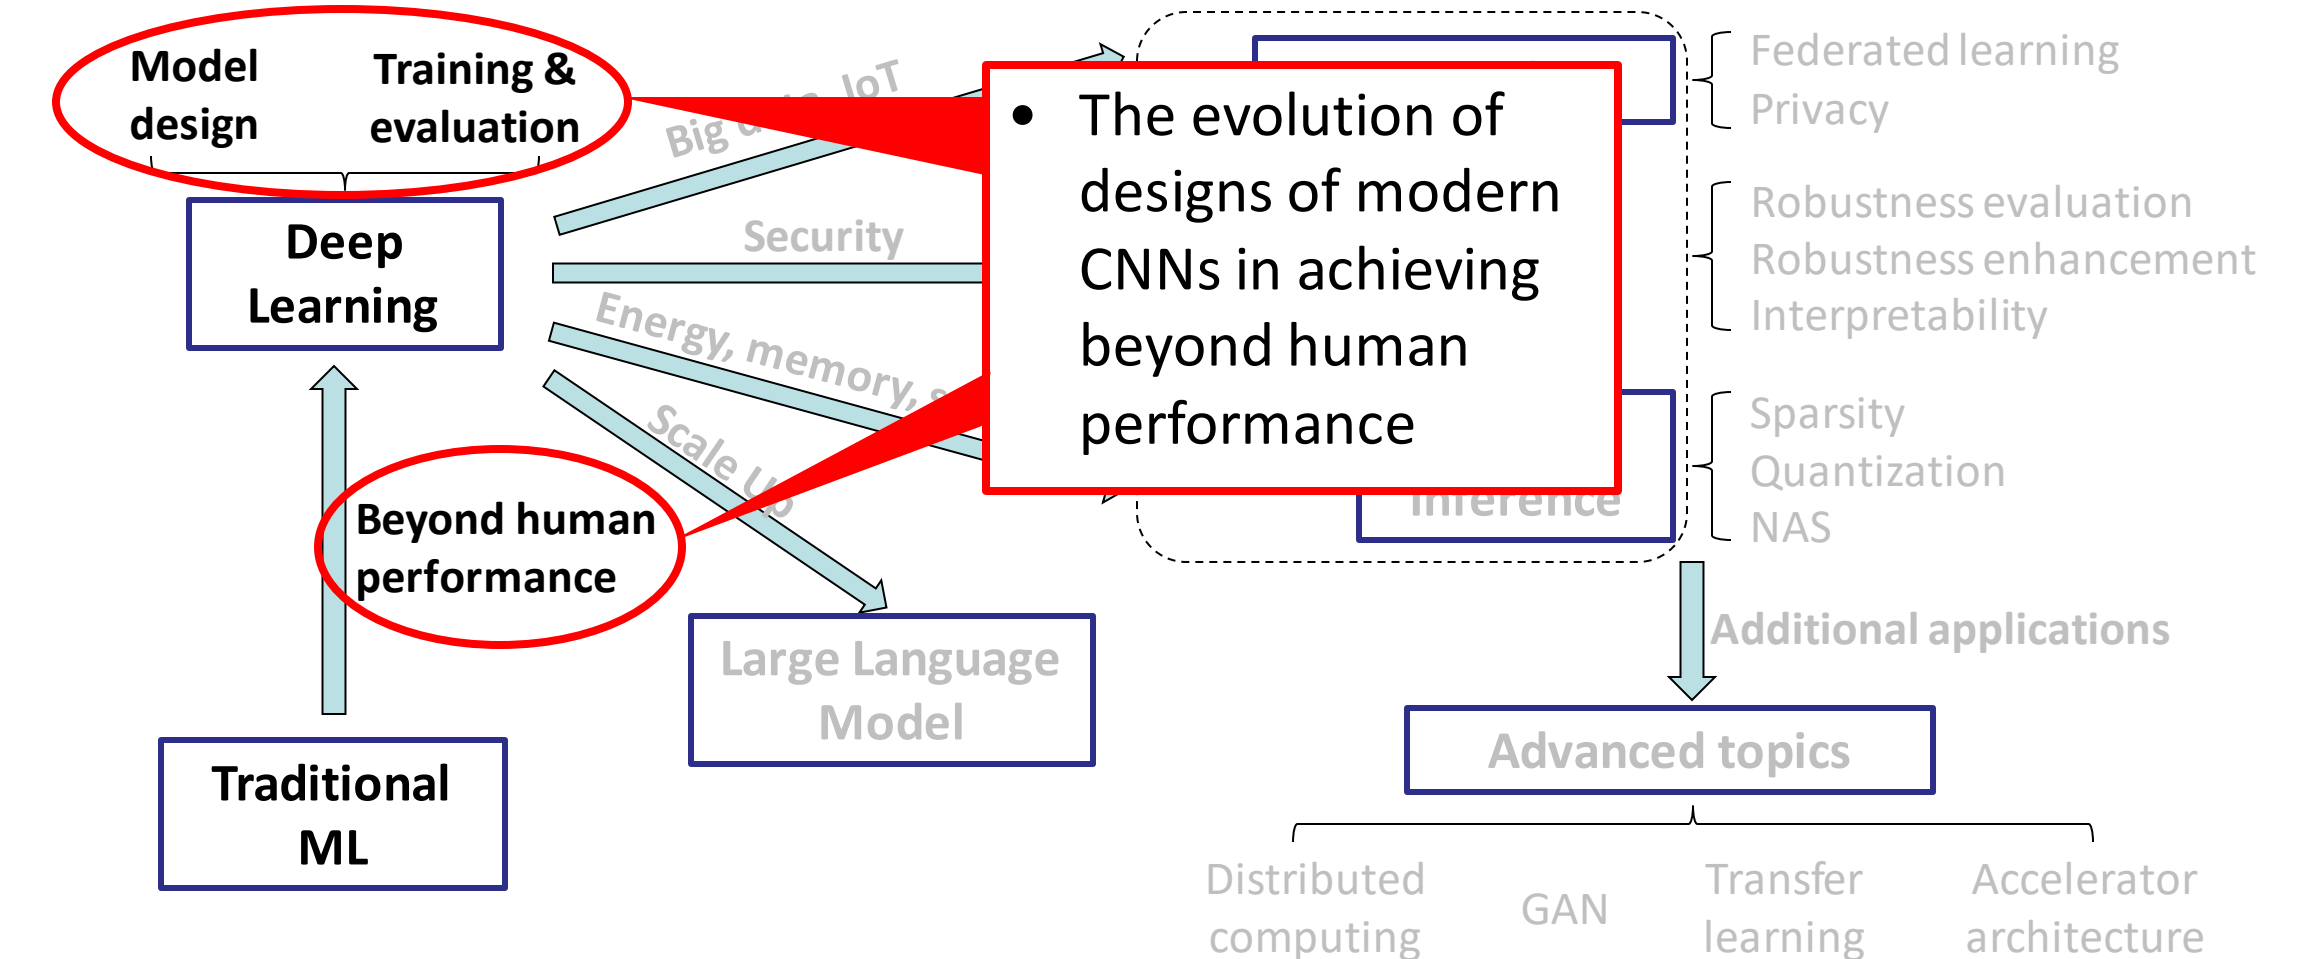
ECE 661 COMP ENG ML & DEEP NEURAL NETS

7. CNN ARCHITECTURES


The previous lecture

Applying machine learning into the real world





Recap: LeNet-5



Features:

- LeNet-5 is the first one to stack CONV-POOL-CONV-POOL structures.
- LeNet-5 achieves good results on MNIST dataset.
- LeNet-5 uses tanh activation, which has serious gradient vanishing problem.

CNN architectures

AlexNet: the **1st** deep convolution neural network for image classification

VGG: **very deep** convolution neural network for image recognition


GoogLeNet (Inception): going **deeper** with convolutions

ResNet: deep **residual** learning for image recognition

DenseNet: **densely connected** convolutional network


Overview: ILSVRC winners

- Large Scale Visual Recognition Competition (ILSVRC) is a world-famous challenge in image classification.
- CNNs are doing better as new architectures emerge.



AlexNet

- AlexNet is the first CNN model which achieves great success in image classification tasks.
- Due to limited computing power, AlexNet is spread over 2 GPUs with each one containing half of the channels.




Channels only
cross-talk at
certain layers.

Features:

- Use very large kernels (11x11) in the first layer
- The first architecture using ReLU as non-linear activations
- Use local response normalization (LRN) to speedup training
- **ImageNet top-5 error: 16.4%**

AlexNet

- AlexNet is the first CNN model which achieves great success in image classification tasks.
- Due to limited computing power, AlexNet is spread over 2 GPUs with each one containing half of the channels.



Question: What is the computation cost (MACs) for the first layer?

Answer: $11 \times 11 \times 3 \times 55 \times 55 \times 96 = 105.4M$

Features:

- Use very large kernels (11x11) in the first layer
- The first architecture using ReLU as non-linear activations
- Use local response normalization (LRN) to speedup training
- ImageNet top-5 error: 16.4%

AlexNet

Local response normalization (LRN)

- AlexNet places LRN after the ReLU non-linearity in some layers.


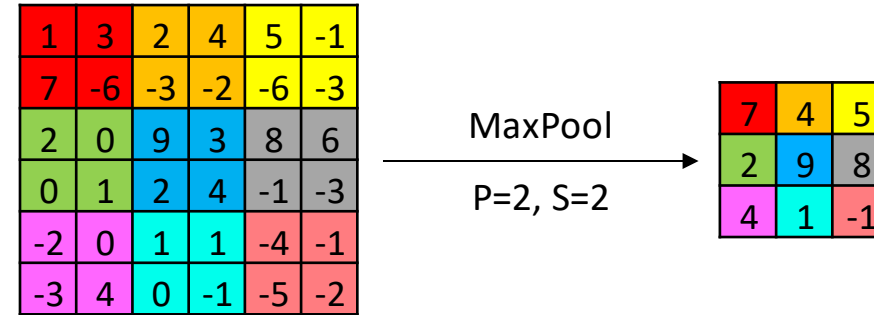
$$b_{x,y}^i = a_{x,y}^i / \left(k + \alpha \sum_{j=\max(0, i-n/2)}^{\min(N-1, i+n/2)} (a_{x,y}^j)^2 \right)^\beta$$

- LRN normalizes the activations to speed up the training of neural networks. It can also be viewed as an early version of batch normalization which is conducted within channels.
- In practice, AlexNet determines the hyperparameters k, n, α, β by running experiments on a validation set.
 - $k = 2, n = 5, \alpha = 10^{-4}$, and $\beta = 0.75$.

AlexNet

Overlapping Pooling

- AlexNet sets a larger pool size in each pooling layer to conduct overlapping pooling. This means that each pooling window **has a slight overlapping**.
- Overlapping pooling makes the model more difficult to overfit as it considers the magnitude of values over a joint pooling region.



Larger activation will dominate the down-sampled feature map.

AlexNet

Training AlexNet

- Batch size 128, initial learning rate 0.01, weight decay $5e-4$
 - Use SGD Momentum with momentum 0.9
 - Use normalization layers (LRN)
 - Dropout 0.5 for FC layers (except for final FC)
 - Use aggressive data augmentation (shifting, flipping, etc.)
- Later training approaches inherit most of the AlexNet approach.

VGG

VGG creates very deep convolutional neural networks (11-19 layers) for image classification.


Features:

- Use 3×3 convolution to replace larger kernels (e.g., 11×11 in AlexNet)
- ImageNet top-5 error: 7.3%
- Until now, VGG model is still popular among a variety of tasks other than ImageNet (e.g., object detection)
- VGG is very large
 - 140/144 Million parameters and 16G/20G MACs in VGG 16/19
 - Hard to deploy for real-time applications



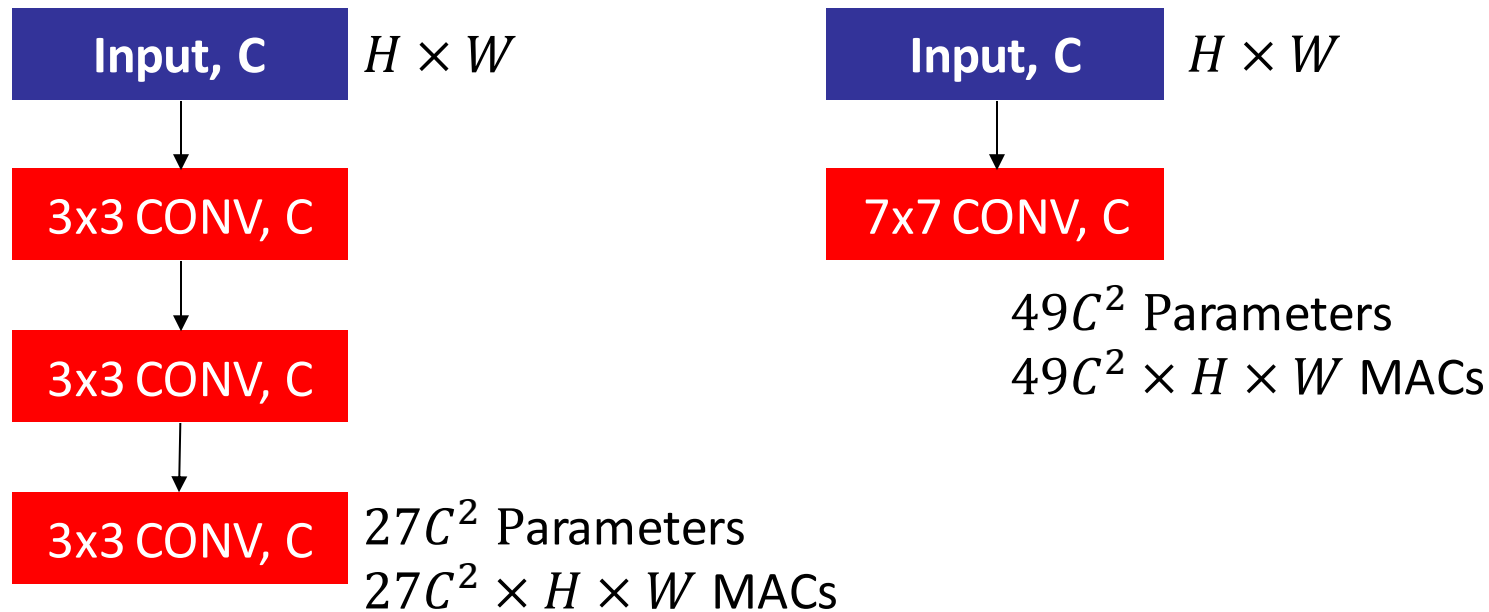
VGG: 3x3 convolution

- Receptive field is the size of the region in the input that produces the feature. VGG enlarges the receptive field by stacking **ONLY** 3x3 convolutions.
 - Two 3x3 convolutions has the same receptive field as one 5x5 convolution, however, with fewer parameters and MACs
 - However, two 3x3 convolutions perform better as it makes the model deeper



VGG: 3x3 convolution

- Receptive field is the size of the region in the input that produces the feature. VGG enlarges the receptive field by stacking **ONLY** 3x3 convolutions.
 - Three 3x3 convolution has the same receptive field as one 7x7 convolution with fewer parameters and MACs
 - However, three 3x3 convolutions perform better as it makes the model deeper



VGG

Training VGG

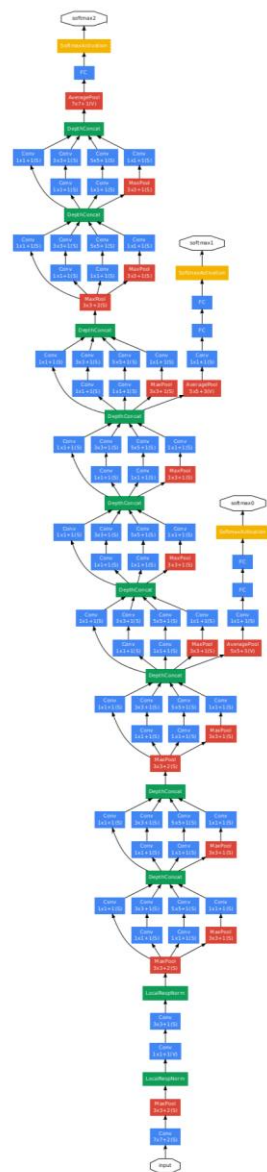
- Batch size 256, initial learning rate is 0.01, weight decay $5e-4$
- Decay the learning rate by a factor of 0.1 if the validation accuracy plateaus
- Use SGD Momentum optimizer with momentum 0.9
- Dropout 0.5 for FC layers (except final FC that produces output)
- Data preprocessing is a bit different from AlexNet. For example, VGG-16 uses **multi-crop** (10-Crop) evaluation for better inference performance. See the VGG paper for more details
- Use **model ensembling** to reach better results. This is common in ILSVRC competition



Designing CNNs in a nutshell.

Fun fact, this meme was referenced in the first inception net paper.





Designing CNNs in a nutshell.
Fun fact, this meme was referenced in the first
inception net paper.

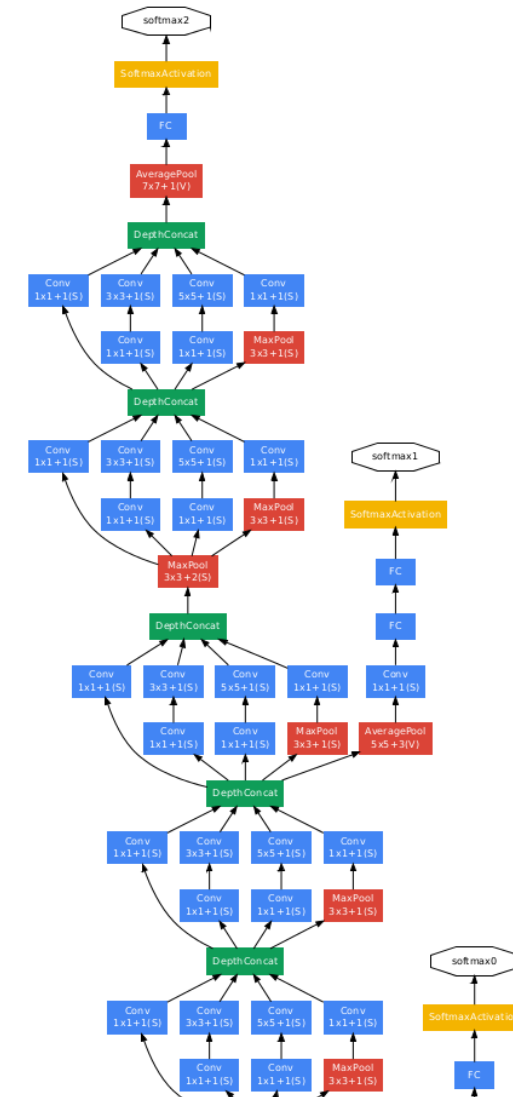


GoogLeNet (Inception)

- From GoogLeNet, CNNs not only go deeper but evolve into multiple branches.

Features:

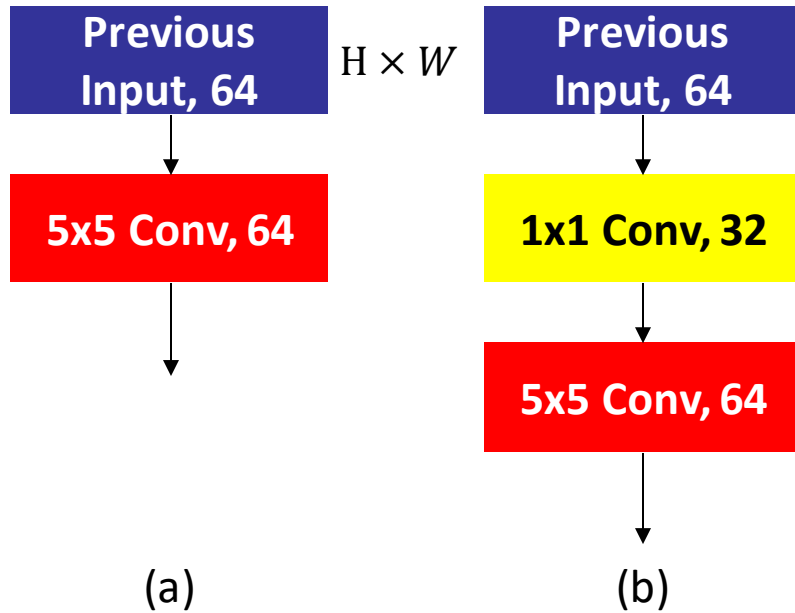
- Build a deeper model with a mixture of convolutions (3x3, 5x5, 1x1).
- Extract information from different parts of CNNs.
- Uses **local response normalization (LRN)** to speed up training.
- Use **bottleneck** structure to reduce parameters.
 - Yet, Inception is still time and memory-consuming.
 - The training takes about a week to reach convergence using a few high-end GPUs.
- ImageNet top-5 error: 6.7%



GoogLeNet: bottleneck structure


The number of filters is reduced before convolving large filters (e.g., 5x5). This structure is also called **bottleneck** structure.

Question: Calculate the weight parameter and MAC numbers for models in (a) & (b).



GoogLeNet: bottleneck structure

The number of filters is reduced before convolving large filters (e.g., 5x5). This structure is also called **bottleneck** structure.



Question: Calculate the weight parameter and MAC numbers for models in (a) & (b).

Answer:

(a) $5 \times 5 \times 64 \times 64 = 102,400$ weight parameters

$5 \times 5 \times 64 \times 64 \times H \times W = 102,400 \times H \times W$ MACs

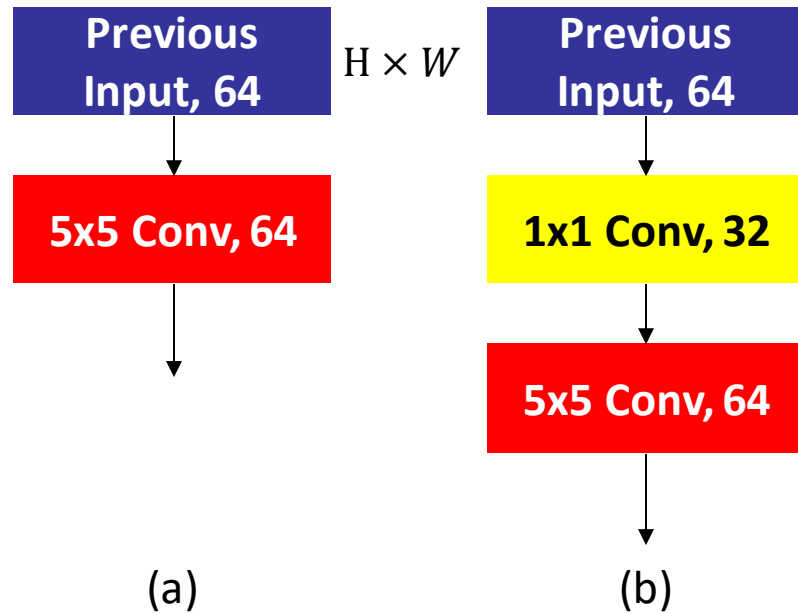
(b) $1 \times 1 \times 64 \times 32 + 5 \times 5 \times 32 \times 64 = 53,248$ weight parameters

$(1 \times 1 \times 64 \times 32 + 5 \times 5 \times 32 \times 64) \times H \times W = 53,248 \times H \times W$ MACs

GoogLeNet: bottleneck structure

Why is bottleneck structure useful?

- It seems that the bottleneck structures cut down the number of feature maps. Does it affect the performance of GoogLeNet?

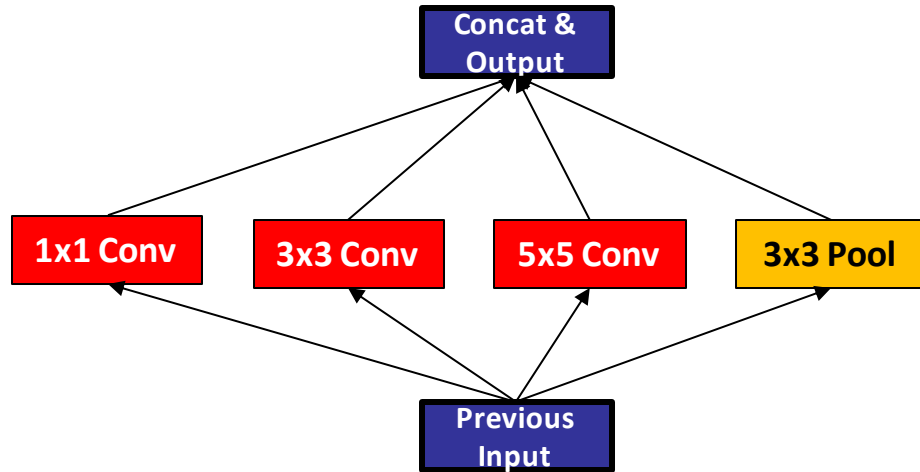


- Bottleneck structures project (compress) high-dimensional inputs into low-dimensional space.
- This 'projection' may even act as a form of regularization to prevent overfitting, which improves the performance of GoogLeNet.

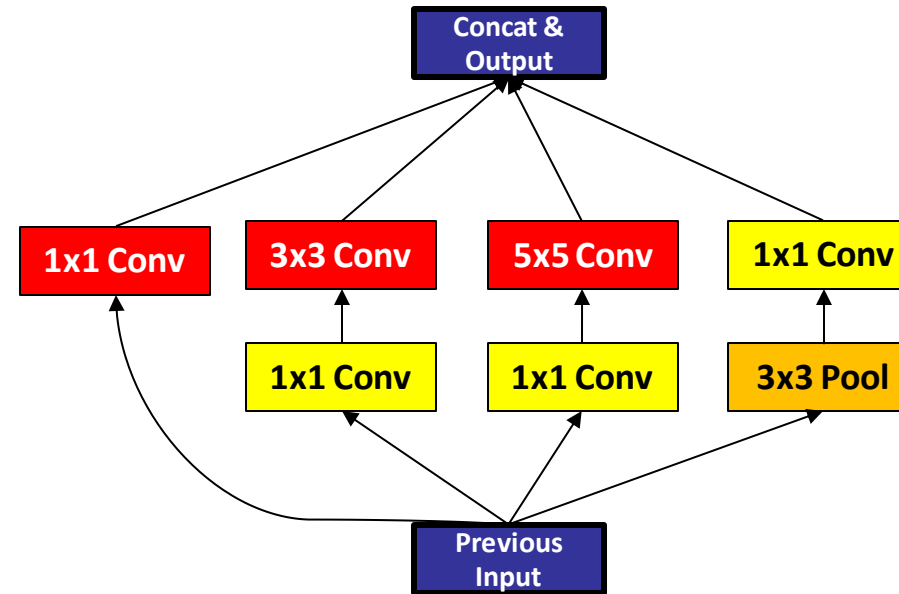
GoogLeNet: bottleneck structure

Where do we usually apply bottleneck structure?

- Bottleneck structure can be applied to any **branch** within the GoogLeNet structure.



Naïve GoogLeNet



GoogLeNet with dimension reduction

GoogLeNet

Train GoogLeNet

- Batch size is 256, initial learning rate is 0.01, decay the learning rate by a factor of 0.96 every 8 epochs
- Use **auxiliary tower** connected to intermediate layers as ways of combatting the overfitting problem. However, this is not common in later training approaches
- 40% dropout before the final FC layer
- Use SGD Momentum optimizer with momentum 0.9
- Preprocessing is a bit different from VGG. See the GoogLeNet-V1 (V2) paper for details
- Average predictions over multiple crops of the same input image

GoogLeNet-V2 (Inception-V2)

- GoogLeNet-V2 introduces **Batch Normalization (BN)** and adds it before each nonlinearity function in the Inception architecture.




Figure 2: *Single crop validation accuracy of Inception and its batch-normalized variants, vs. the number of training steps.*

Features:

- Remove local response normalization (LRN) layers in previous architecture
- Introduce Batch Normalization (BN)
- GoogLeNet-V2 is trained faster than its GoogLeNet-V1 counterpart
- ImageNet top-5 error: 4.9%

ResNet: deep residual learning

Even with batch normalization, very deep neural networks are difficult to train without **residual learning**.

- **Optimization difficulty** grows with the number of parameters in a deep neural network
- **Gradient explosion/vanishing** problems occur when the networks go deeper
- It is desirable to learn **identity mappings** for generalization



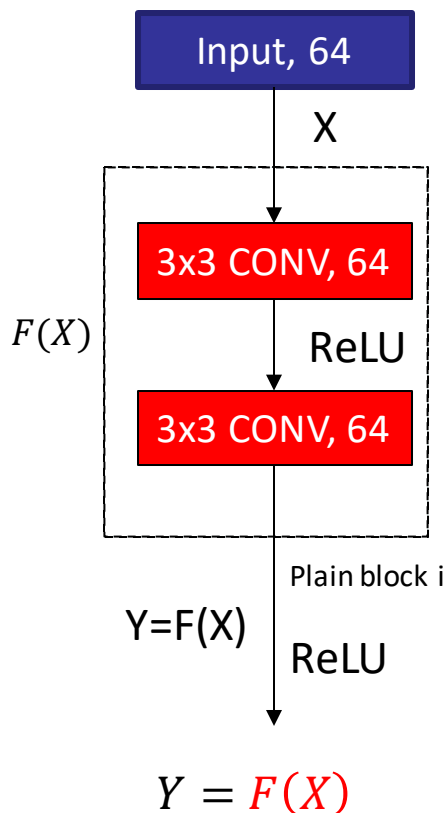


Figure 1. Training error (left) and test error (right) on CIFAR-10 with 20-layer and 56-layer “plain” networks. The deeper network has higher training error, and thus test error. Similar phenomena on ImageNet is presented in Fig. 4.

Residual learning can do it!

ResNet

- During DNN training, we are trying to learn a function $F(X)=Y$.
- ResNet adds a shortcut connection to fit a residual mapping $F(X)=Y-X$.
This is called **residual learning**.



$$Y = F(x, \{W_i\}) + X$$

Residual
mapping

Identity
mapping



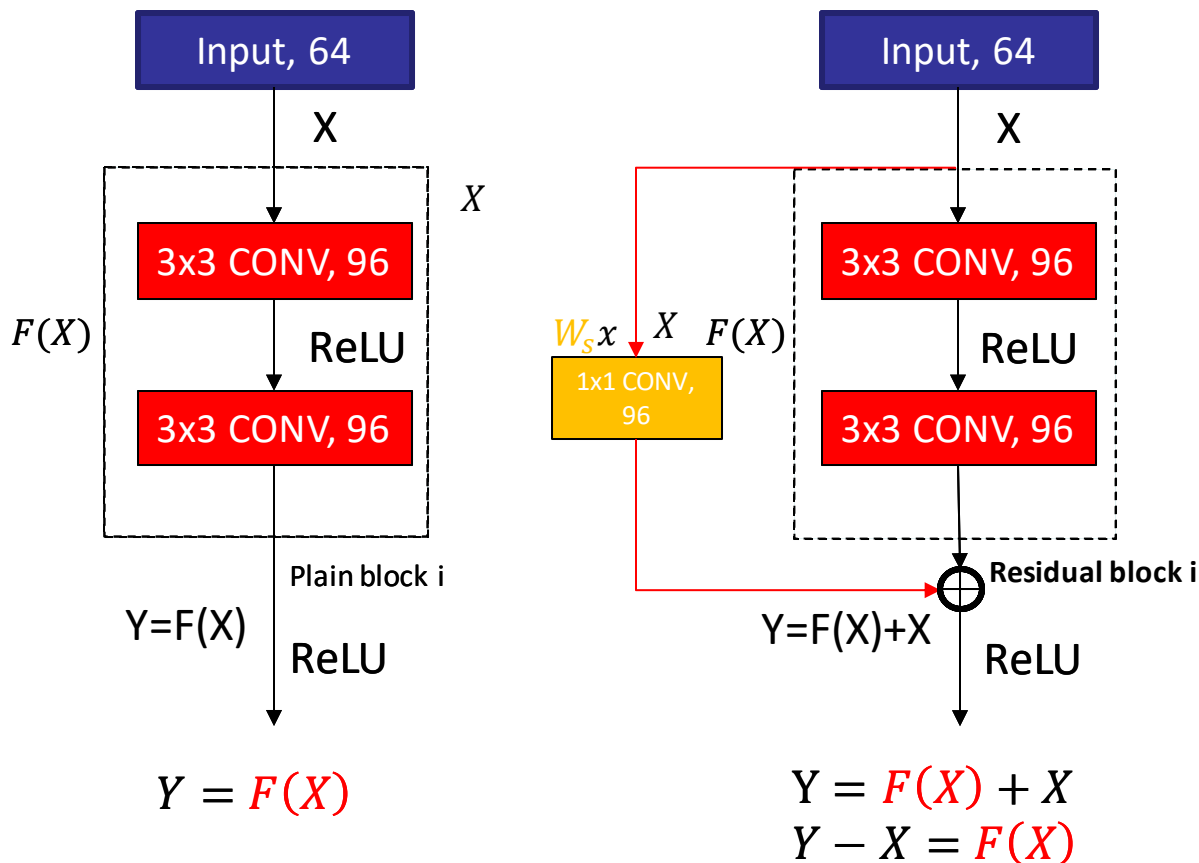
Kaiming He
Associate Professor, EECS, MIT
Verified email at mit.edu - [Homepage](#)
Computer Vision Machine Learning



ResNet

What if the number of output filters in Y changes after passing function $F(X)$?

- Solution: Use 1x1 convolution to match dimension.



$$Y = F(x, \{W_i\}) + W_S x$$


Original mapping Residual mapping

$W_S \in \mathbb{R}^{1 \times 1 \times 64 \times 96}$ is used for identity mapping.


ResNet: visualization of loss surface

Residual learning make the optimization process easier!

- Shortcut connections make the loss surface smoother and easier to optimize.




Loss surface **without** shortcut connections. Noisy & curvy. Optimization process is easy to get stuck into local minimum.




Loss surface **with** shortcut connections. Smooth & easy to optimize. Optimization process is easy to find sharp minima.

ResNet: visualization of loss values


- The accuracy of plain networks is upper-bounded by the total number of layers.
- The accuracy of ResNet grows with the number of residual layers.



Plain networks



ResNets

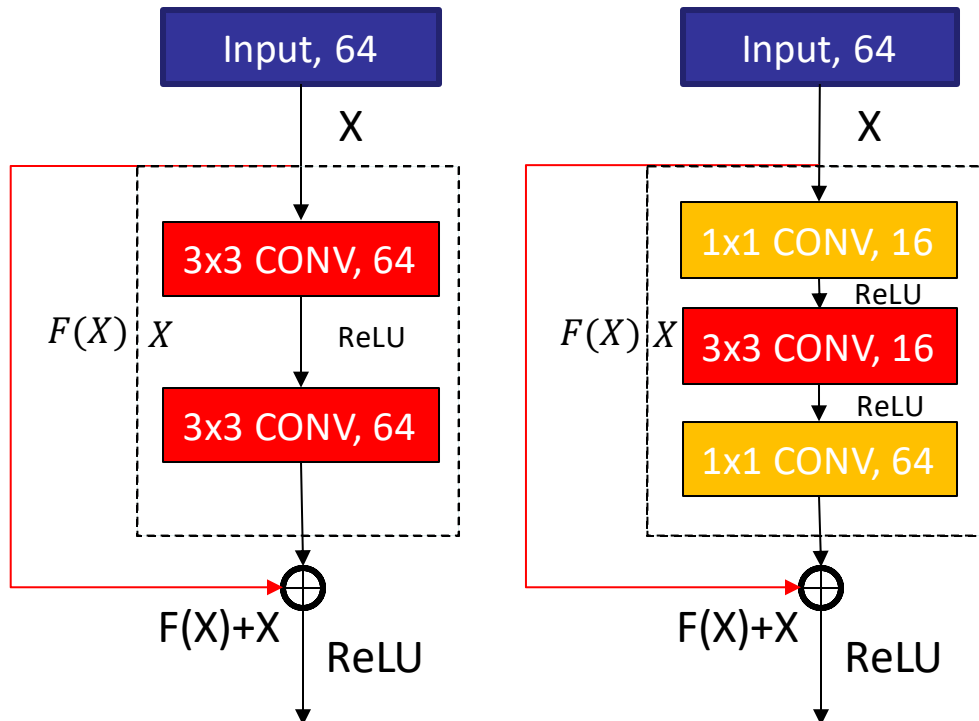


Very deep ResNets

- However, we cannot indefinitely add the depth of neural networks: a ResNet with more than 1000 layers is still difficult to optimize!

ResNet: residual block with bottleneck

- Similar to GoogLeNet, ResNet incorporates bottleneck layers into residual blocks to improve the efficiency of the model. This is called **ResNet Bottleneck**.
- Bottleneck layers make residual blocks deeper and more efficient.



16 1x1 filters **project** to 16 feature maps.

64 1x1 filters **expand** back to 64 feature maps.

ResNet: residual block with bottleneck

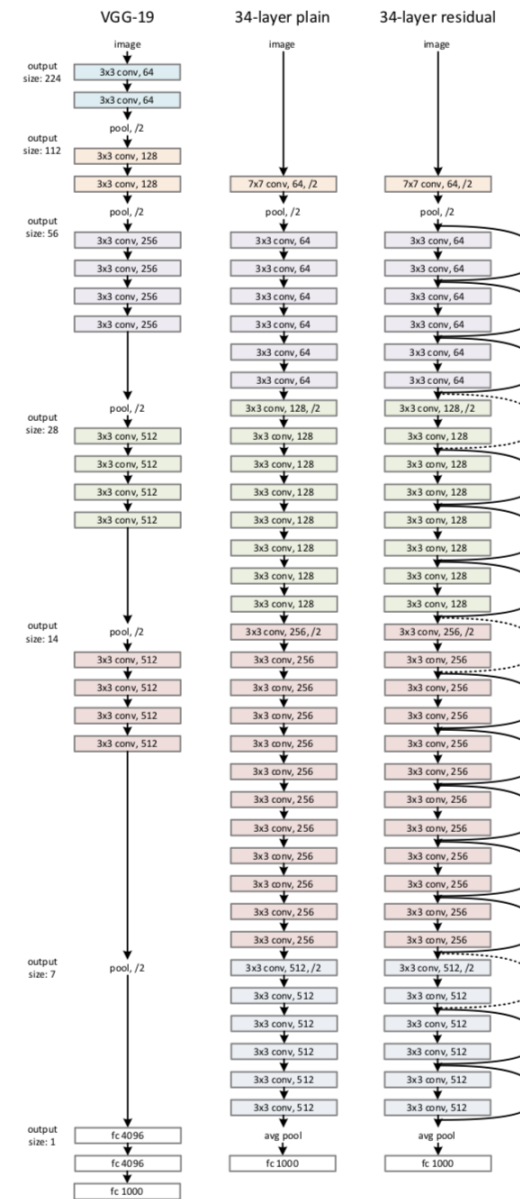
- Similar to GoogLeNet, ResNet incorporates bottleneck layers into residual blocks to improve the efficiency of the model. This is called **ResNet Bottleneck**.
- Bottleneck layers make residual blocks deeper and more efficient.

| layer name | output size | 18-layer | 34-layer | 50-layer | 101-layer | 152-layer |
|------------|-------------|---|---|--|---|---|
| conv1 | 112×112 | 7×7, 64, stride 2 | | | | |
| conv2_x | 56×56 | 3×3 max pool, stride 2 | | | | |
| | | $\begin{bmatrix} 3\times 3, 64 \\ 3\times 3, 64 \end{bmatrix} \times 2$ | $\begin{bmatrix} 3\times 3, 64 \\ 3\times 3, 64 \end{bmatrix} \times 3$ | $\begin{bmatrix} 1\times 1, 64 \\ 3\times 3, 64 \\ 1\times 1, 256 \end{bmatrix} \times 3$ | $\begin{bmatrix} 1\times 1, 64 \\ 3\times 3, 64 \\ 1\times 1, 256 \end{bmatrix} \times 3$ | $\begin{bmatrix} 1\times 1, 64 \\ 3\times 3, 64 \\ 1\times 1, 256 \end{bmatrix} \times 3$ |
| conv3_x | 28×28 | $\begin{bmatrix} 3\times 3, 128 \\ 3\times 3, 128 \end{bmatrix} \times 2$ | $\begin{bmatrix} 3\times 3, 128 \\ 3\times 3, 128 \end{bmatrix} \times 4$ | $\begin{bmatrix} 1\times 1, 128 \\ 3\times 3, 128 \\ 1\times 1, 512 \end{bmatrix} \times 4$ | $\begin{bmatrix} 1\times 1, 128 \\ 3\times 3, 128 \\ 1\times 1, 512 \end{bmatrix} \times 4$ | $\begin{bmatrix} 1\times 1, 128 \\ 3\times 3, 128 \\ 1\times 1, 512 \end{bmatrix} \times 8$ |
| conv4_x | 14×14 | $\begin{bmatrix} 3\times 3, 256 \\ 3\times 3, 256 \end{bmatrix} \times 2$ | $\begin{bmatrix} 3\times 3, 256 \\ 3\times 3, 256 \end{bmatrix} \times 6$ | $\begin{bmatrix} 1\times 1, 256 \\ 3\times 3, 256 \\ 1\times 1, 1024 \end{bmatrix} \times 6$ | $\begin{bmatrix} 1\times 1, 256 \\ 3\times 3, 256 \\ 1\times 1, 1024 \end{bmatrix} \times 23$ | $\begin{bmatrix} 1\times 1, 256 \\ 3\times 3, 256 \\ 1\times 1, 1024 \end{bmatrix} \times 36$ |
| conv5_x | 7×7 | $\begin{bmatrix} 3\times 3, 512 \\ 3\times 3, 512 \end{bmatrix} \times 2$ | $\begin{bmatrix} 3\times 3, 512 \\ 3\times 3, 512 \end{bmatrix} \times 3$ | $\begin{bmatrix} 1\times 1, 512 \\ 3\times 3, 512 \\ 1\times 1, 2048 \end{bmatrix} \times 3$ | $\begin{bmatrix} 1\times 1, 512 \\ 3\times 3, 512 \\ 1\times 1, 2048 \end{bmatrix} \times 3$ | $\begin{bmatrix} 1\times 1, 512 \\ 3\times 3, 512 \\ 1\times 1, 2048 \end{bmatrix} \times 3$ |
| | 1×1 | average pool, 1000-d fc, softmax | | | | |
| FLOPs | | 1.8×10^9 | 3.6×10^9 | 3.8×10^9 | 7.6×10^9 | 11.3×10^9 |

ResNet

Features

- ResNet constructs the neural network by stacking residual blocks. Each residual block has 2 3×3 convolutions.
- ResNet enables the design of very deep convolutional networks (from ResNet-18 to ResNet-152).
- By using bottleneck architectures, ResNet can be deeper and more efficient.
- ResNet achieves 3.57% top-5 error on ImageNet dataset.




ResNet

Training ResNet

- Batch size is 256, use learning rate 0.01 to train for a few warmup epochs, then switch to initial learning rate 0.1
- Decay the learning rate by 0.1 when the validation accuracy plateaus
- Use SGD momentum with momentum 0.9
- Use $1e-5$ L2 regularization
- Use similar preprocessing methods as inception

DenseNet

- DenseNet formulates **dense connectivity** between layers.




Features

- DenseNet encourages feature reuse, sharing, and interaction between **different positions (layers)** of the network.
- DenseNet uses a **transition** layer to do down-sampling on images.
- With respect to feature maps of similar depth, DenseNet reduces the weight parameter by using **filter concatenation**. However, DenseNet can lead to high memory consumption.
- ImageNet top-5 error: 5.3%

DenseNet: dense connections


Dense connections is achieved by **filter concatenation**.



Standard ConvNet Concept

Standard ConvNet:


The output of the current layer is directly passed to the next layer.



+ : Element-wise addition

ResNet:

The output of the current layer is first added to the identity of the current residual block, then passed to the next layer.



⊕ : Channel-wise concatenation

One Dense Block in DenseNet

DenseNet:

The output of the previous layers in the current dense block is first concatenated and then passed to the next layer.

DenseNet

Training DenseNet

- Similar to ResNet
- The authors adopted the same training methodology as ResNet for fair comparison

Summary of CNN architectures

| Architecture | Single-Path? | Use small (≤ 3) convolutional kernel? | Depth? |
|--------------|--------------|--|----------------------------|
| AlexNet | ✓ | ✗ | Shallow (8 layers) |
| VGG | ✓ | ✓ | Deep (16-19 layers) |
| GoogLeNet | ✗ | ✓ (V3) ✗ (V1,V2,V4) | Deeper (22 layers) |
| ResNet | ✗ | ✓ | Very deep (34-152 layers) |
| DenseNet | ✗ | ✓ | Very deep (121-264 layers) |

| Architecture | Input Size | Parameter Memory | MACs |
|--------------|------------|------------------|-------------|
| AlexNet | 224x224 | 233 MB | 727 M |
| VGG-16/19 | 224x224 | 528 MB / 548 MB | 16 G / 20 G |
| GoogLeNet | 224x224 | 51 MB | 2 G |
| ResNet-18/34 | 224x224 | 45 MB / 83 MB | 2 G / 4 G |
| DenseNet-121 | 224x224 | 31 MB | 3 G |

Summary of CNN architectures

| Metrics | LeNet 5 | AlexNet | Overfeat fast | VGG 16 | GoogLeNet v1 | ResNet 50 |
|--|-------------------|-----------|------------------|-----------|-----------------|--------------|
| Top-5 error [†] | n/a | 16.4 | 14.2 | 7.4 | 6.7 | 5.3 |
| Top-5 error (single crop) [†] | n/a | 19.8 | 17.0 | 8.8 | 10.7 | 7.0 |
| Input Size | 28×28 | 227×227 | 231×231 | 224×224 | 224×224 | 224×224 |
| # of CONV Layers | 2 | 5 | 5 | 13 | 57 | 53 |
| Depth in # of CONV Layers | 2 | 5 | 5 | 13 | 21 | 49 |
| Filter Sizes | 5 | 3,5,11 | 3,5,11 | 3 | 1,3,5,7 | 1,3,7 |
| # of Channels | 1, 20 | 3-256 | 3-1024 | 3-512 | 3-832 | 3-2048 |
| # of Filters | 20, 50 | 96-384 | 96-1024 | 64-512 | 16-384 | 64-2048 |
| Stride | 1 | 1,4 | 1,4 | 1 | 1,2 | 1,2 |
| Weights | 2.6k | 2.3M | 16M | 14.7M | 6.0M | 23.5M |
| MACs | 283k | 666M | 2.67G | 15.3G | 1.43G | 3.86G |
| # of FC Layers | 2 | 3 | 3 | 3 | 1 | 1 |
| Filter Sizes | 1,4 | 1,6 | 1,6,12 | 1,7 | 1 | 1 |
| # of Channels | 50, 500 | 256-4096 | 1024-4096 | 512-4096 | 1024 | 2048 |
| # of Filters | 10, 500 | 1000-4096 | 1000-4096 | 1000-4096 | 1000 | 1000 |
| Weights | 58k | 58.6M | 130M | 124M | 1M | 2M |
| MACs | 58k | 58.6M | 130M | 124M | 1M | 2M |
| Total Weights | 60k | 61M | 146M | 138M | 7M | 25.5M |
| Total MACs | 341k | 724M | 2.8G | 15.5G | 1.43G | 3.9G |
| Pretrained Model Website | [56] [‡] | [57, 58] | n/a | [57–59] | [57–59] | [57–59] |

The original training setup for ImageNet

| | AlexNet | VGG | GoogLeNet | ResNet |
|----------------------------------|-----------------------|-----------------------|---------------------------------------|------------------------------------|
| Year | 2012 | 2014 | 2014 | 2015 |
| Layer # | 8 | 16-19 | 22 | 34-152 |
| Batch size | 128 | 256 | 256 | 256 |
| LRN vs. BN | LRN | * | V1: LRN V2: BN | BN |
| Learning rate | 0.01 | 0.01, decay 0.1 | 0.01, decay 0.96 every 8 epochs | 0.1 (0.01 warmup), decay 0.1 |
| Optimizer | SGD w/o momentum | SGD w/o momentum =0.9 | | |
| Weight decay (regularization) | 5e-4 | 5e-4 | | 1e-5 |
| Dropout | 0.5 | 0.5 | 0.4 | No dropout |
| Data preprocessing | Shifting, flipping | + Multi-crop | + Brightness, aspect ratio distortion | |
| Model ensemble | No | Yes | Yes | Yes |

Reading Materials

- AlexNet:** Alex Krizhevsky, Ilya Sutskever, and Geoffrey E. Hinton. "Imagenet classification with deep convolutional neural networks." (2012)
- VGG:** Karen Simonyan and Andrew Zisserman. "Very deep convolutional networks for large-scale image recognition." (2014)
- GoogLeNet (Inception):** Christian Szegedy, et al. "Going deeper with convolutions." (2015)
Sergey Lofte, et al., Batch Normalization: Accelerating Deep Network Training by Reducing Internal Covariate Shift (2015)
- ResNet:** Kaiming He, et al. "Deep residual learning for image recognition." (2015)
- DenseNet:** Gao Huang, et al. "Densely connected convolutional networks." (2017)