



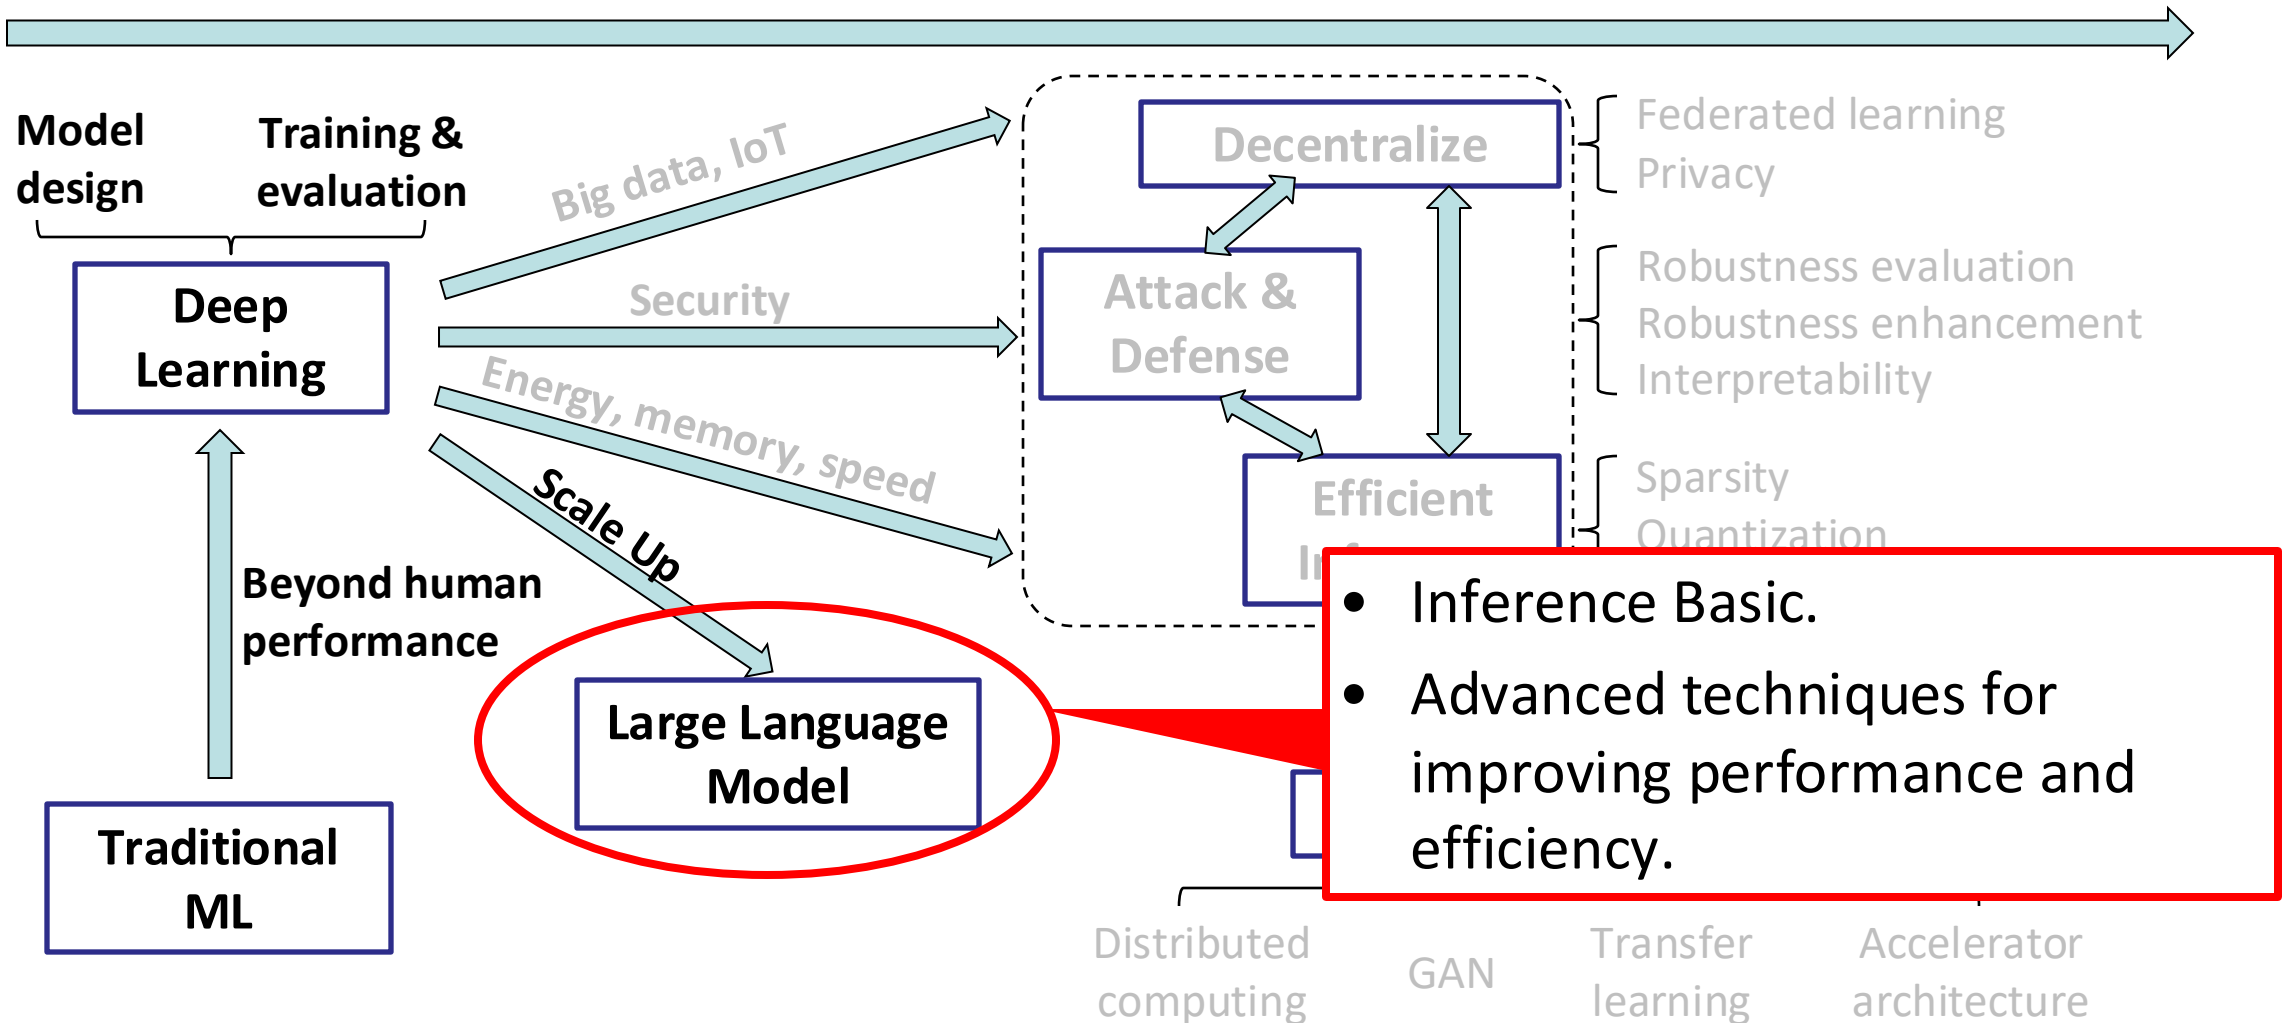
ECE 661 COMP ENG ML & DEEP NEURAL NETS

**11. LARGE LANGUAGE MODELS INFERENCE**

HAI "HELEN" LI, SPRING 2025 <sub>1</sub>

# This lecture

Applying machine learning into the real world



# Outline

---

## Lecture 11: Large Language Model Inference

- **Inference Basic**
- Additional Transformer Designs
  - KV cache
  - Attention mechanisms optimization
- Advanced Inference Systems
  - FlashAttention
  - vLLM
- LLM Inference Randomness

# LLM Inference

---

- What is Inference in Machine Learning?
  - The process of inputting new data into a trained machine learning model to generate a prediction.
  - Example: Inputting an image into a CNN model to recognize its class
- What is Inference in context of LLMs?
  - The process of generating text outputs based on input prompts, by iteratively predicting the next token in a sequence.

# LLM Inference Procedure

---

- Loading Weight to GPU
- Tokenizing the input text sequence (Prompt)
- Prefill Phase
  - Decoding Phase
- Detokenize output tokens

## Key Phases

# LLM Inference Procedure

---

- Loading Weight to GPU
  - LLaMa-2-7B (FP32 ~ 28GB)
- Tokenizing the input text sequence (Prompt)
  - Tokenizer breaks down text into tokens (e.g word, subword, characters)
  - Tokens are converted into vectors that model can understand
  - Text -> tokens -> vector

**What is LLM inference?**



What is LLM inference?

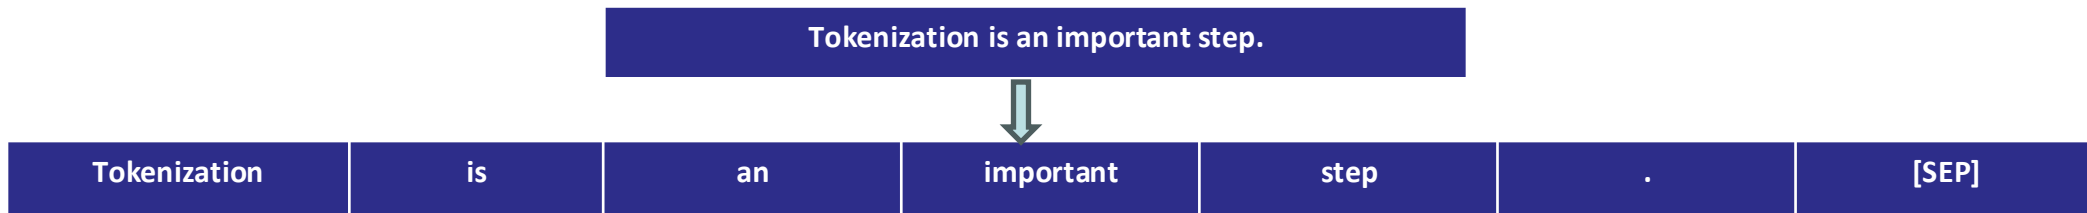


[3923, 374, 445, 11237, 45478, 30]

# Tokenization

---

- Tokenization is the process of dividing text into smaller units called tokens, which are typically words or sub-words.
- Tokens are mapped to vectors for use in neural networks.



Two Approaches :

- **Top-Down (Rule-based tokenization)** uses predefined rules to segment text into tokens, typically based on grammar and syntax, e.g., splitting sentences at punctuation marks or spaces.
- **Bottom-up (Subword tokenization)** breaks down words into smaller units, such as subwords or characters, allowing for the handling of unknown words and variations, e.g., **Byte Pair Encoding used in BERT and GPT.**

# Byte-Pair Encoding

---

Byte Pair Encoding is a compression-based tokenization method that iteratively merges the most frequent character pairs to create subword units.

**Step 1:** Start with a vocabulary containing the individual characters present in the training corpus.

**Step 2:** Examine the training corpus and identify the two most frequently adjacent symbols.

**Step 3:** Add a new merged symbol representing the combined pair to the vocabulary. Replace every instance of the adjacent pair in the corpus with the new merged symbol.

**Step 4:** Continue counting and merging the most frequent pairs. Repeat until you've performed  $k$  merges, resulting in  $k$  novel tokens.

**Step 5:** The final vocabulary consists of the original set of characters plus the  $k$  new symbols created through merging.



# Byte-Pair Encoding

---

Initial vocabulary:  
characters



Split each word  
into characters

Words in the data:

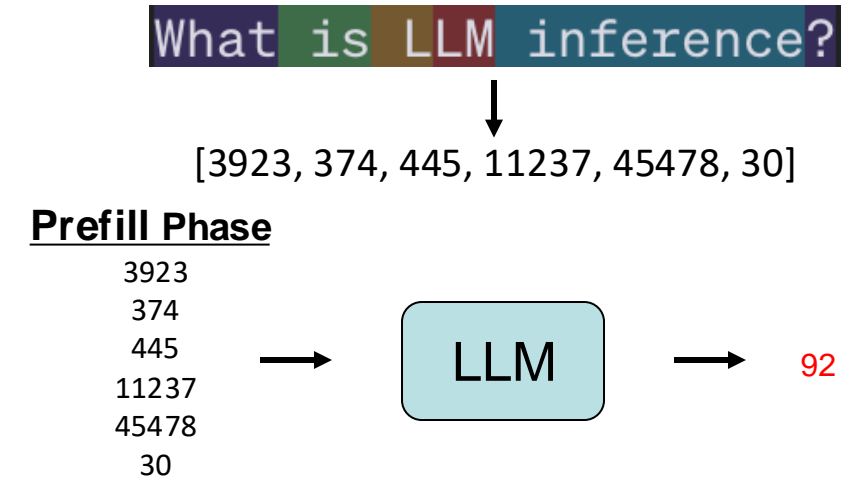
word	count
c a t	4
m a t	5
m a t s	2
m a t e	3
a t e	3
e a t	2

Current merge table:

(empty)

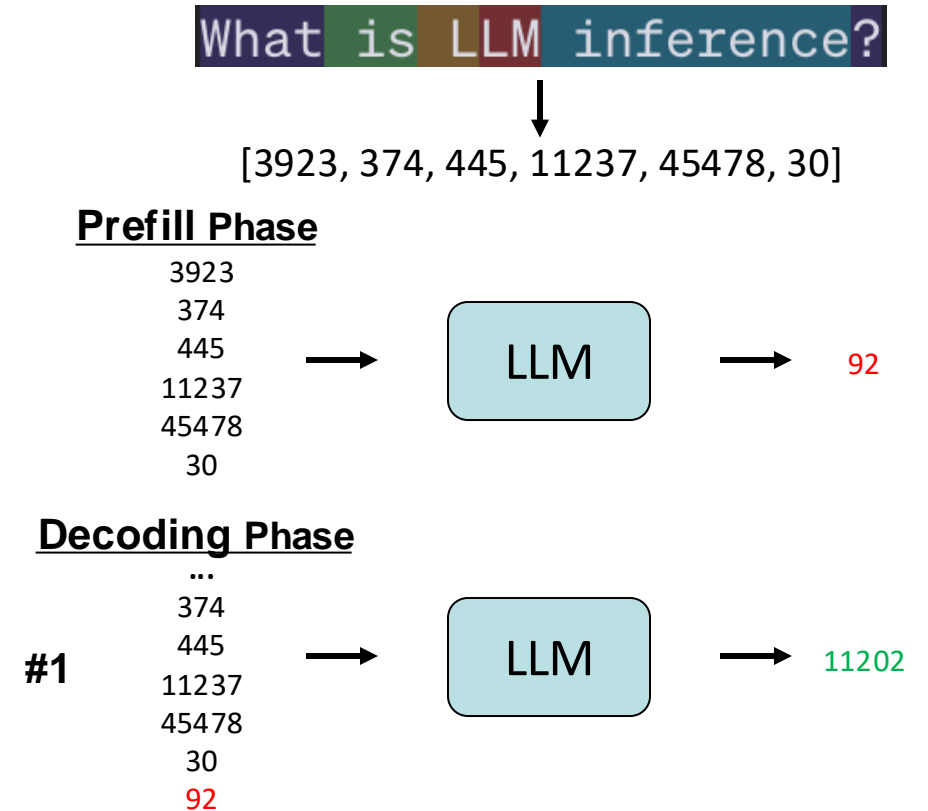
# LLM Inference Procedure

- Prefill Phase (Single-step Phase)
  - Running the tokenized prompt through the LLM Model to generate the first token



# LLM Inference Procedure

- Prefill Phase (Single-step Phase)
  - Running the tokenized prompt through the LLM Model to generate the first token
- Decoding Phase (Multi-step Phase)
  - Appending the generated token to the sequence of input tokens and using it as a new input to generate the next token

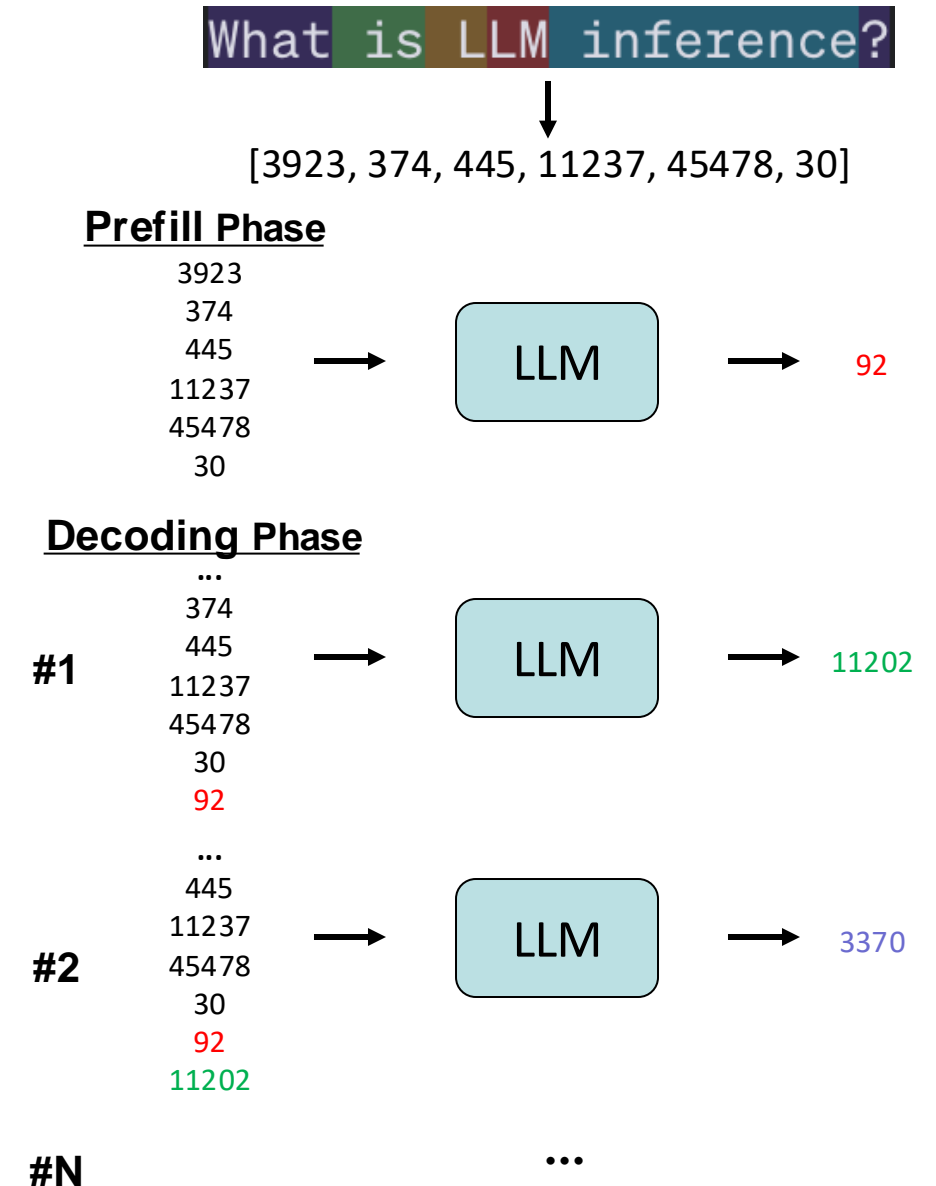


# LLM Inference Procedure

- **Prefill Phase** (Single-step Phase)
  - Running the tokenized prompt through the LLM Model to generate the first token
- **Decoding Phase** (Multi-step Phase)
  - Appending the generated token to the sequence of input tokens and using it as a new input to generate the next token

Repeat decoding until meeting a stopping criteria

- Generating end-of-sequence token
- Reaching maximum sequence length



# LLM Inference Scenarios

---

- **Inference** - Fewer request, offline traffic, latency

Take a series of tokens as inputs, and generate subsequent tokens autoregressively until they meet a stopping criteria

- Prefill Phase (Process the input)
- Decoding Phase (Generate the output)



- **Serving** - Many requests, online traffic, cost-per-query

- Co-locate the two phases and batch the computation of prefill and decoding across all users and requests



# Outline

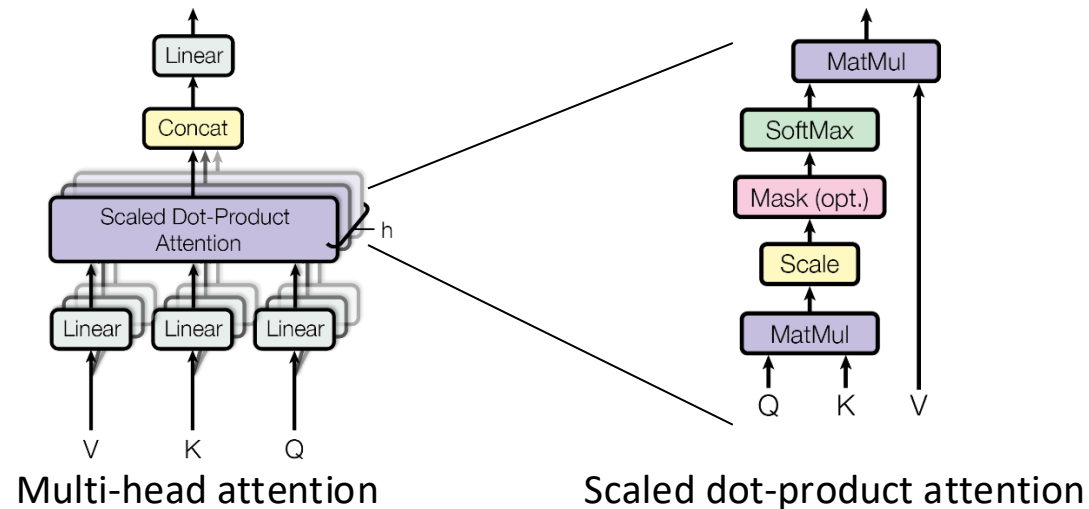
---

## Lecture 11: Large Language Model Inference

- Inference Basic
- Additional Transformer Designs
  - **KV cache**
  - Attention mechanisms optimization
- Advanced Inference Systems
  - FlashAttention
  - vLLM
- LLM Inference Randomness

# KV Cache

- Recaps Attention Function

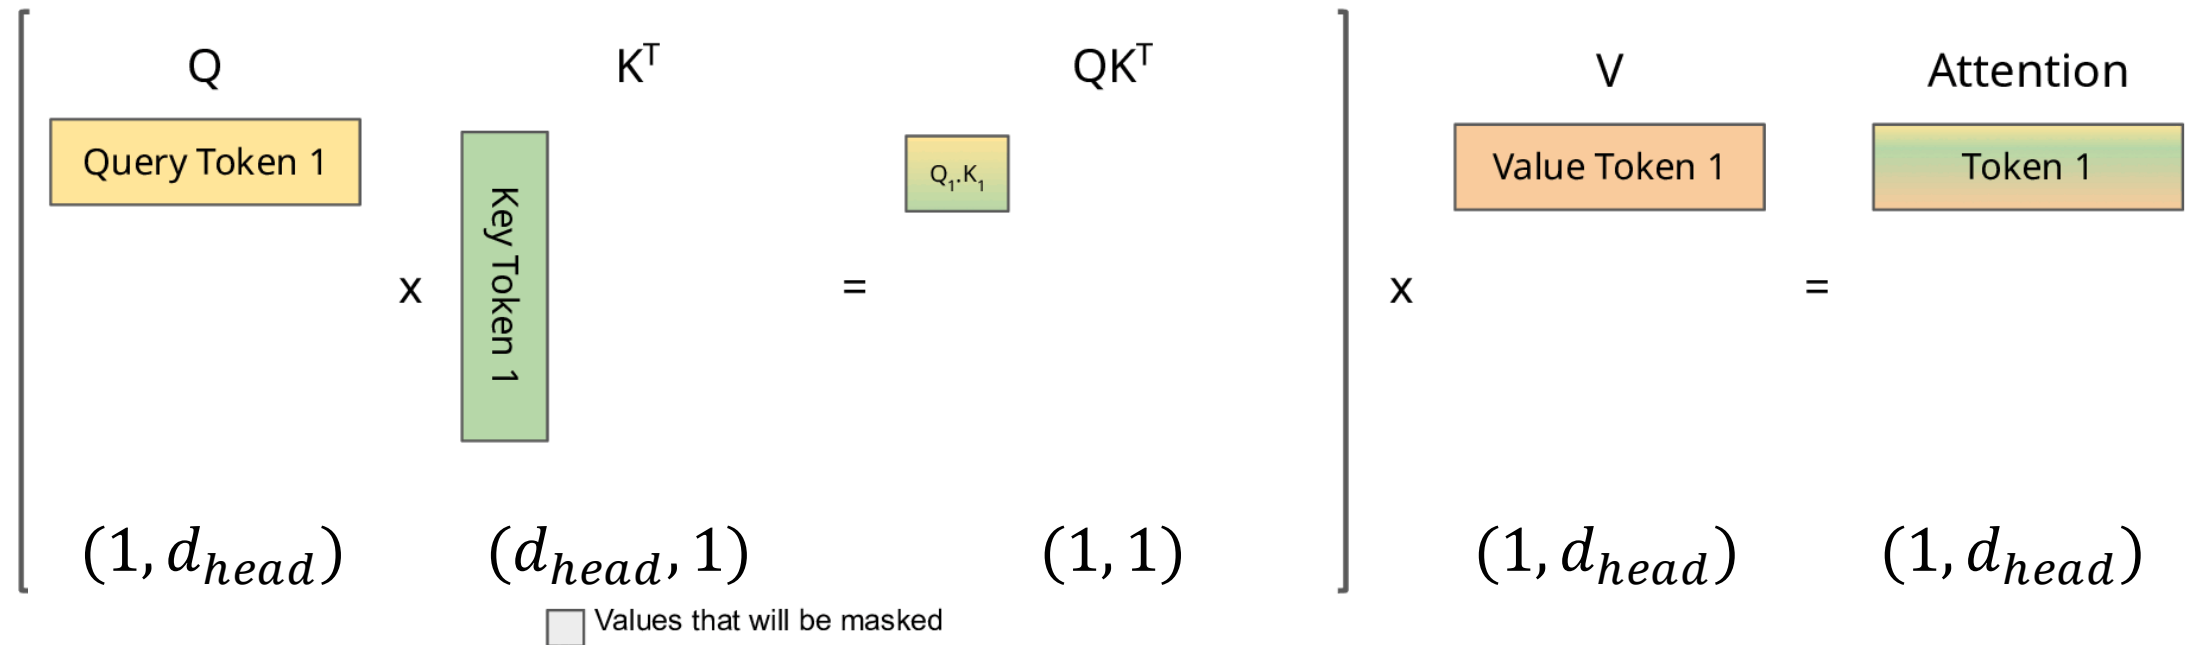


- Token generation (Attention computation)
  - Keys and Values of all preceding tokens
  - Query from current token
  - Recalculating previous tokens' attention
- Token generation only occurs in Decoder

# KV cache

- Without Cache

## Step 1



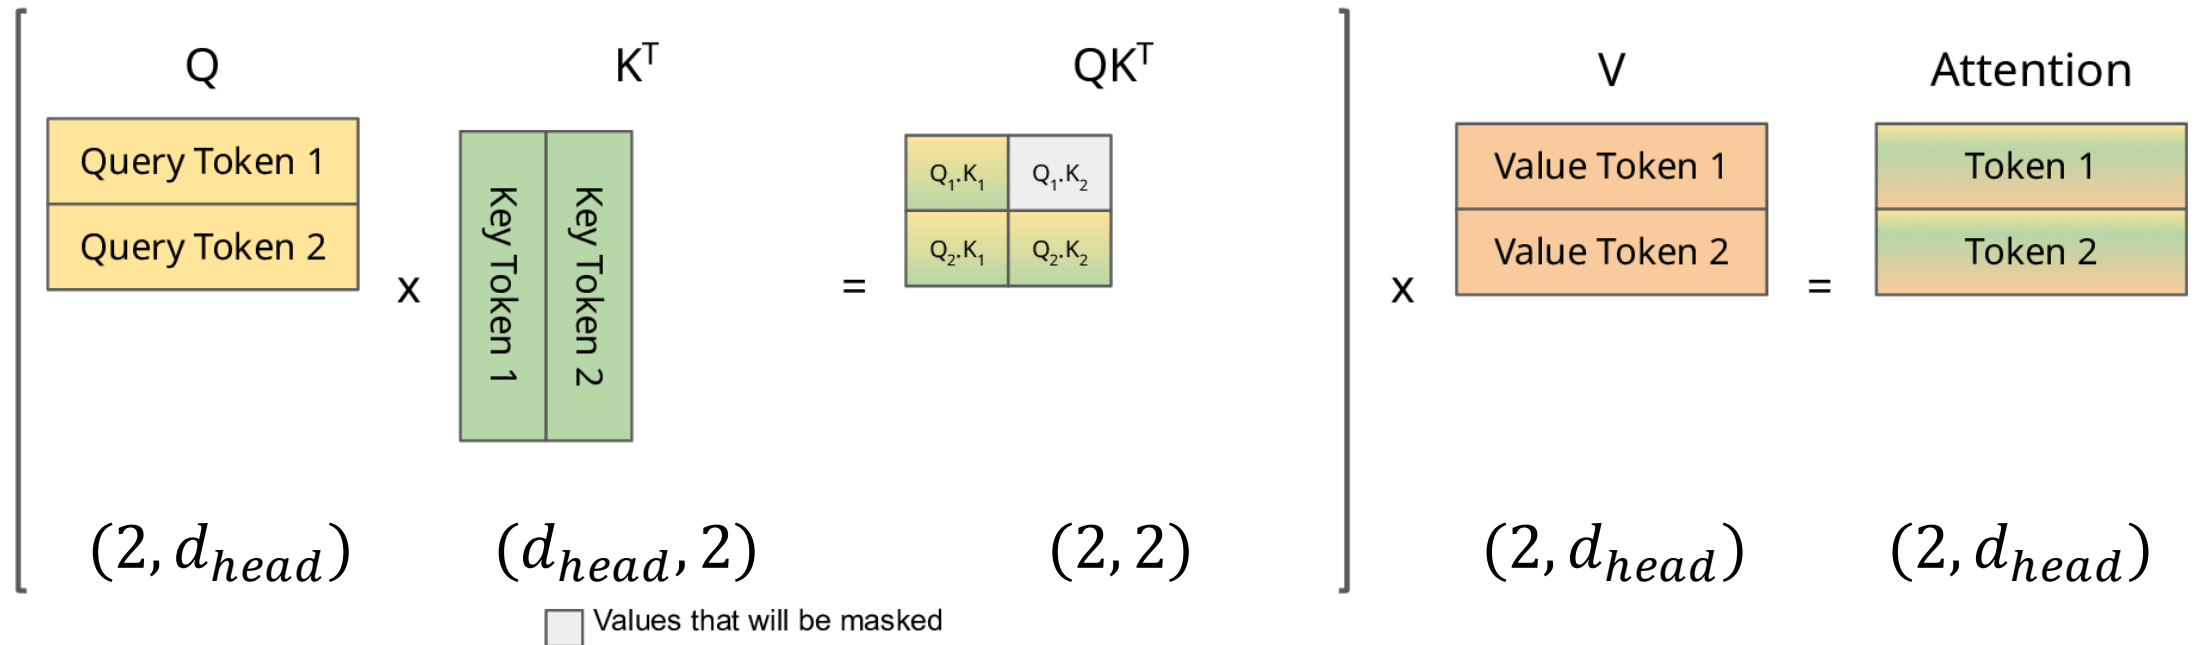
$d_{head}$ : the hidden dimension of the attention head



# KV cache

- Without Cache

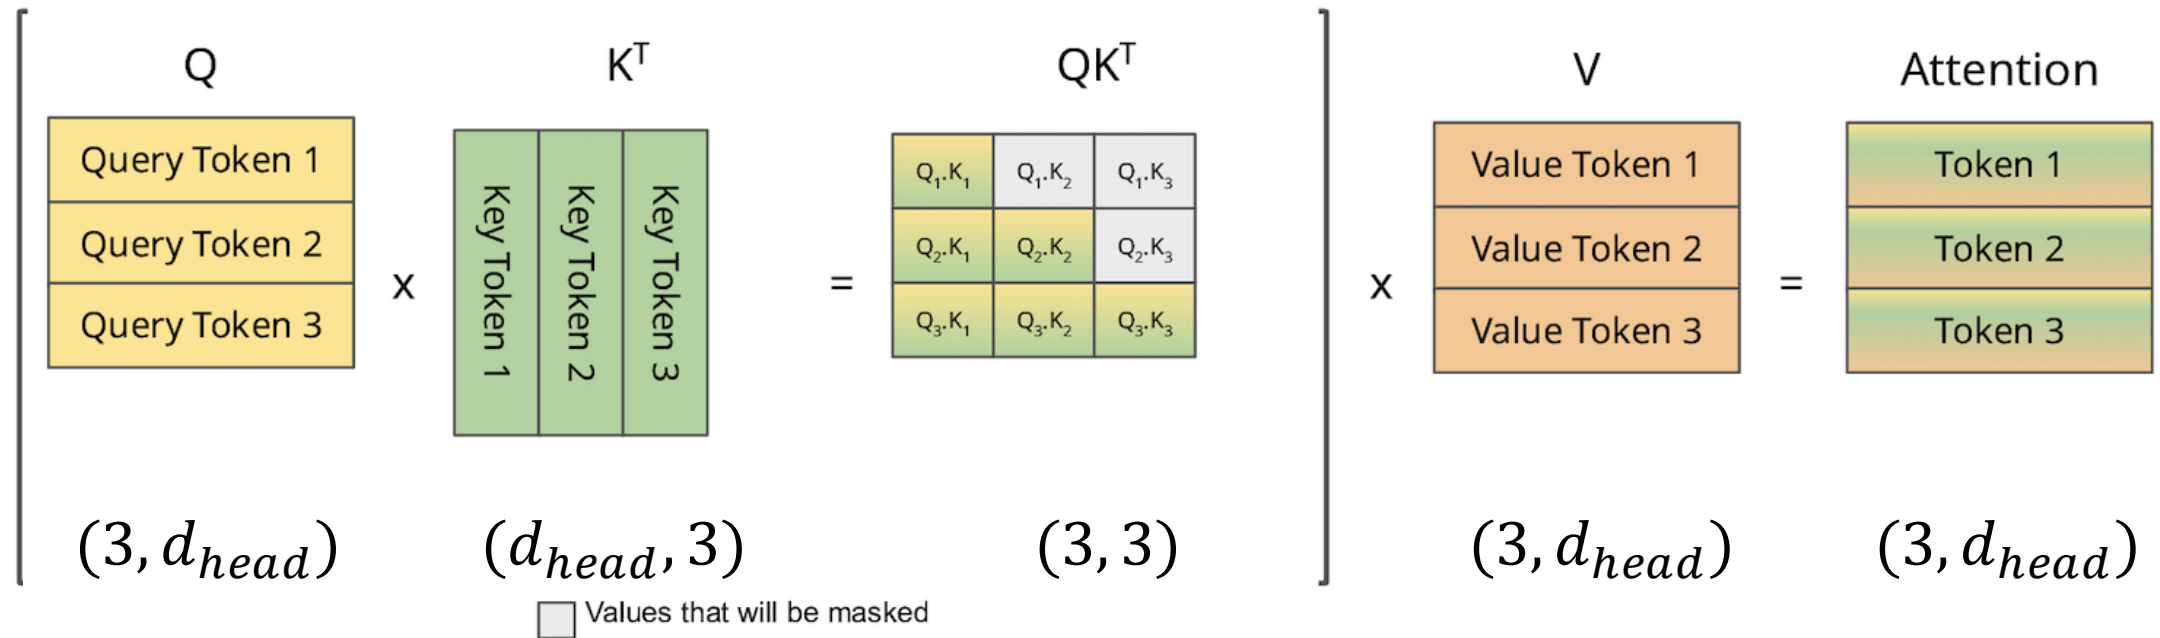
## Step 2



# KV cache

- Without Cache

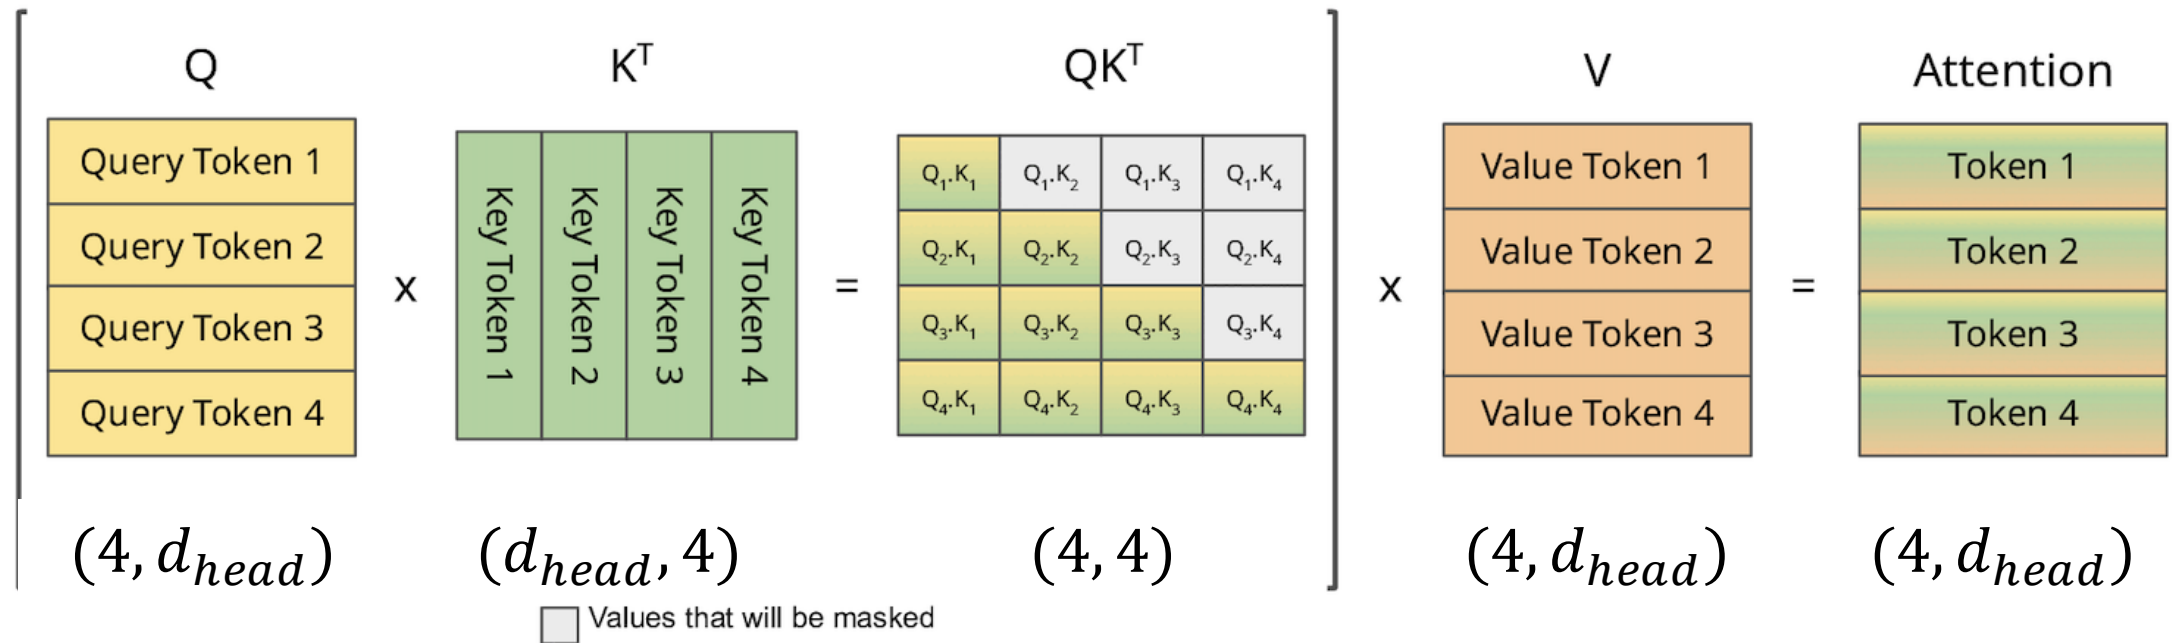
## Step 3



# KV cache

- Without Cache

## Step 4



# KV cache

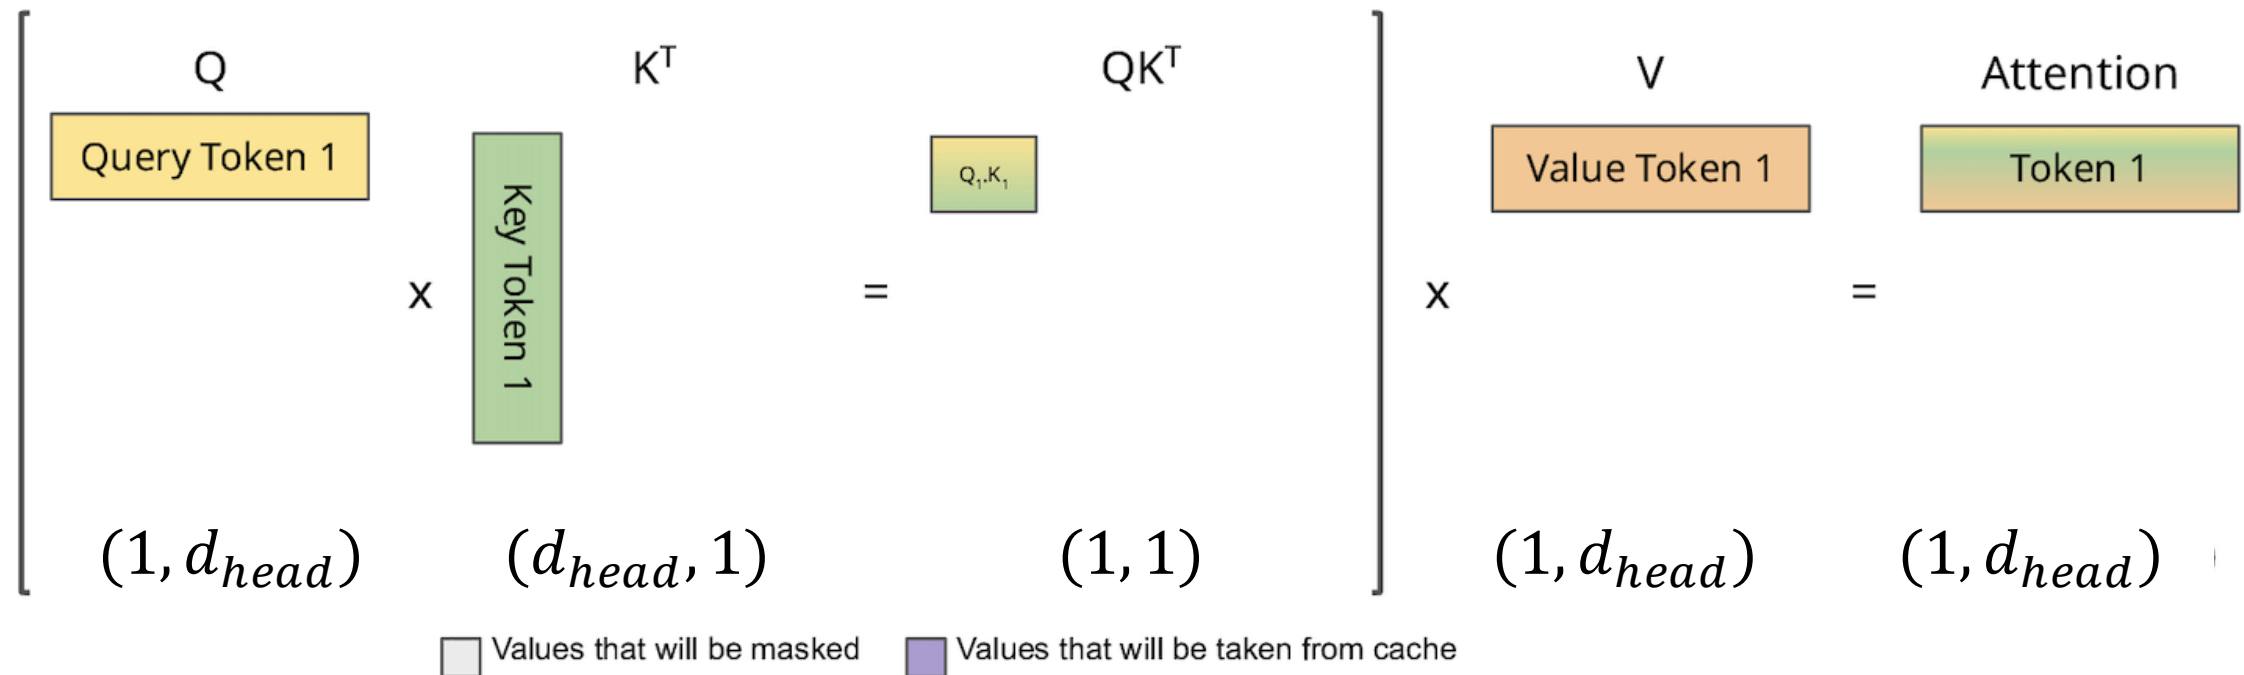
---

- Function
  - Storing previously calculated Keys and Values
- Benefits
  - Reducing the size of the matrices involved (compute attention only for new tokens.)
  - Leading to faster matrix multiplication and overall improved efficiency.
- Drawbacks
  - Requiring **EXTRA** memory to store the KV cache

# KV cache

- With Cache

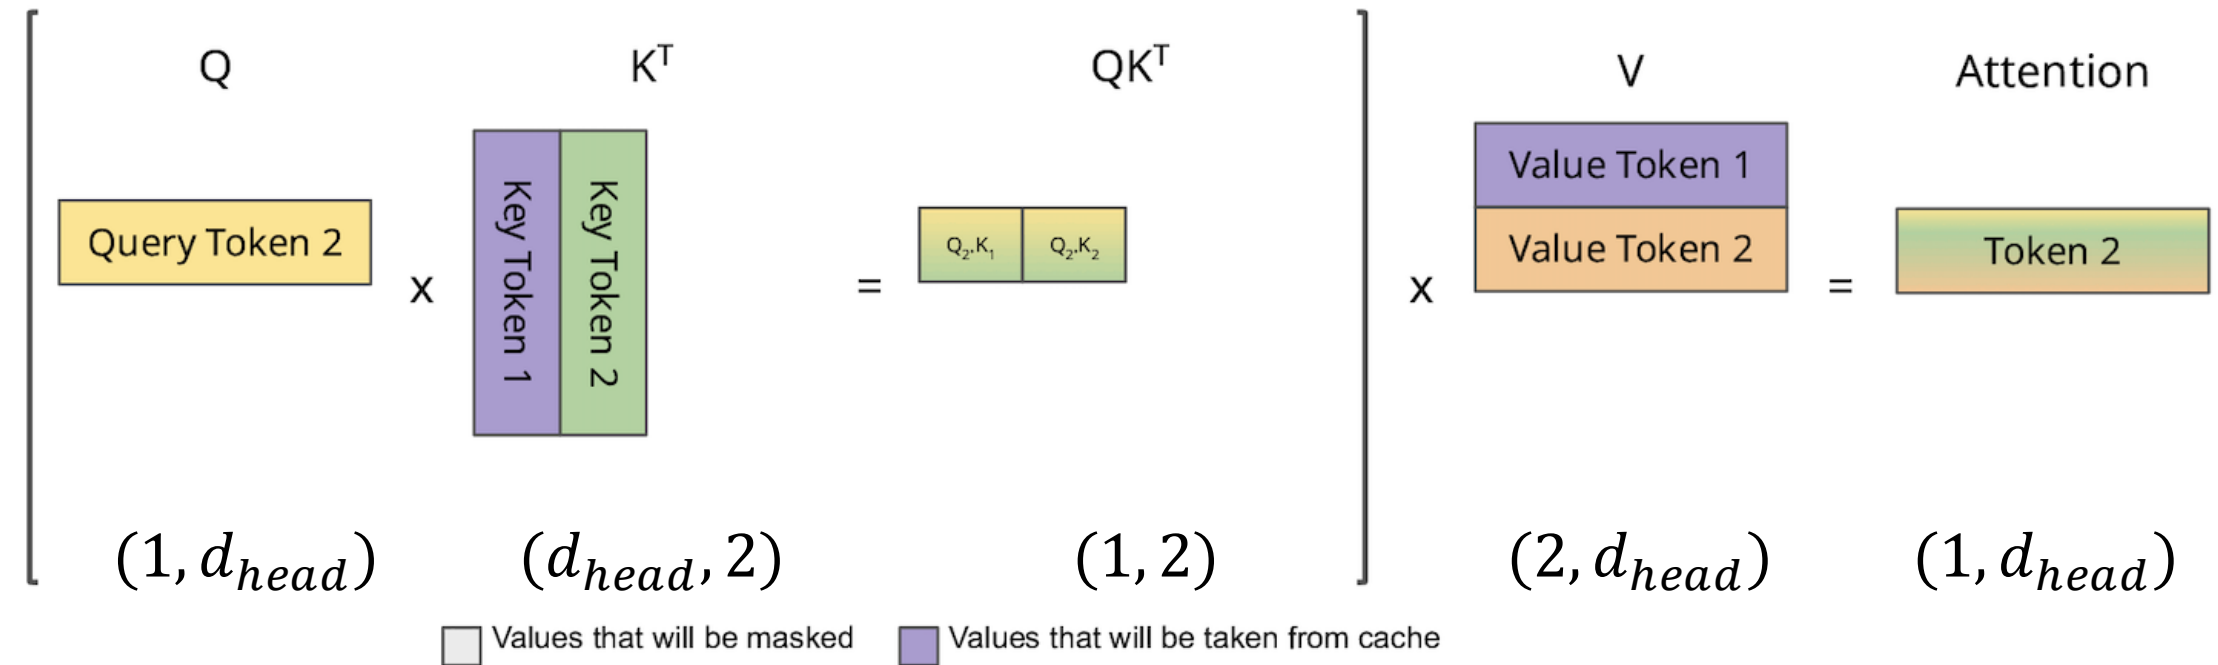
## Step 1



# KV cache

- With Cache

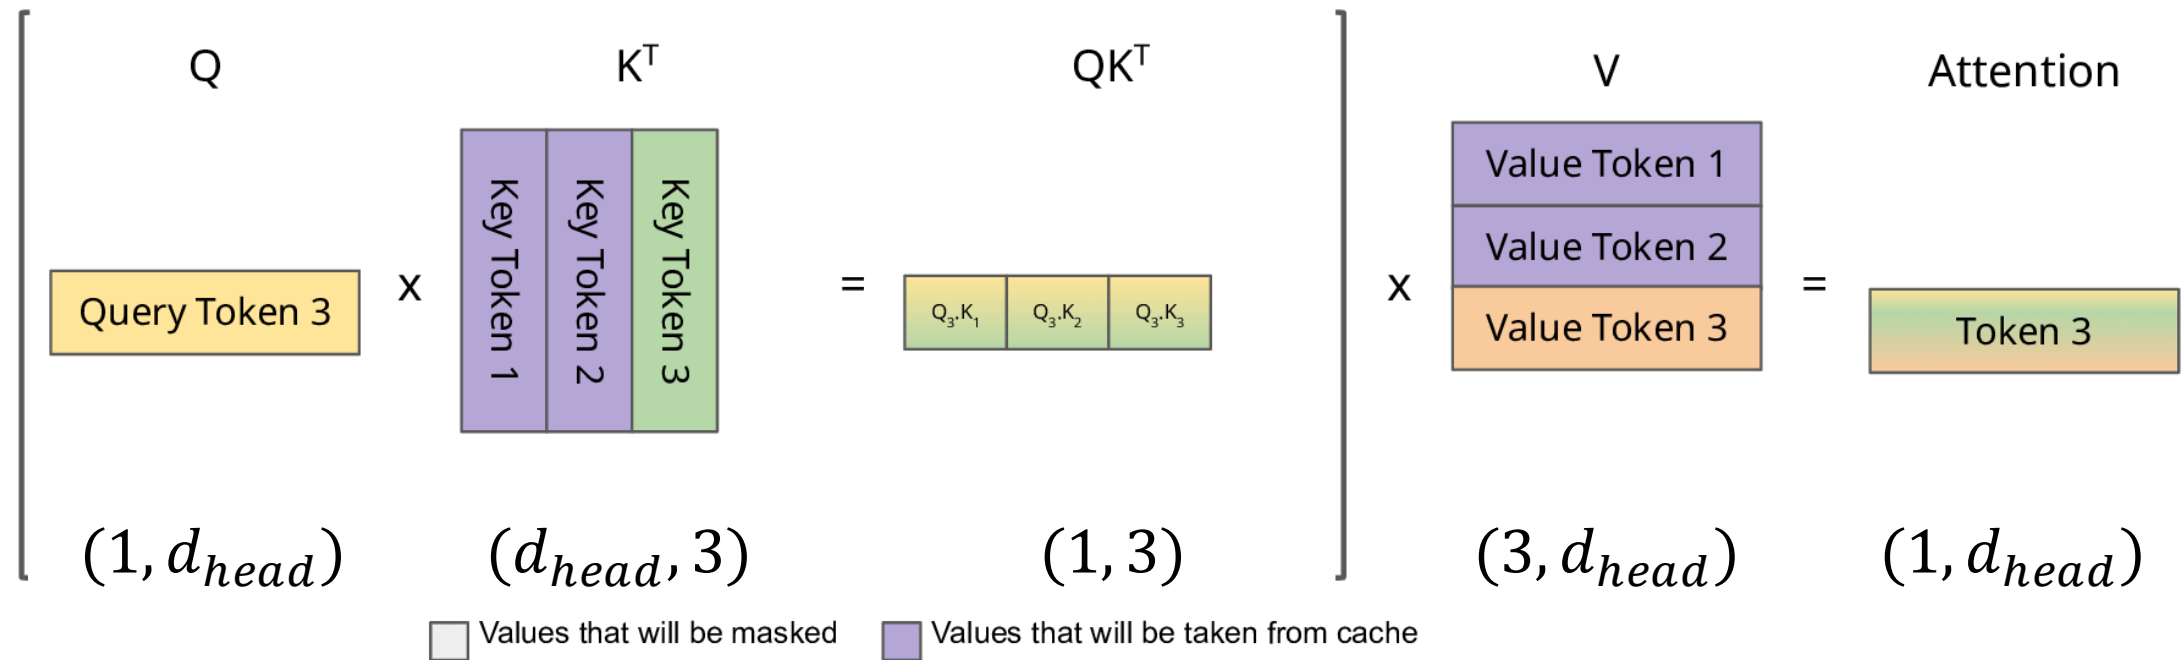
## Step 2



# KV cache

- With Cache

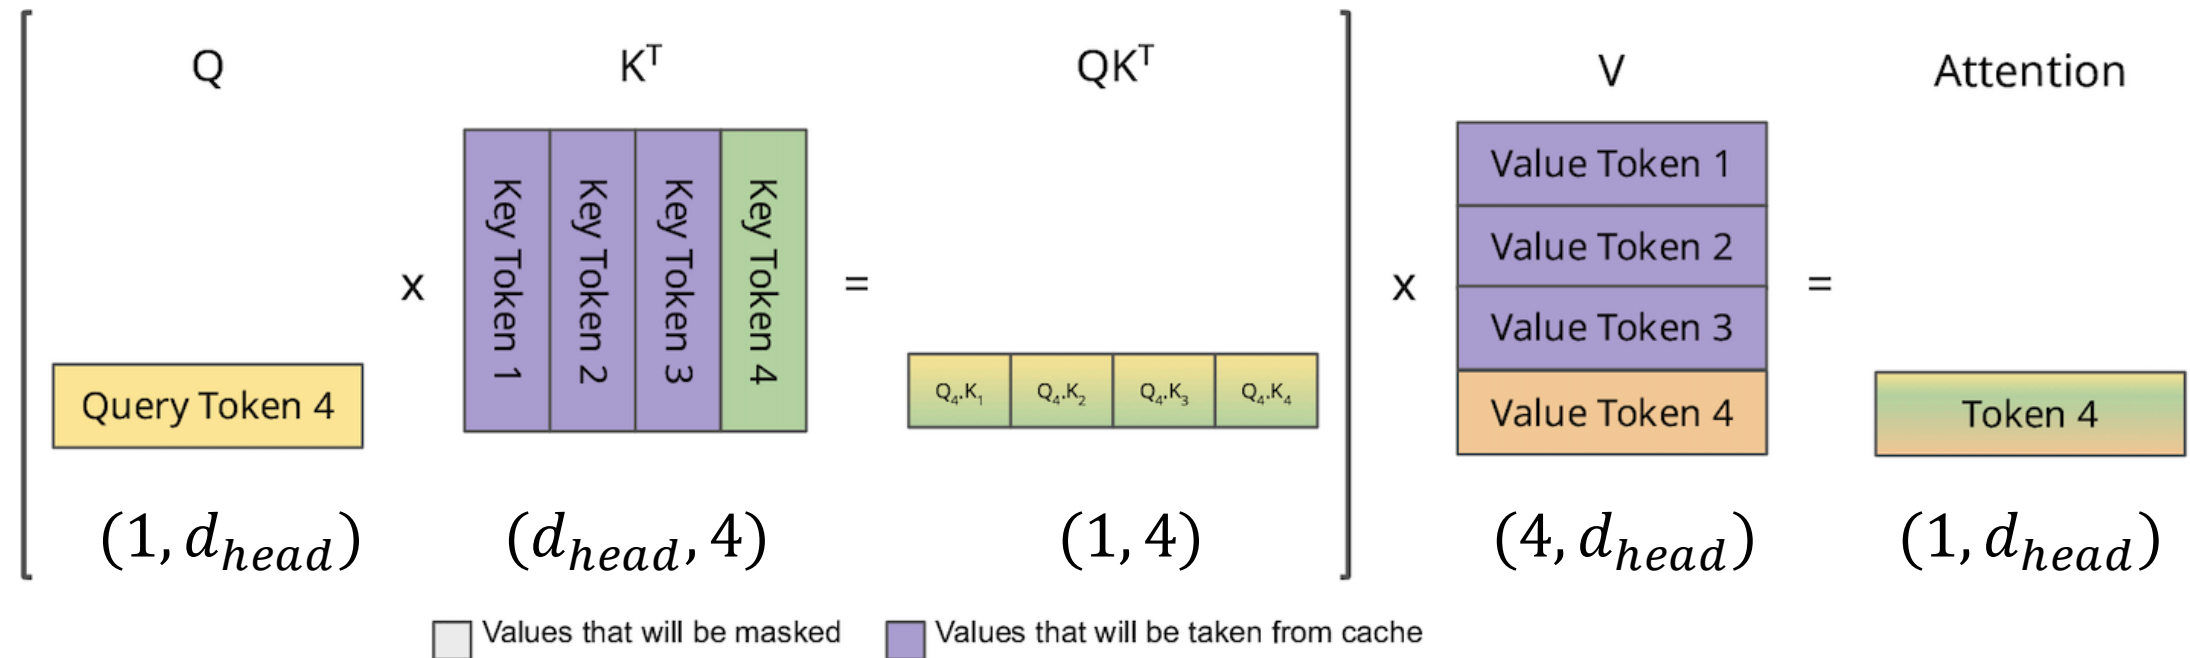
## Step 3



# KV cache

- With Cache

## Step 4





# Outline

---

## Lecture 11: Large Language Model Inference

- Inference Basic
- Additional Transformer Designs
  - KV cache
  - **Attention mechanisms optimization**
- Advanced Inference Systems
  - FlashAttention
  - vLLM
- LLM Inference Randomness

# KV cache Memory Usage

---

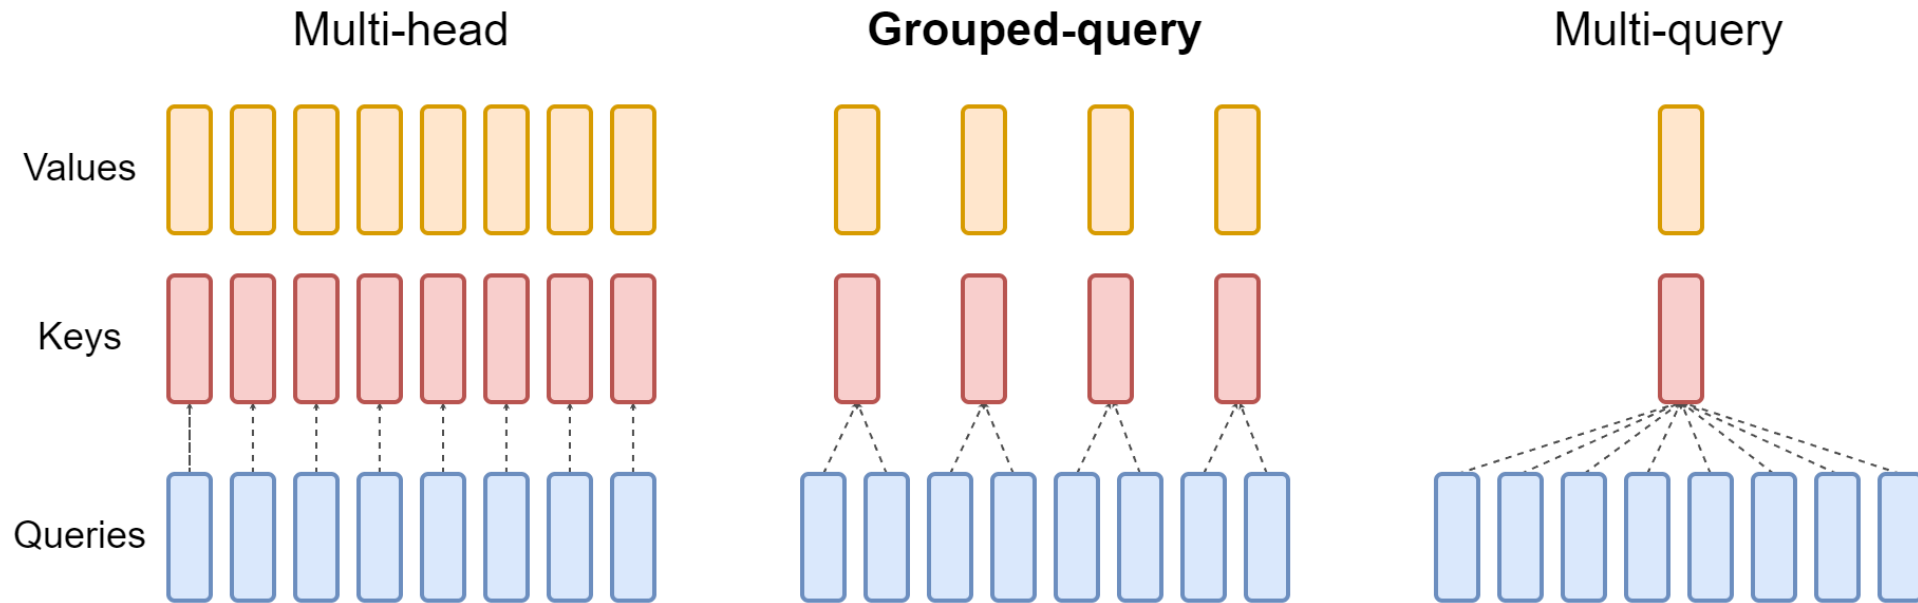
- How Big is KV cache?
  - Total size of KV cache (FP16 = 2 bytes):

$$Size_{KV} = 2 * n_{batch} * n_{seq} * n_{layers} * (n_{heads} * d_{head}) * 2(bytes)$$

- $n_{batch}$ : batch size,
  - $n_{seq}$ : total sequence length
  - $n_{layers}$ : the number of decoder attention layers,
  - $n_{heads}$ : the number of attention heads per attention layer
  - $d_{head}$ : the hidden dimension of the attention head
  - $(n_{heads} * d_{head})$ : generally called embedding dimension “ $d$ ”
- Llama-2-7B:  $n_{layers} = 32, n_{heads} = 32, d_{head} = 128$ 
    - $n_{batch} = 1, n_{seq} = 100 \rightarrow Size_{KV} = 0.05GB$
    - $n_{batch} = 16, n_{seq} = 100 \rightarrow Size_{KV} = 0.8GB$
    - $n_{batch} = 16, n_{seq} = 10000 \rightarrow Size_{KV} = 80GB$
- A100 GPU Memory = 80GB**

# Attention mechanisms optimization

- Reduce KV Cache Memory Usage with  $n_{heads}$



- Multi-head Attention (MHA):  $N$  head for Query, Key, Value
- Grouped-query attention (GQA):  $N$  head for Query,  $G$  head for Key and Value
- Multi-query attention (MQA):  $N$  head for Query, 1 head for Key and Value

# Outline

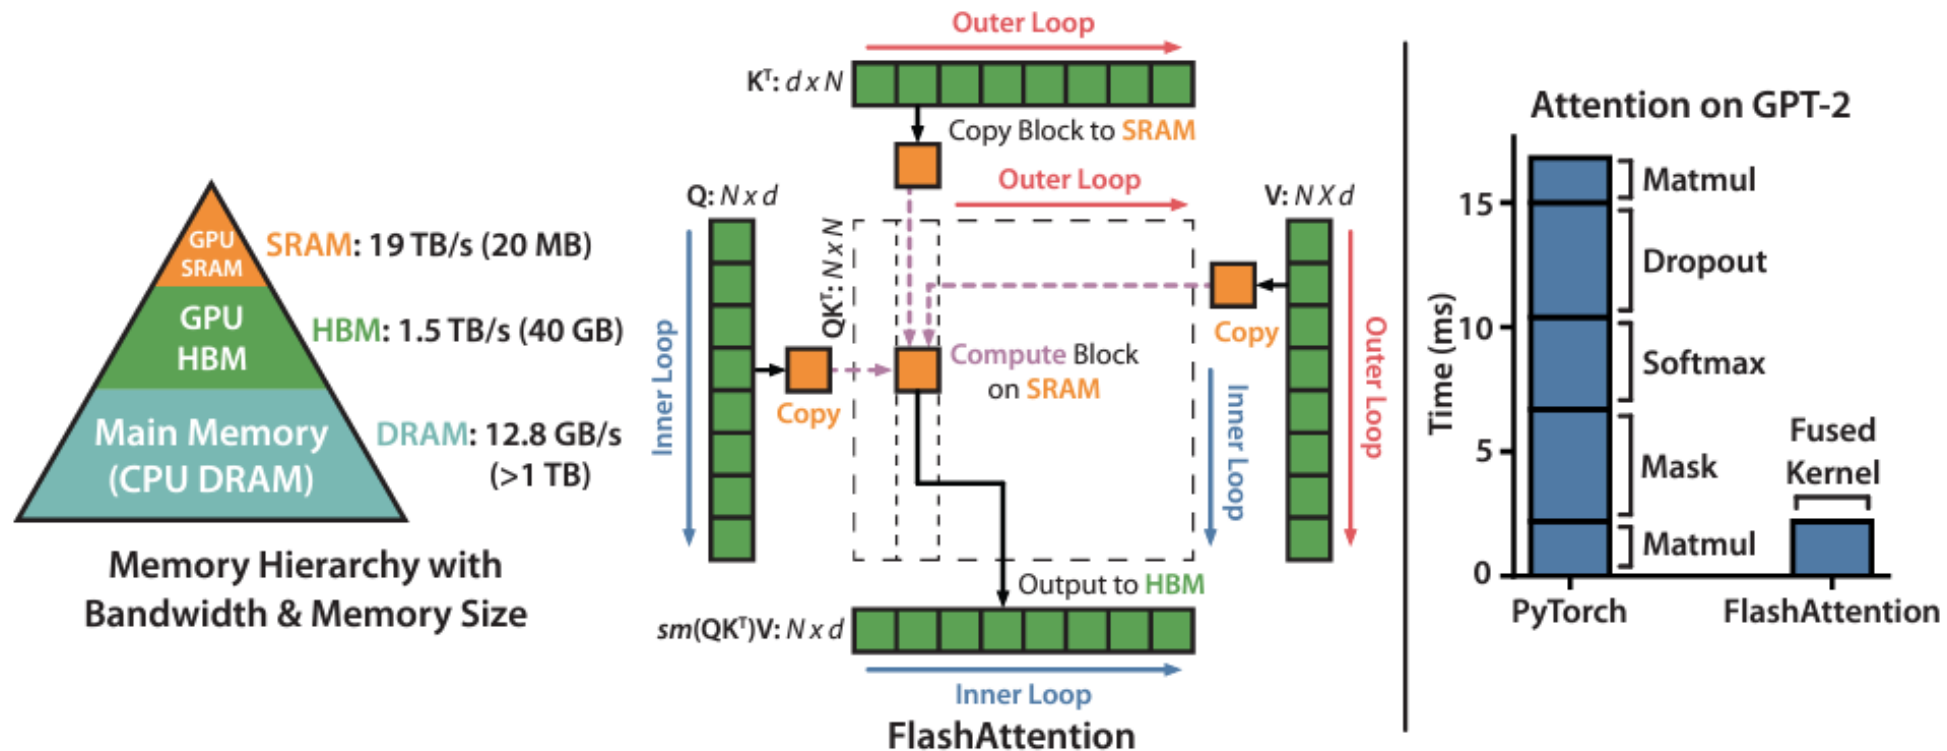
---

## Lecture 11: Large Language Model Inference

- Inference Basic
- Additional Transformer Designs
  - KV cache
  - Attention mechanisms optimization
- Advanced Inference Systems
  - **FlashAttention**
  - vLLM
- LLM Inference Randomness

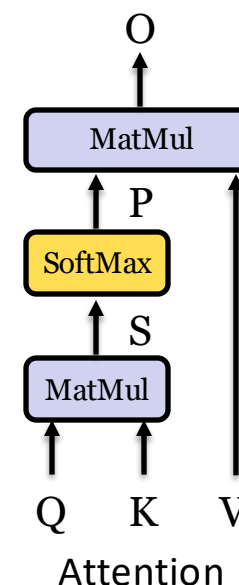
# FlashAttention

- Uses tiling to **reduce the number of memory reads/writes** between GPU high bandwidth memory (HBM) and GPU on-chip SRAM



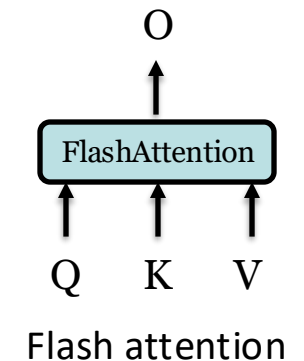
# Attention Memory Access

- Input stage
  - Load  $Q, K$  ( $n_{seq} * d$ ) from High Bandwidth Memory (HBM)
  - Calculate and write back  $S$  ( $n_{seq} * n_{seq}$ ) to HBM
- Intermediate stage
  - Load  $S$  ( $n_{seq} * n_{seq}$ ) from HBM
  - Calculate and write back  $P$  ( $n_{seq} * n_{seq}$ ) to HBM.
- Output stage
  - Load  $P$  ( $n_{seq} * n_{seq}$ ) and  $V$  ( $n_{seq} * d$ ) from HBM,
  - Calculate and write back  $O$  ( $n_{seq} * d$ ) to HBM
- HBM Memory Access Complexity
  - $\Theta(n_{seq} * d + n_{seq} * n_{seq})$



# FlashAttention Memory Access

- Key Idea
  - Breaking down the large attention matrix into smaller sub-matrices (tiles).
  - Tile fits within SRAM (access faster than HBM)
  - **Significantly reducing the need to access HBM during computation**
- Breakdown of FlashAttention
  - Load a tile (part of **Q, K, V**) from HBM into SRAM (**size: M**)
  - Perform all operations for the given tile in SRAM
  - Eliminate the need to load and write back of **S** and **P** to HBM
  - Write **O** back to HBM once the computation is complete

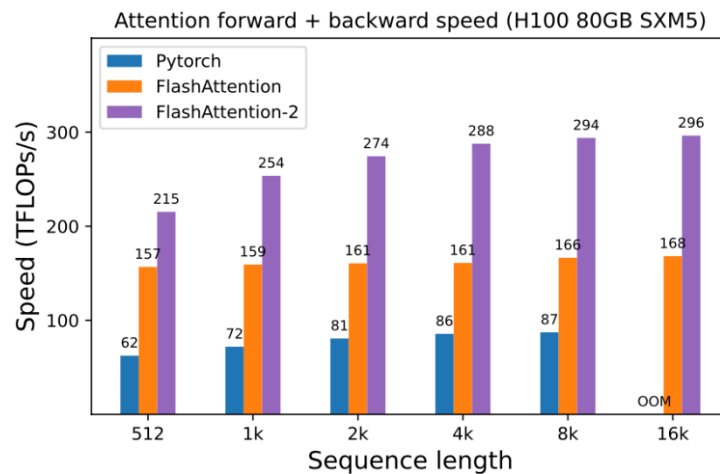


- HBM Memory Access Complexity

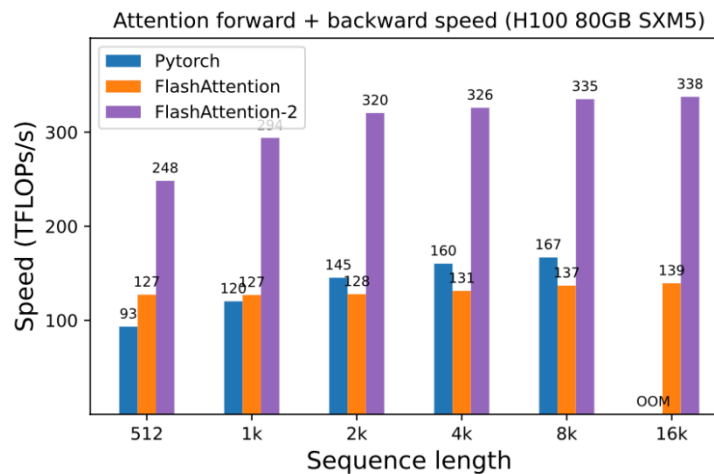
- $\Theta\left(\frac{n_{seq}^2 * d^2}{M}\right)$

# FlashAttention

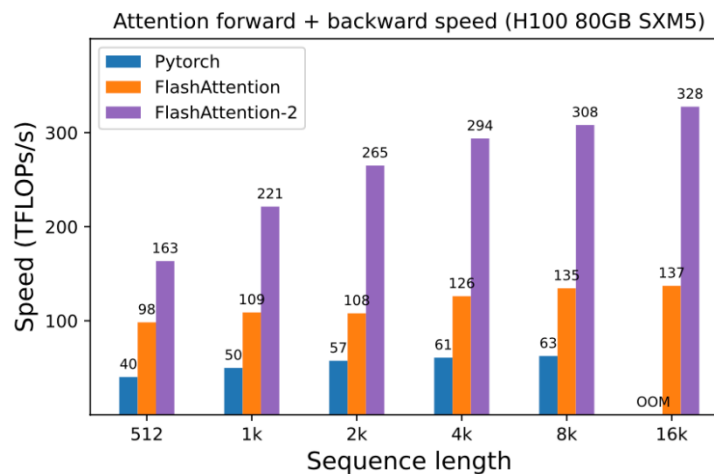
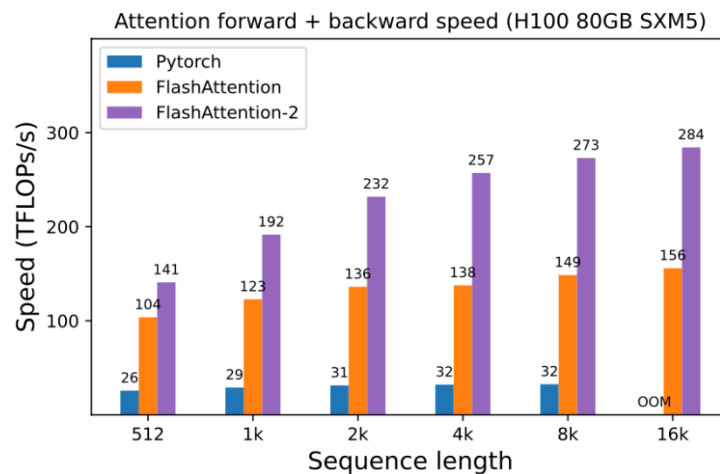
- Results on NVIDIA H100



(a) Without causal mask, head dimension 64



(b) Without causal mask, head dimension 128





# Outline

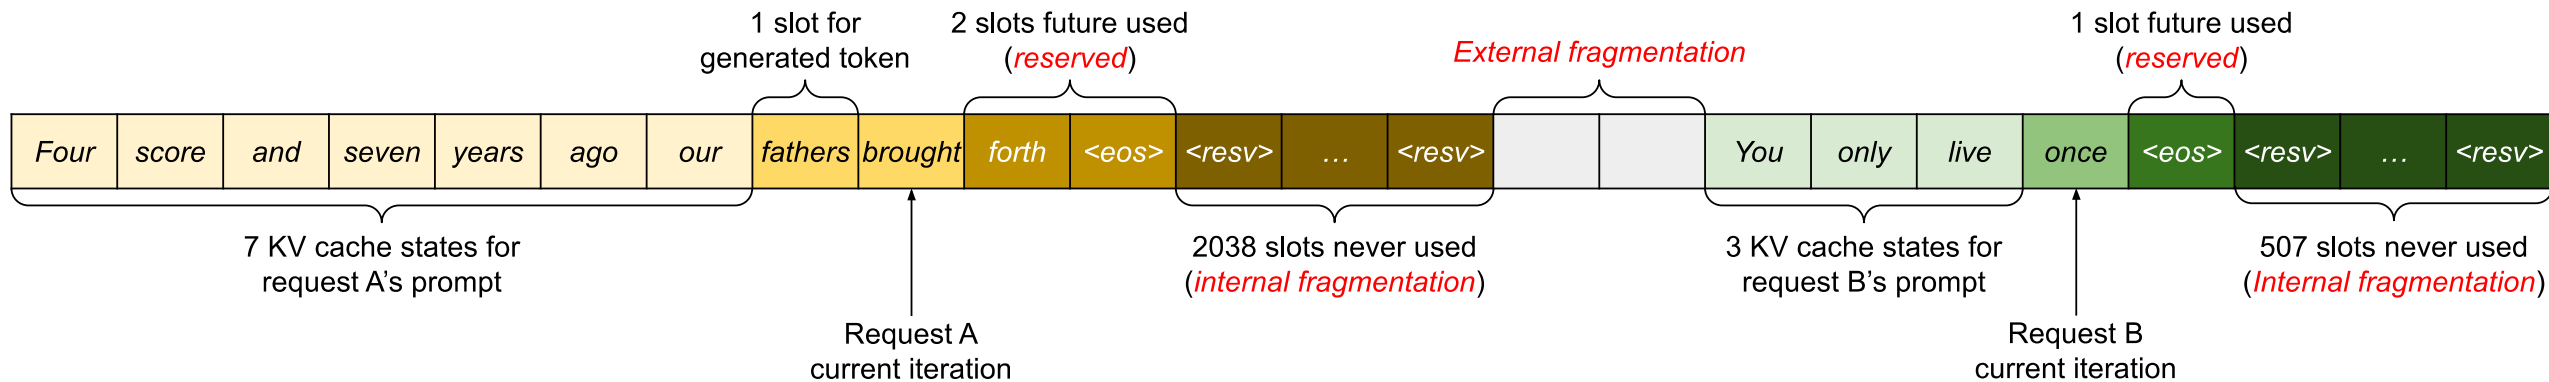
---

## Lecture 11: Large Language Model Inference

- Inference Basic
- Additional Transformer Designs
  - KV cache
  - Attention mechanisms optimization
- Advanced Inference Systems
  - FlashAttention
  - **vLLM**
- LLM Inference Randomness

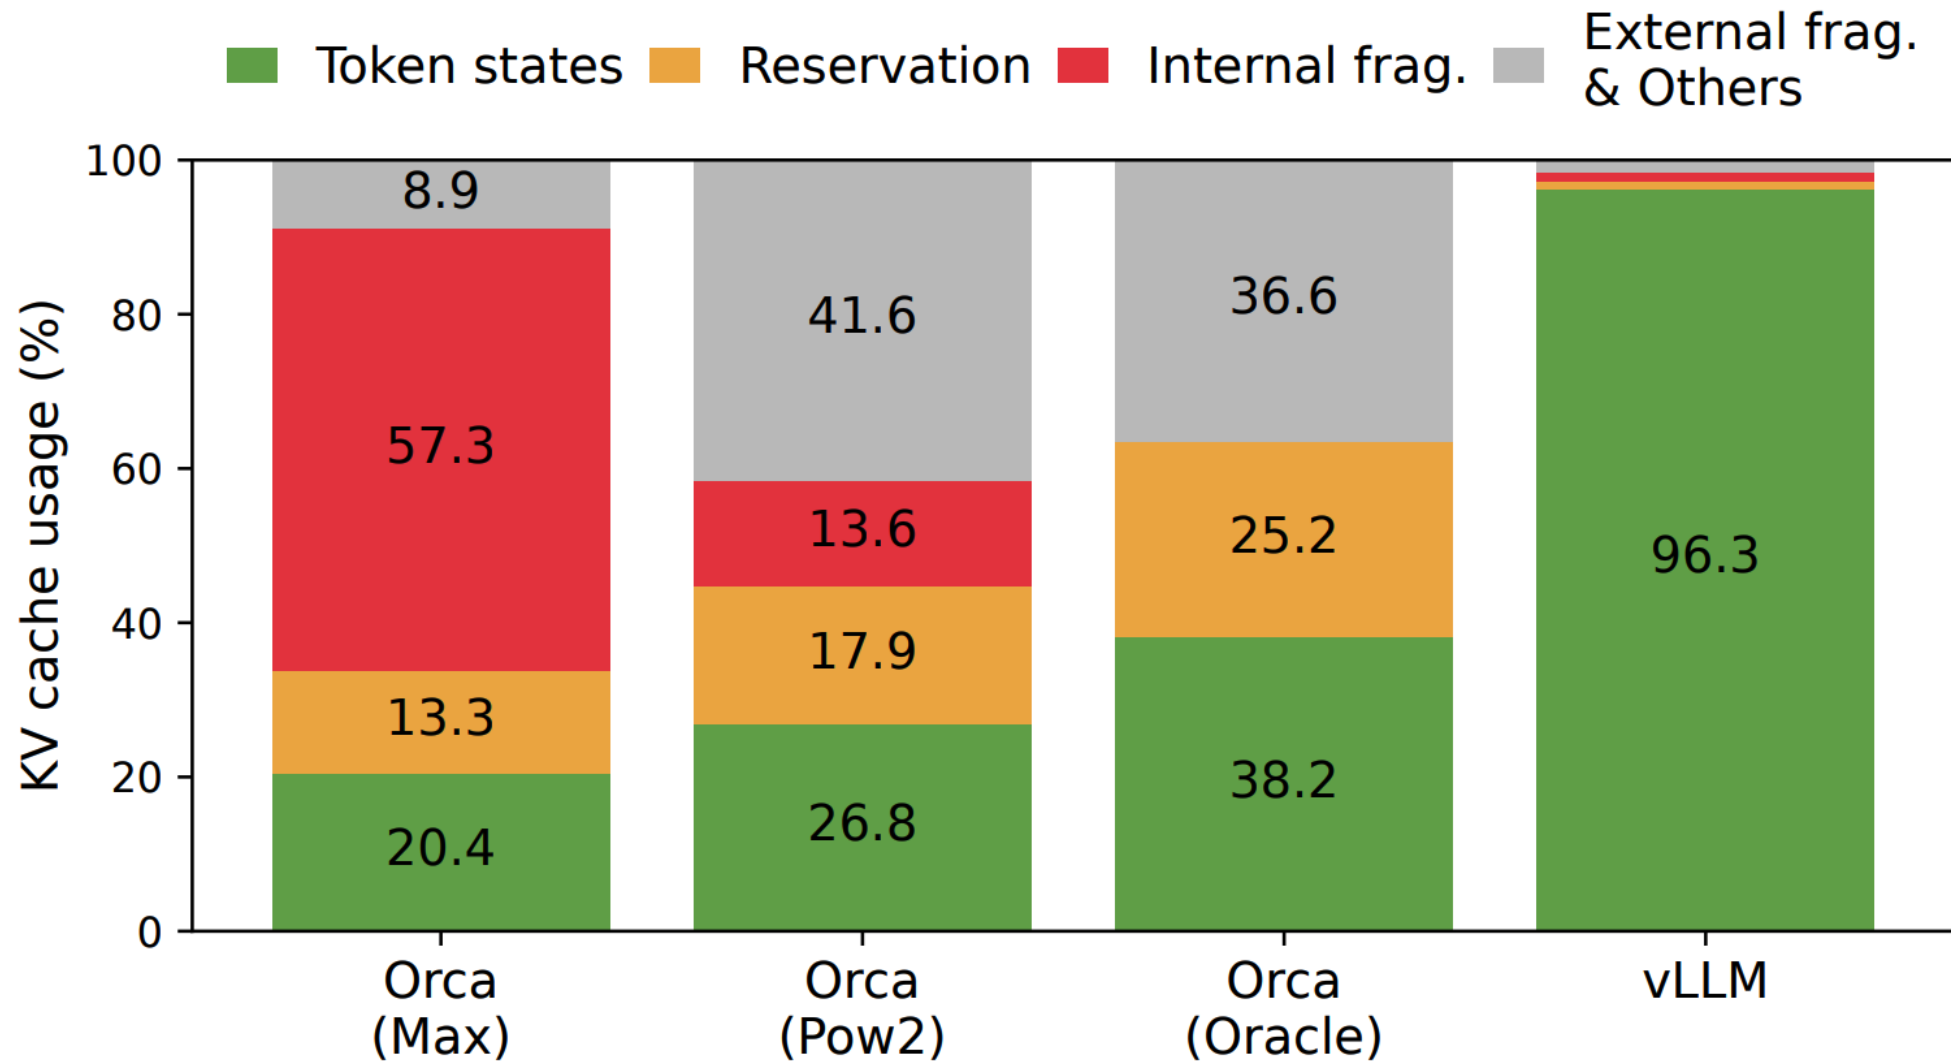
# vLLM

- A high-throughput and memory-efficient inference and serving engine for LLMs
- Motivation
  - KV cache Memory Usage problem



- **Reservation:** not used at the current step, but used in the future
- **Internal fragmentation:** over-allocated due to the unknown output length.
- **External fragmentation:** due to different sequence lengths.

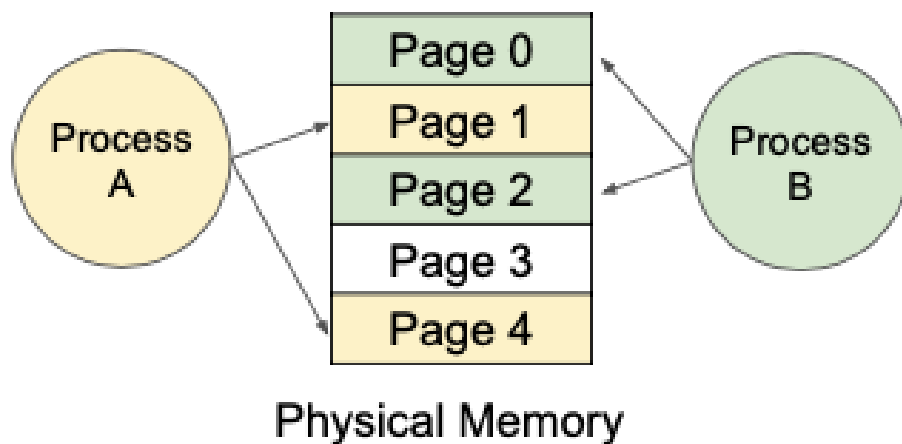
# Memory Wastes in LLM KV Cache



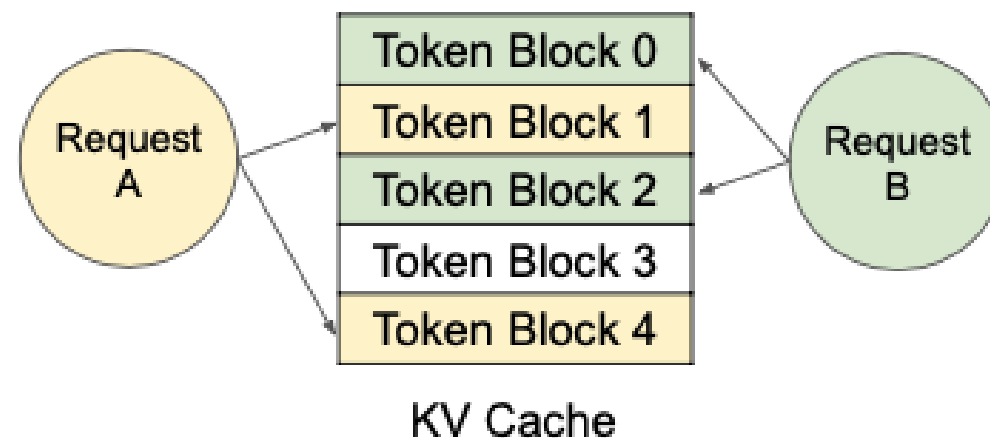
# vLLM Analogy

- Inspired by Virtual Memory Management in Operating system
- Key Algorithm: **PagedAttention**

## Memory management in OS

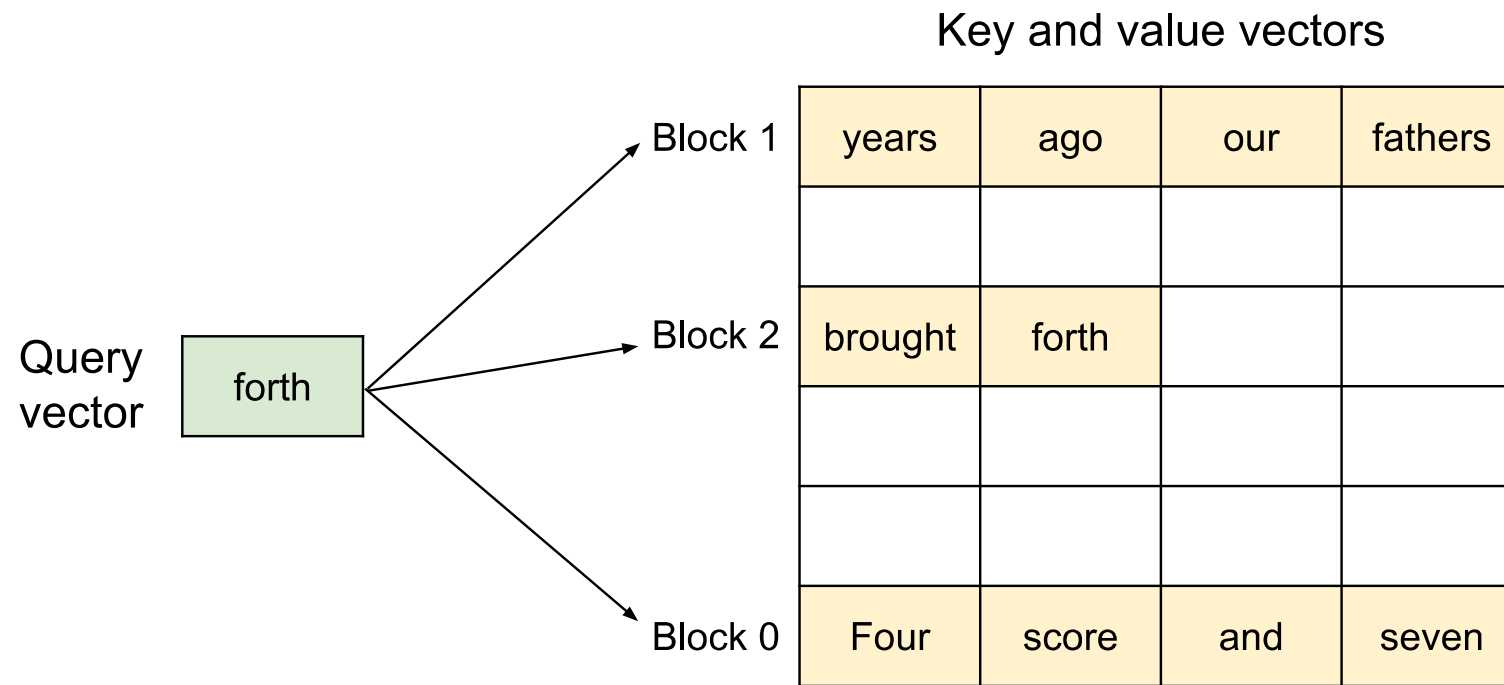


## Memory management in vLLM

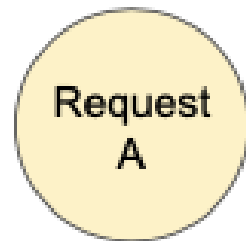


# PagedAttention

- Storing Continuous Keys and Values in non-contiguous memory space



# Logical & Physical KV blocks



Prompt: "Alan Turing is a computer scientist"

**Logical** token blocks

block 0	Alan	Turing	is	a
block 1	computer	scientist		
block 2				
block 3				

**Block table**

Physical block number	# Filled
7	4
1	2
-	-
-	-

**Physical** token blocks  
(KV Cache)

block 0				
block 1	computer	scientist		
block 2				
block 3				
block 4				
block 5				
block 6				
block 7	Alan	Turing	is	a

# Logical & Physical KV blocks

Request  
A

Prompt: "Alan Turing is a computer scientist"  
Completion: "and"

**Logical** token blocks

block 0	Alan	Turing	is	a
block 1	computer	scientist	and	
block 2				
block 3				

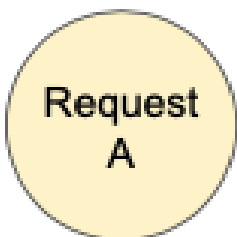
**Block table**

Physical block number	# Filled
7	4
1	3
-	-
-	-

**Physical** token blocks  
(KV Cache)

block 0				
block 1	computer	scientist	and	
block 2				
block 3				
block 4				
block 5				
block 6				
block 7	Alan	Turing	is	a

# Logical & Physical KV blocks



Prompt: "Alan Turing is a computer scientist"  
Completion: "and mathematician"

**Logical** token blocks

block 0	Alan	Turing	is	a
block 1	computer	scientist	and	mathematician
block 2				
block 3				

**Block table**

Physical block number	# Filled
7	4
1	4
-	-
-	-

**Physical** token blocks  
(KV Cache)

block 0				
block 1	computer	scientist	and	mathematician
block 2				
block 3				
block 4				
block 5				
block 6				
block 7	Alan	Turing	is	a



# Logical & Physical KV blocks

Request  
A

Prompt: "Alan Turing is a computer scientist"  
Completion: "and mathematician renowned"

**Logical** token blocks

block 0	Alan	Turing	is	a
block 1	computer	scientist	and	mathem tician
block 2	renowned			
block 3				

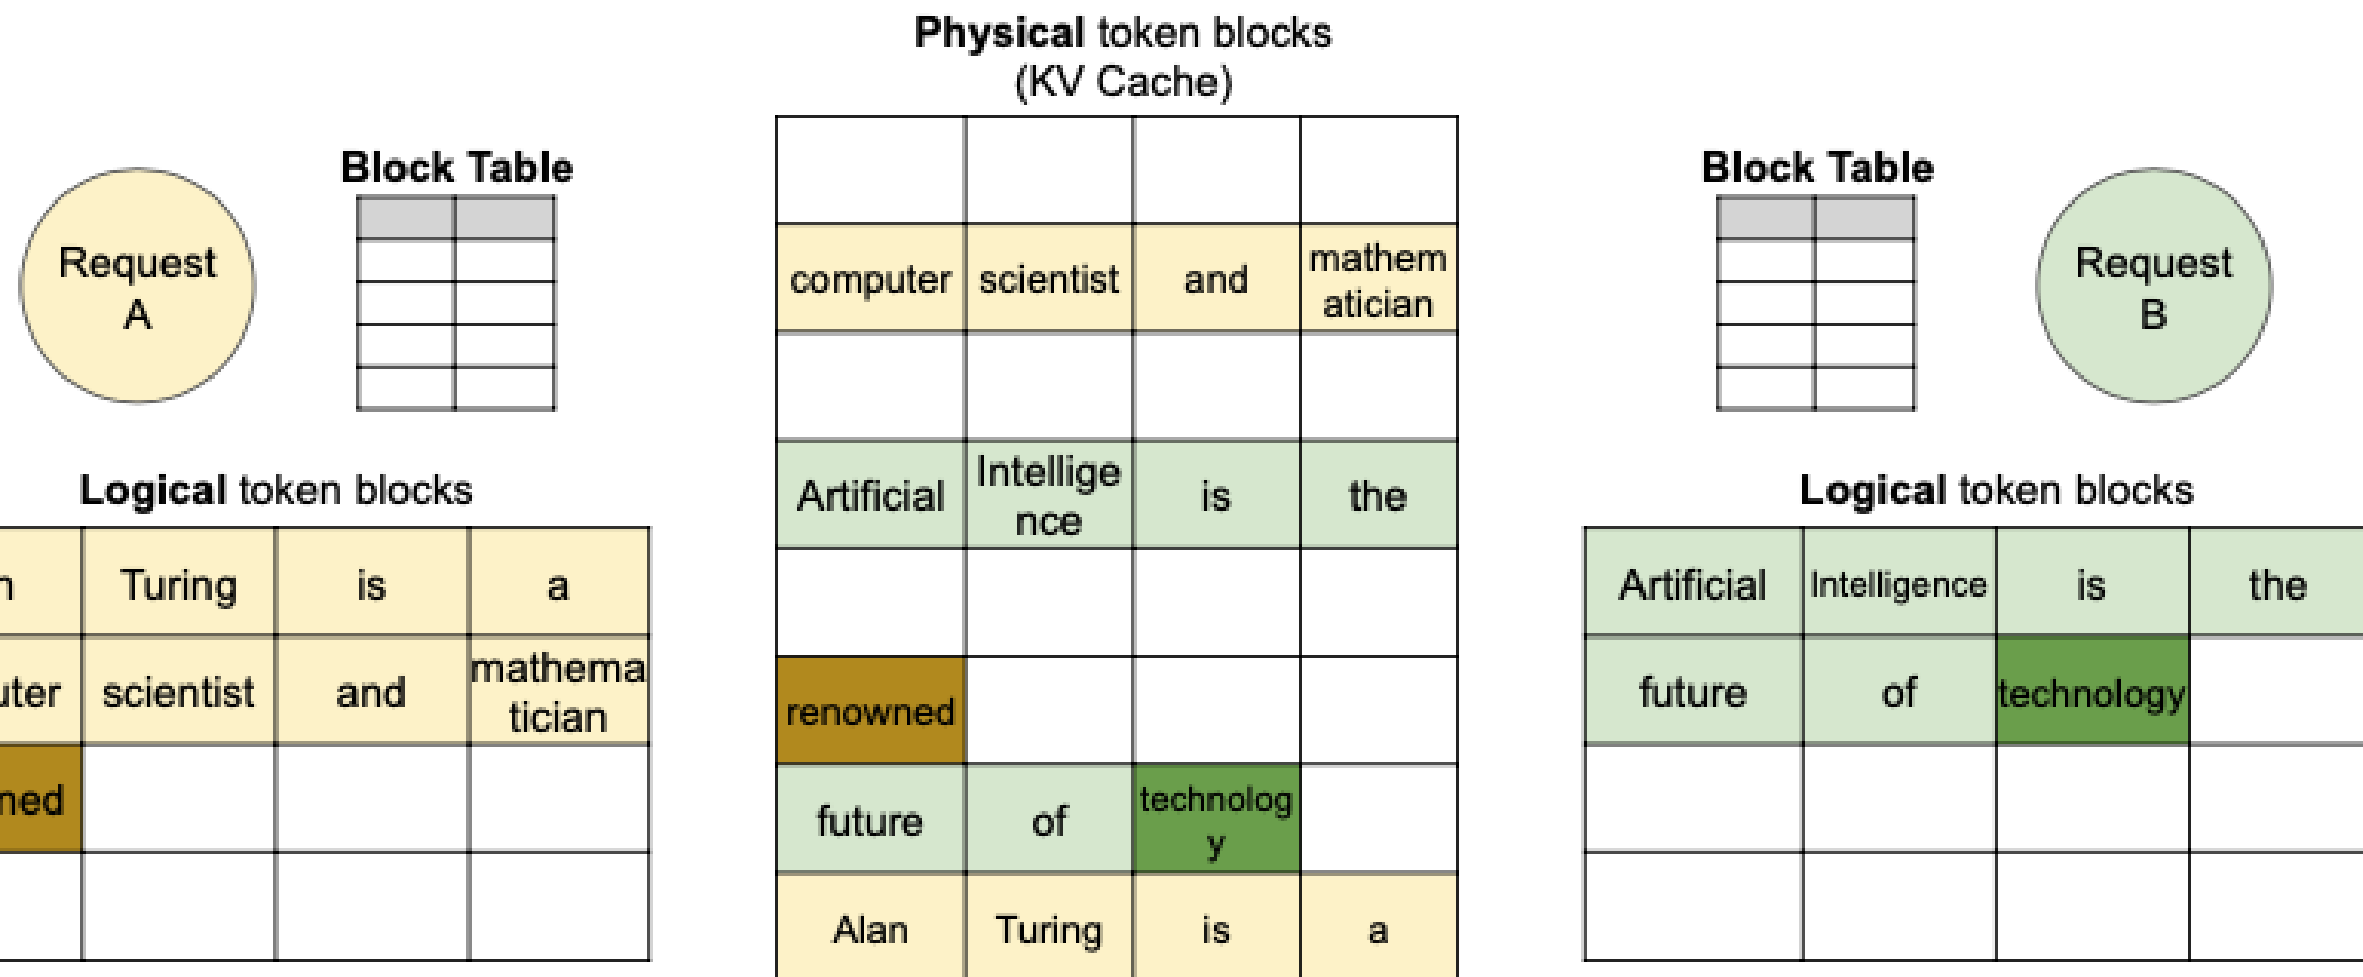
**Block table**

Physical block number	# Filled
7	4
1	4
5	1
-	-

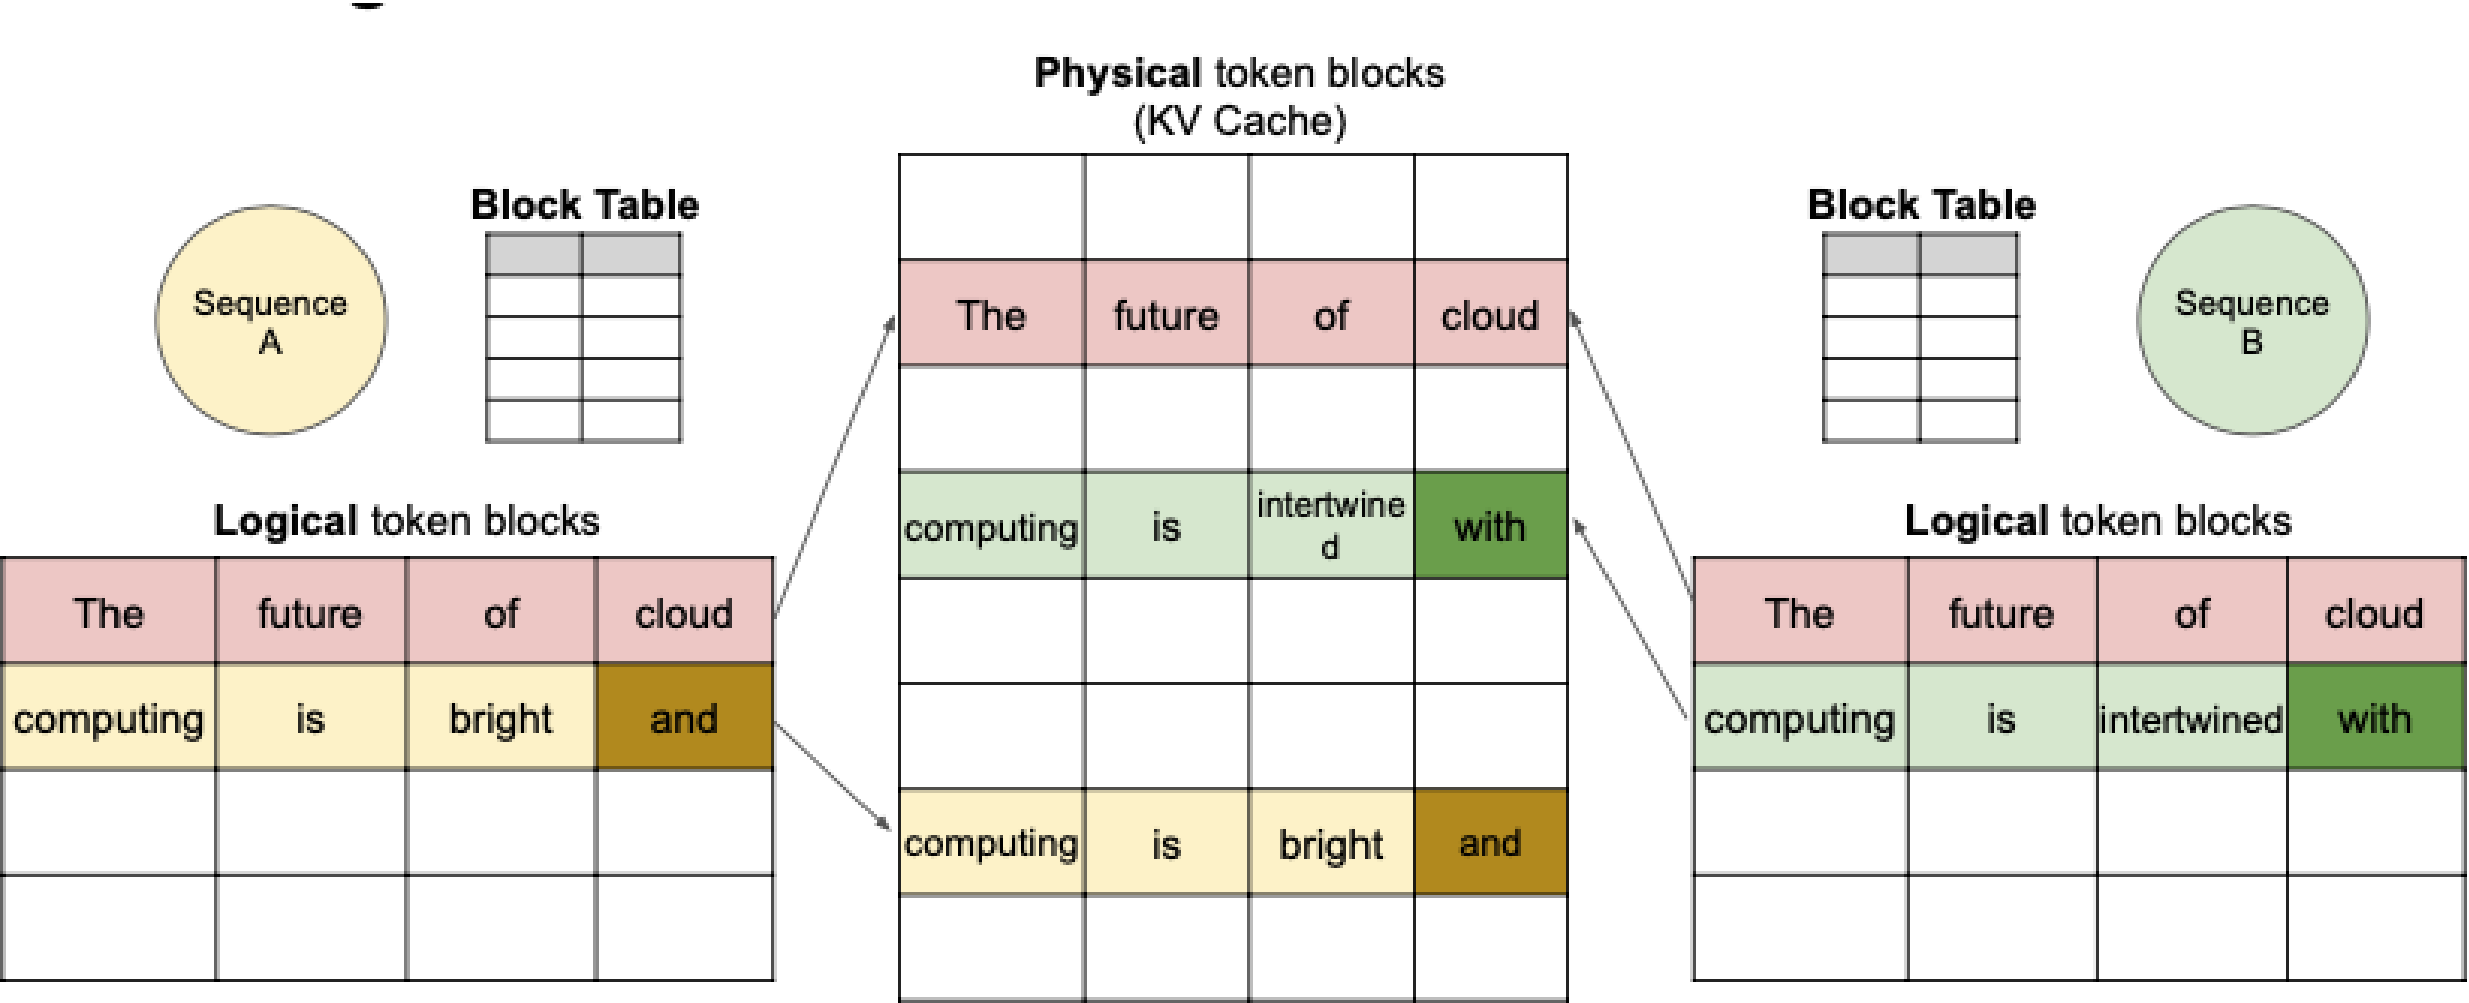
**Physical** token blocks  
(KV Cache)

block 0				
block 1	computer	scientist	and	mathem tician
block 2				
block 3				
block 4	Allocated on demand			
block 5				
block 6				
block 7	Alan	Turing	is	a

# Multiple Request Serving

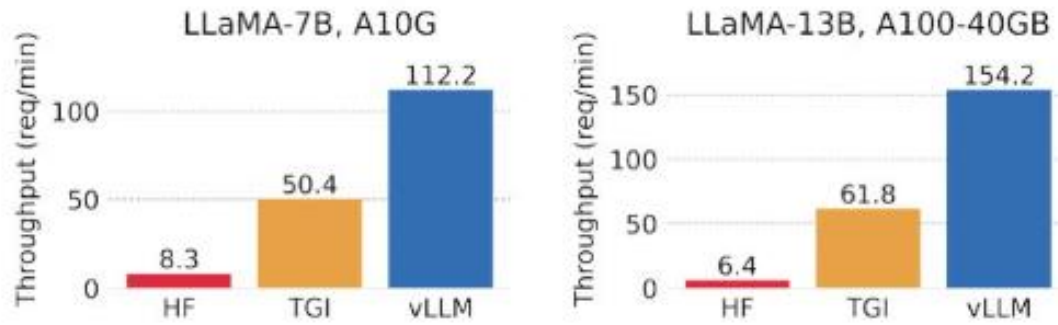


# Token Block Sharing

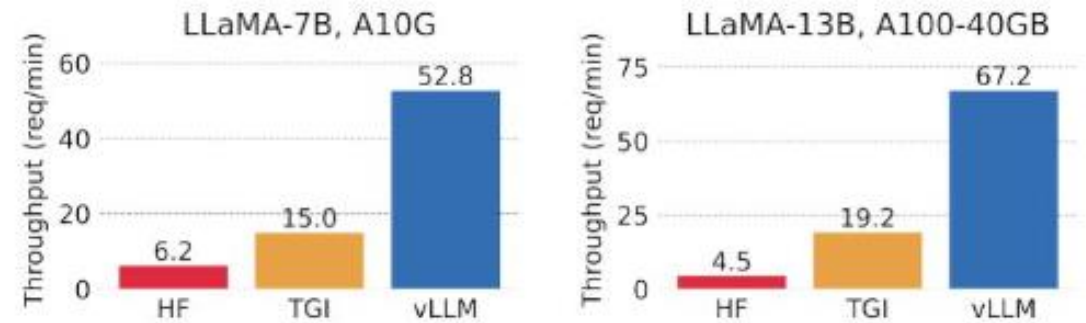


# Performance Comparison with HuggingFace and TGI

- Throughput
  - 24x higher than HuggingFace
  - 3.5x higher than Text Generation Inference (TGI)



Serving throughput when each request asks for 1 output completion.



Serving throughput when each request asks for 3 output completions.

# Outline

---

## Lecture 11: Large Language Model Inference

- Inference Basic
- Additional Transformer Designs
  - KV cache
  - Attention mechanisms optimization
- Advanced Inference Systems
  - FlashAttention
  - vLLM
- **LLM Inference Randomness**

# Temperature

---

- A crucial hyperparameter in fine-tuning the output of LLMs
  - control the randomness and creativity of generated text by adjusting word probability distributions.
- How it is Implemented
  - Softmax (same as Temperature = 1)

$$P_i = \frac{e^{x_i}}{\sum_j e^{x_j}}$$

- Softmax with Temperature (T)

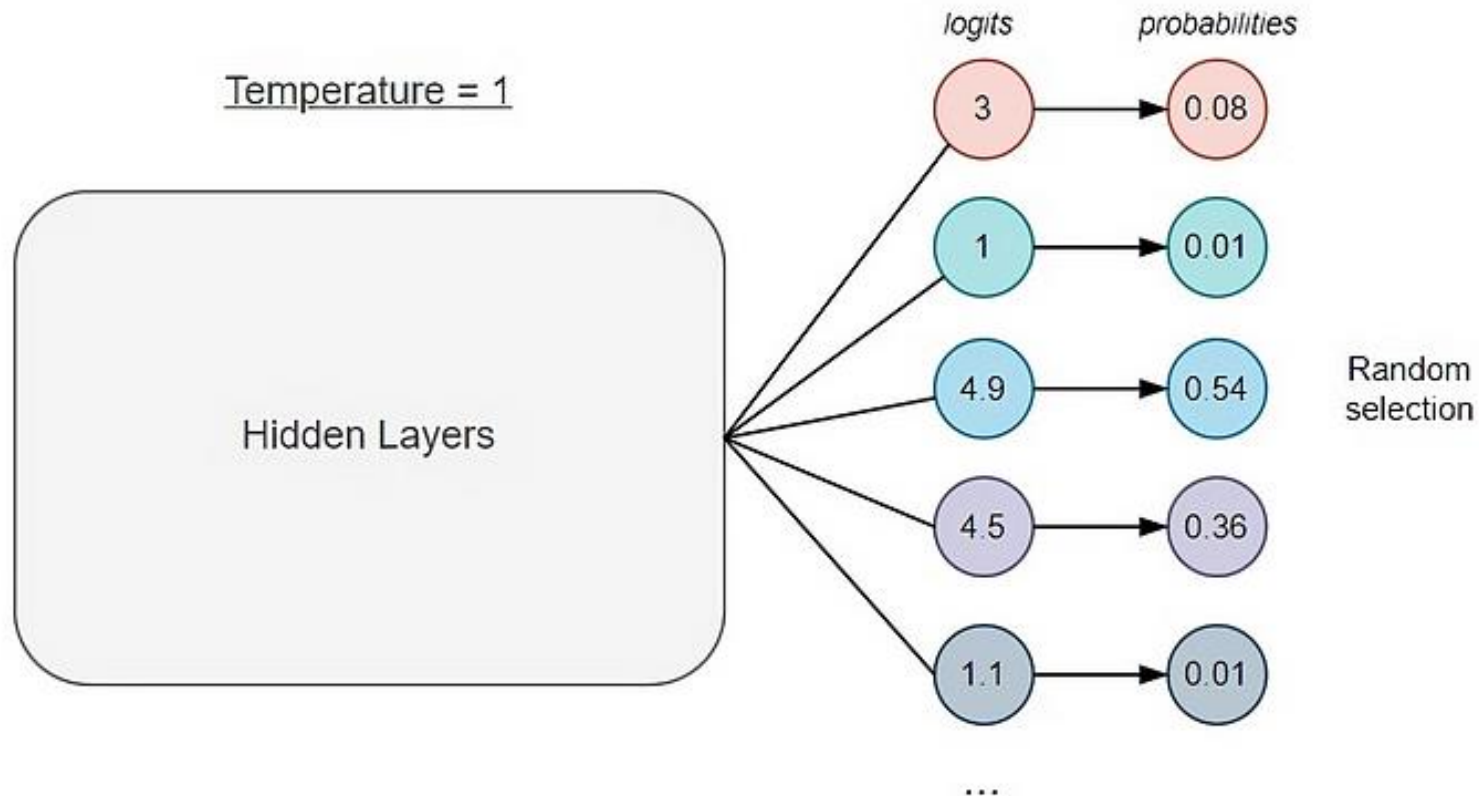
$$P_i = \frac{e^{x_i/T}}{\sum_j e^{x_j/T}}$$

# Temperature (T=10)

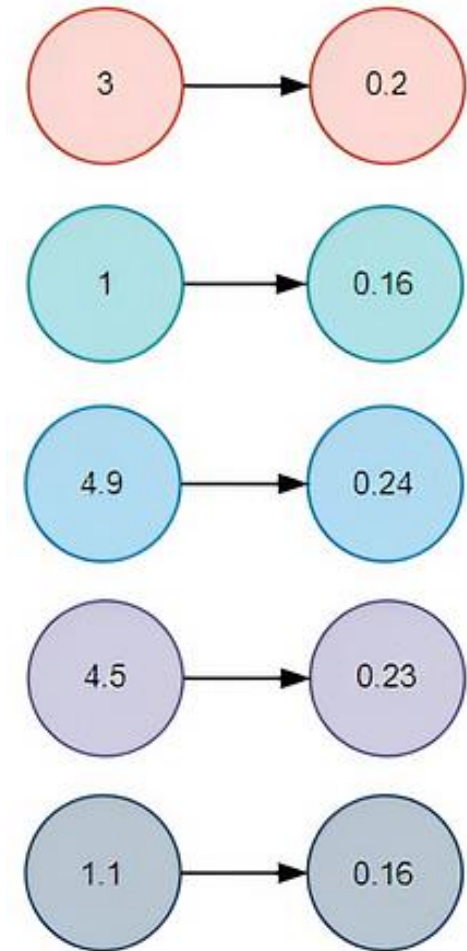
Yesterday I went to the cinema to see a \_\_\_\_



Temperature = 1



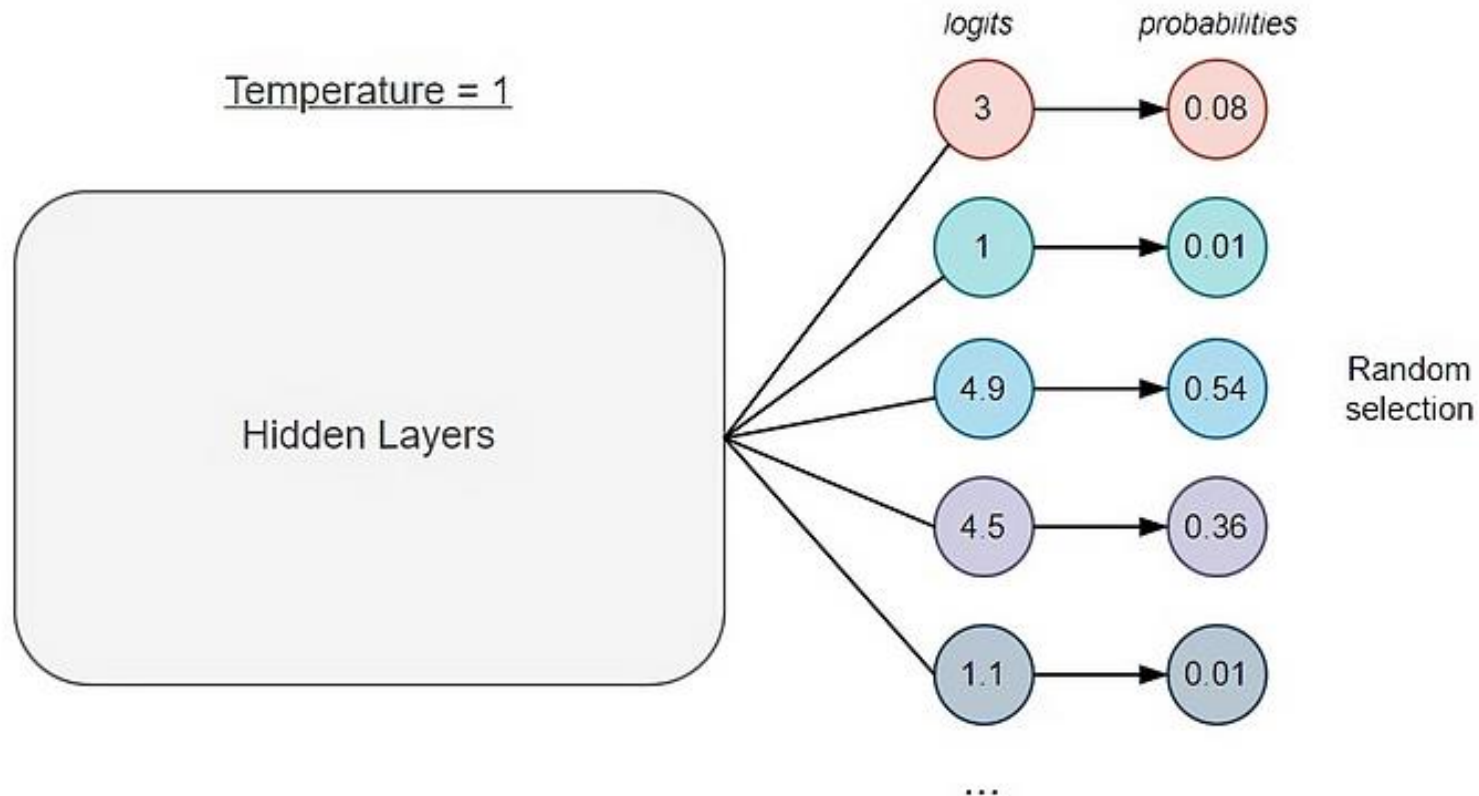
Temperature (T) = 10



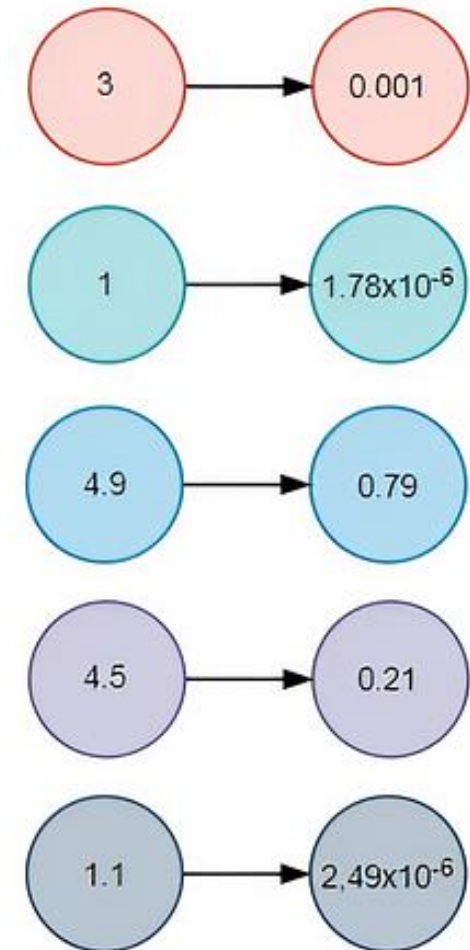
# Temperature (T=0.3)

Yesterday I went to the cinema to see a \_\_\_\_

omelette   like   film   documental   love



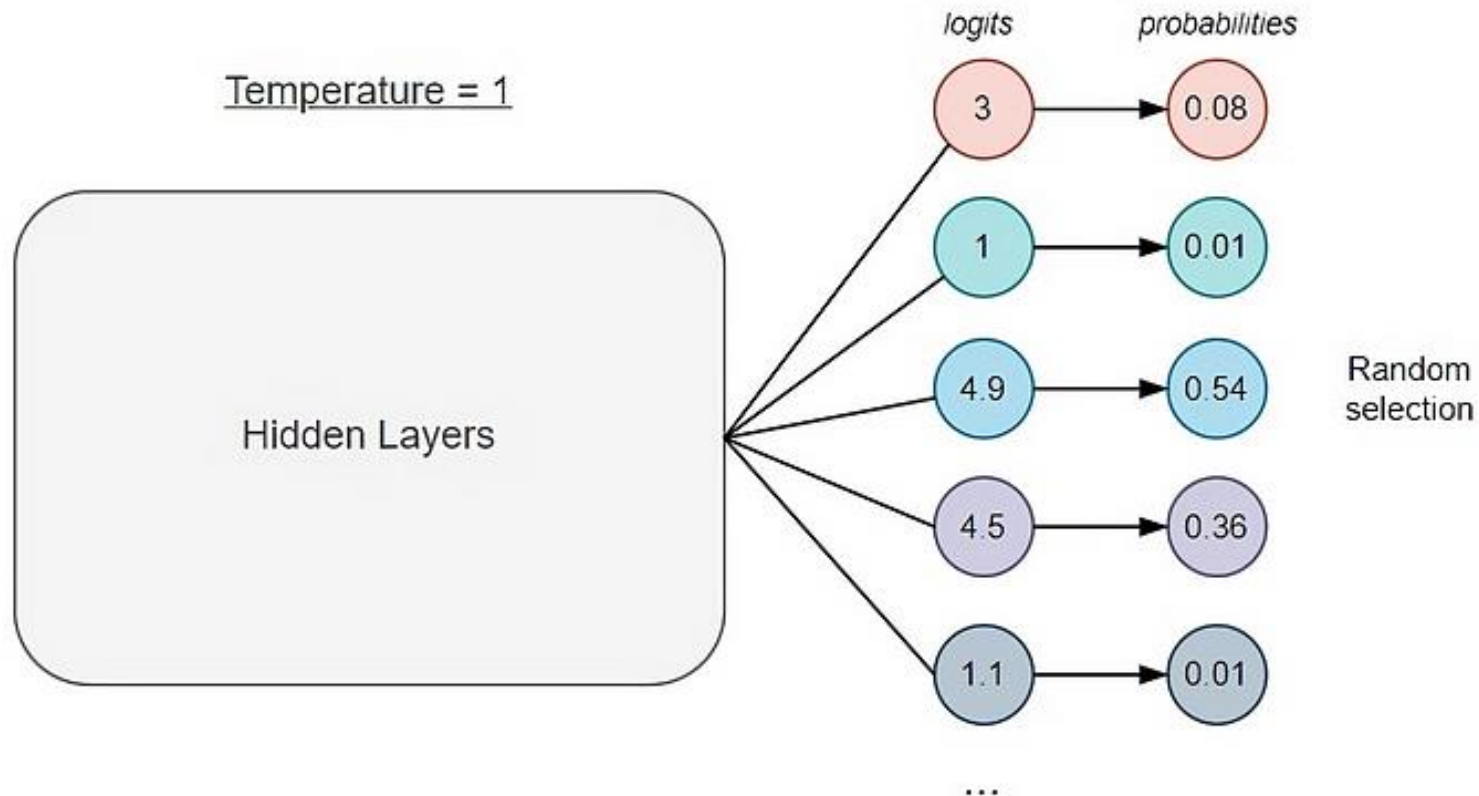
Temperature (T) = 0.3



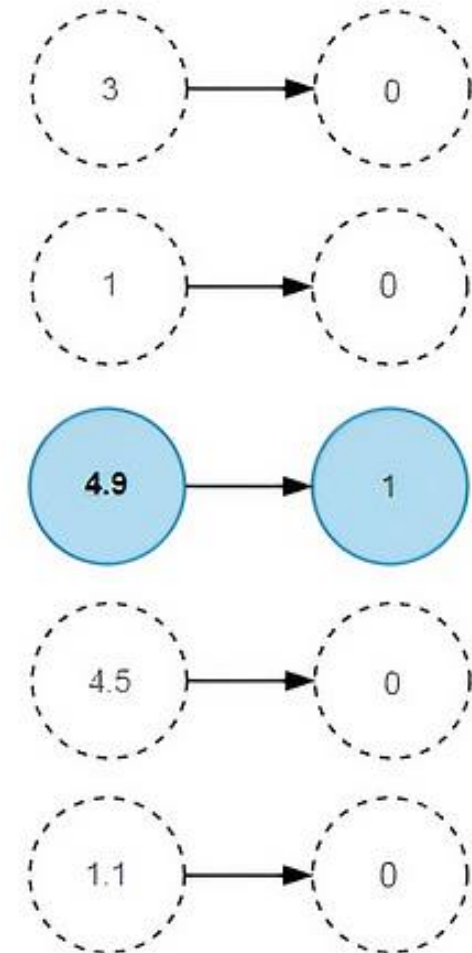


# Temperature (T=0)

Yesterday I went to the cinema to see a \_\_\_\_



Temperature (T) = 0



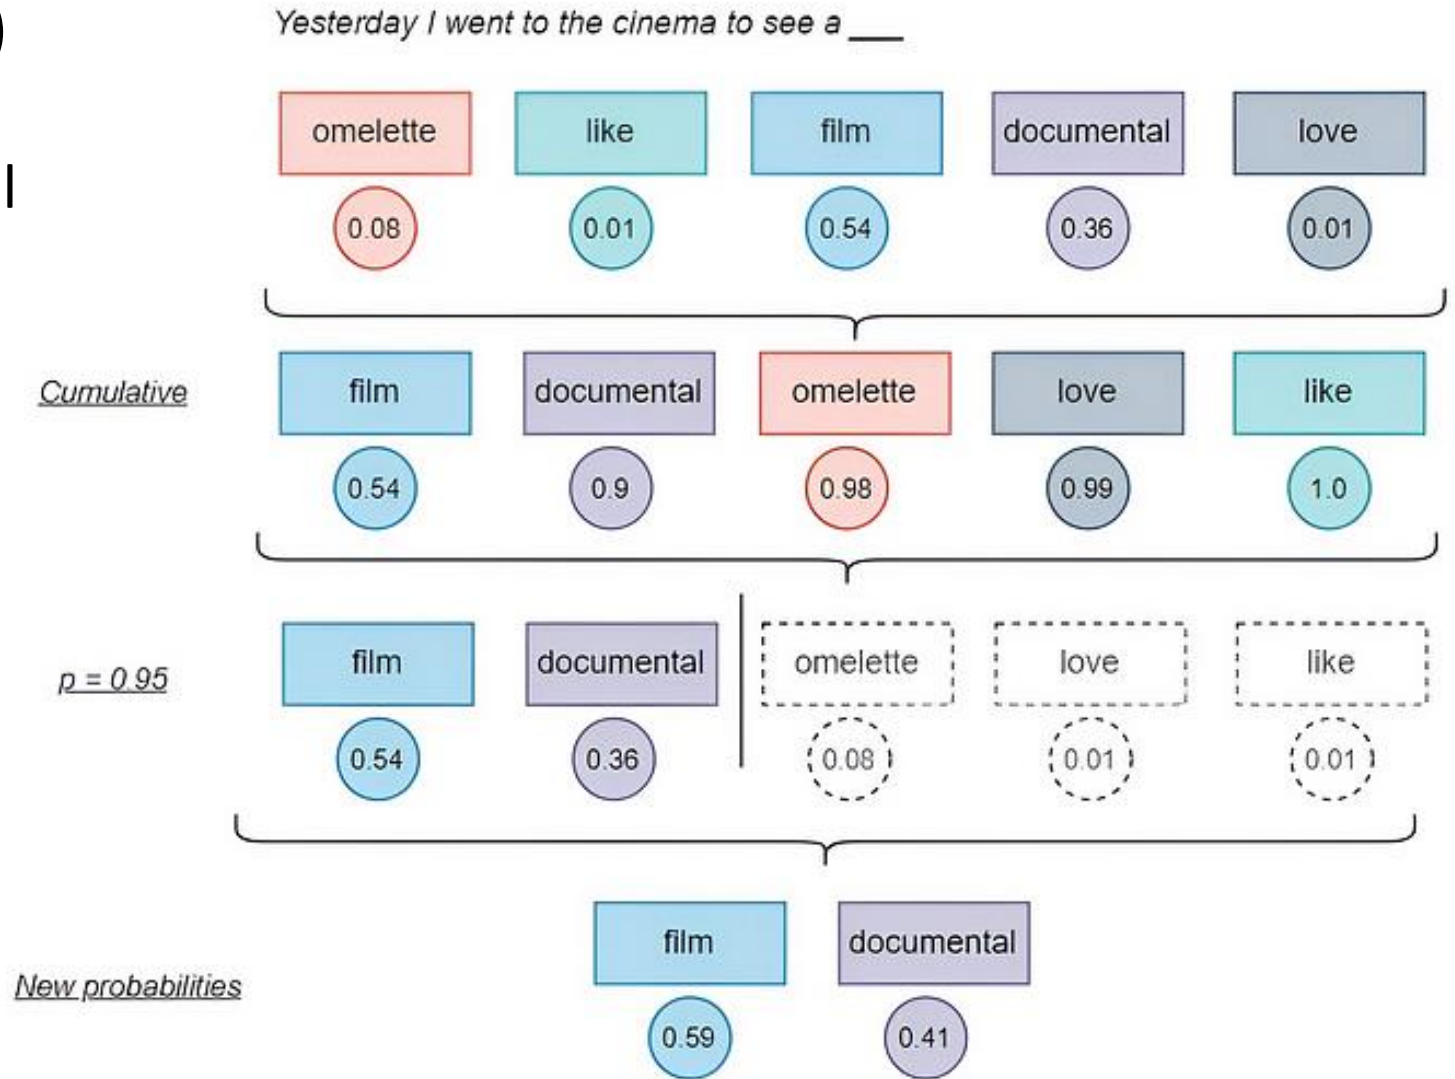
# Temperature Impact

---

- Low Temperature ( $T \rightarrow 0$ )
  - Makes the model more deterministic
  - Higher probability words become even more likely
- High Temperature ( $T > 1$ ):
  - Increases randomness.
  - Less probable words have a higher chance of selection
- Practical Usage:
  - Low T: For precise and factual responses
  - High T: For creative writing and idea generation

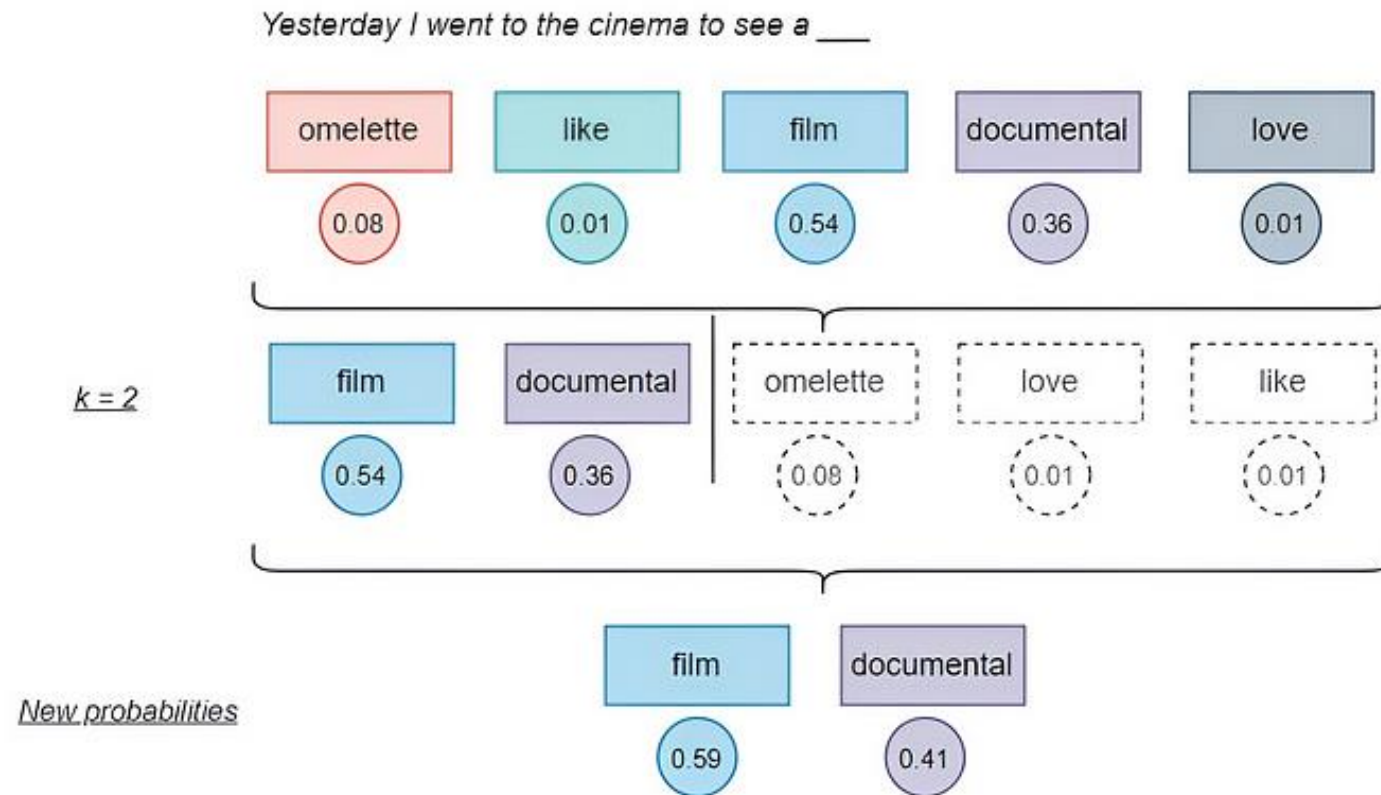
# Top\_p

- Top\_p (Nucleus Sampling)
  - Selects tokens based on cumulative probability until reaching a predefined threshold **p**.
  - Considers the smallest set of top tokens whose probabilities sum up to **p**.



# Top\_k

- Top\_k (Top\_k Sampling)
  - Considers only the top **k** most probable tokens.
  - Ignores all tokens outside the top **k** probabilities.



# In this lecture we learned

---

- LLM Inference
  - Prefill and decoding phase
- Additional Transformer Design
  - KV cache
  - Group Query Attention, Multi-query Attention
- Advanced Inference Systems
  - FlashAttention
  - vLLM
- Temperature, Top-k & Top-p