



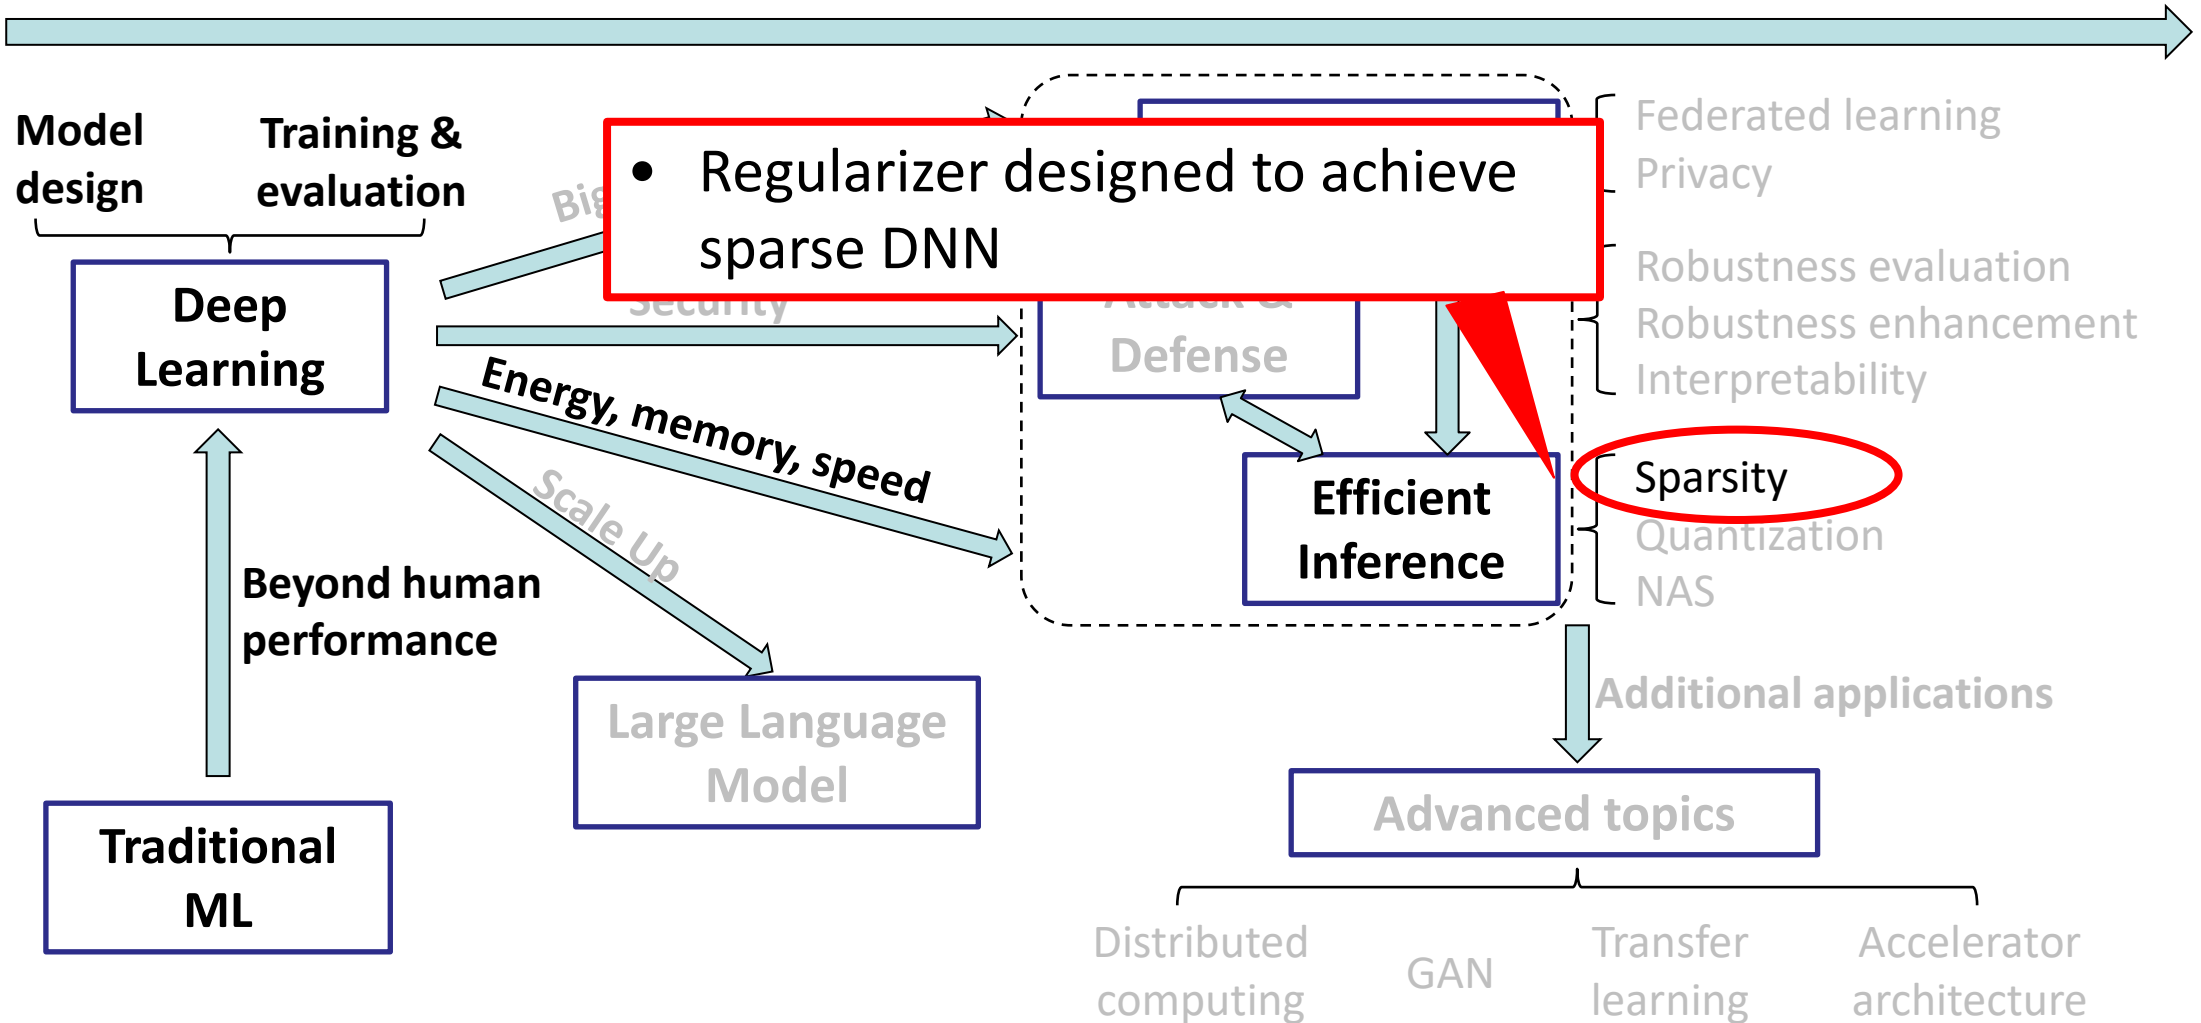
ECE 661 COMP ENG ML & DEEP NEURAL NETS

16. OPTIMIZATION-BASED METHODS FOR SPARSE DNN

HAI "HELEN" LI, SPRING 2025

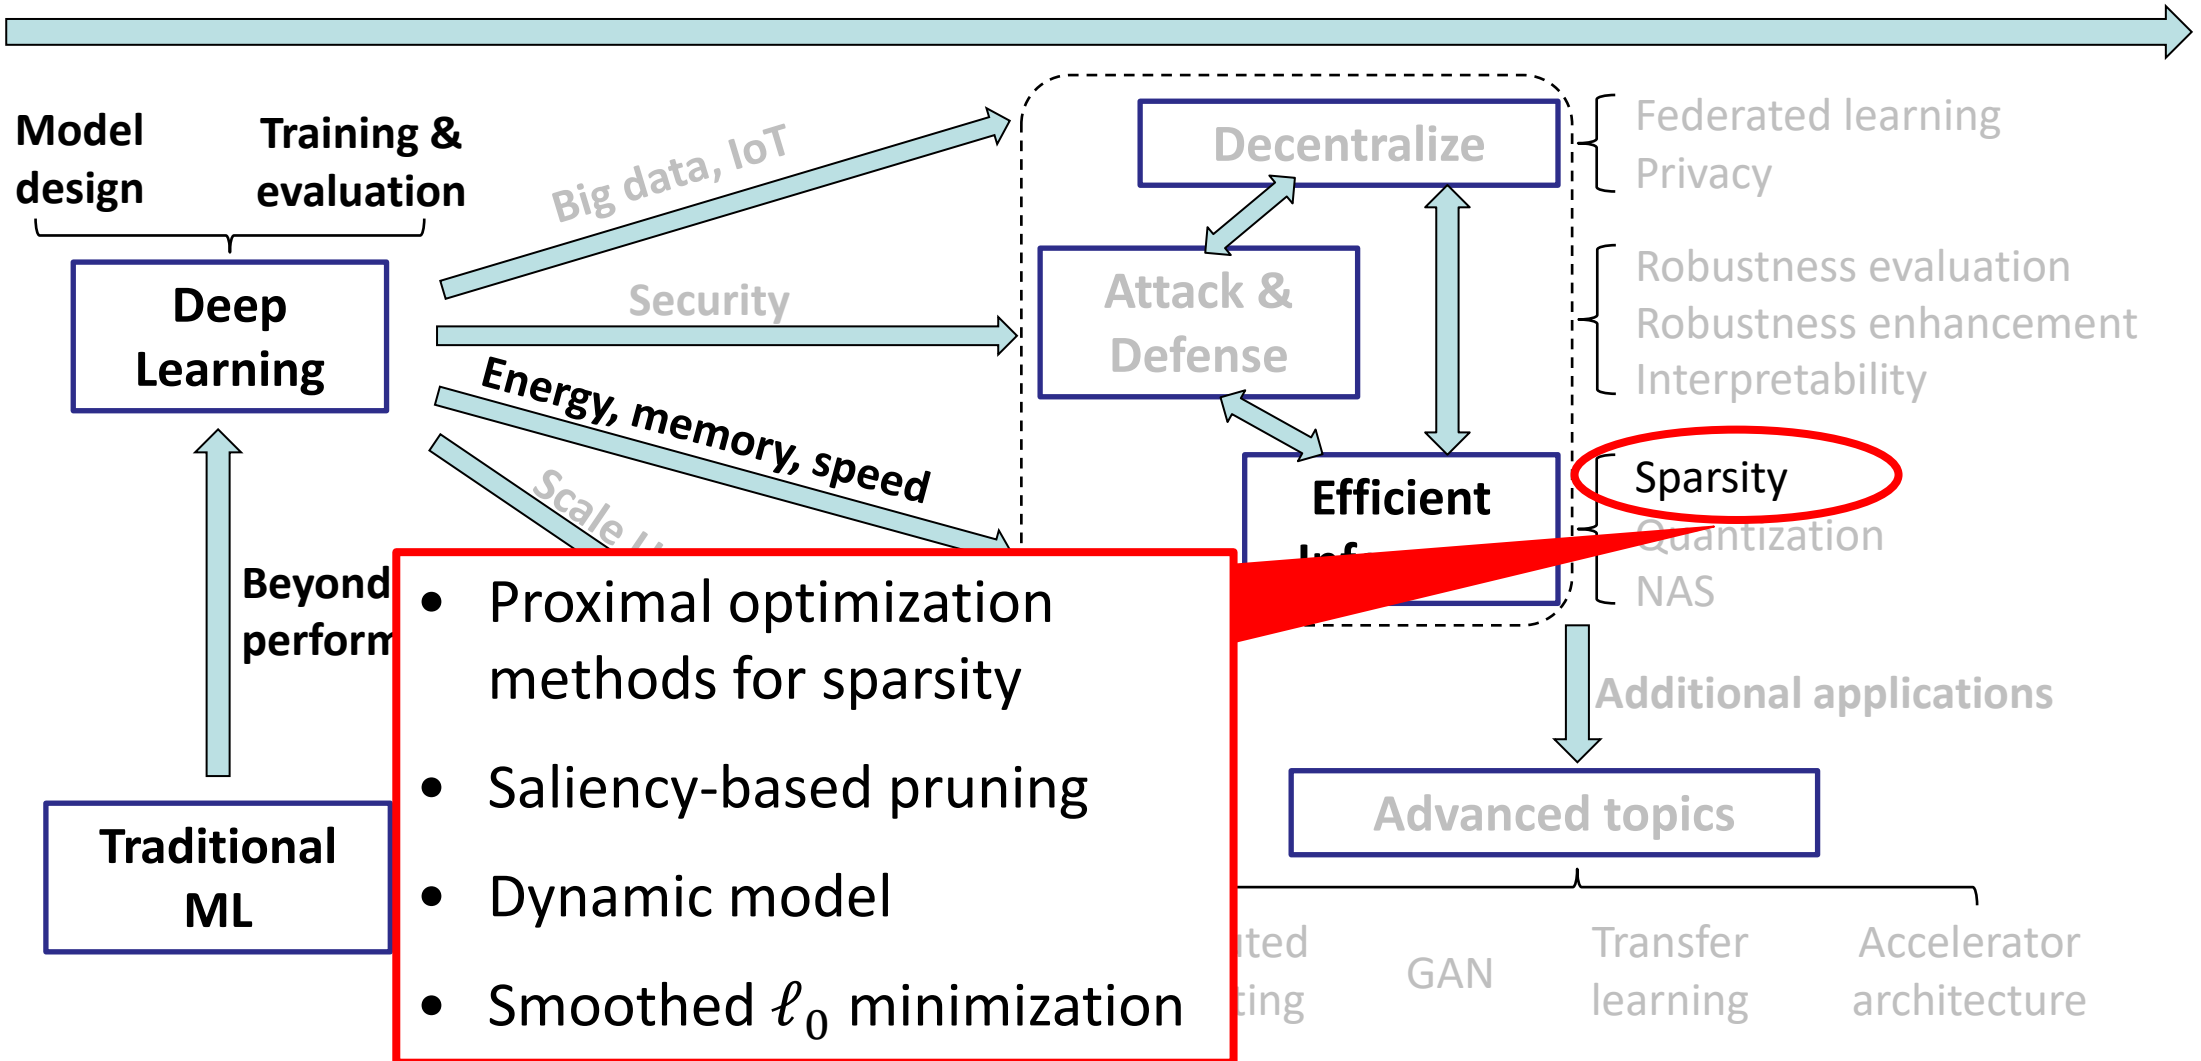
Previously

Applying machine learning into the real world



This lecture

Applying machine learning into the real world



Recall: Regularized training

- In a typical objective of regularized training of DNN, regularizer $R(\cdot)$ is applied to each parameter θ_i , then added to the overall objective with strength λ

$$\min_{\theta} \mathcal{L}(\theta) + \lambda \sum_i R(\theta_i)$$

- If the regularizer is differentiable, we can directly use SGD on the overall objective, but the performance will be sensitive to the choice of λ
 - We tend to choose a larger λ if we want the regularizer to have a stronger effect
 - But large strength will make the gradient of R dominate the overall gradient, affecting the minimization of $\mathcal{L}(\theta)$

Smoothing the regularized optimization

$$\min_{\theta} \mathcal{L}(\theta) + \lambda \sum_i R(\theta_i)$$

- Proximal gradient update
 - Consider $\min_{\theta} \mathcal{L}(\theta)$ as main objective, modify the resulted $\theta^{(l)}$ at each step slightly to minimize the regularizer term

Proximal gradient update

- Consider a general case $\min_{\theta} \mathcal{L}(\theta) + \lambda R(\theta)$, at step l
- Main objective: Optimize $\mathcal{L}(\theta)$

$$\hat{\theta}^{(l)} = \theta^{(l-1)} - \alpha \nabla_{\theta} \mathcal{L}(\theta^{(l-1)})$$

- Proximal operator: Optimize $R(\theta)$
 - Try minimize $R(\cdot)$ in the **proximity of $\hat{\theta}^{(l)}$**

$$\theta^{(l)} := \text{prox}_{\lambda R}(\hat{\theta}^{(l)}) = \underset{z}{\operatorname{argmin}} \left(\lambda R(z) + \frac{1}{2} \|z - \hat{\theta}^{(l)}\|_2^2 \right)$$

- The added proximity term will allow **smoother convergence** of the overall objective

Special case: ℓ_0 as indicator function

$$\hat{\theta} := \operatorname{argmin}_{\theta} \mathcal{L}(\theta) \text{ s.t. } \|\theta\|_0 \leq n$$

- Convert the constraint into an indicator function
 - Let $g(\theta) = \begin{cases} 0, & \|\theta\|_0 \leq n \\ +\infty, & \text{Otherwise} \end{cases}$, we can convert the optimization objective into $\hat{\theta} := \operatorname{argmin}_{\theta} \mathcal{L}(\theta) + g(\theta)$

- Main objective: Gradient descent to optimize $\mathcal{L}(\theta)$

$$\hat{\theta}^{(l)} = \theta^{(l-1)} - \alpha \nabla_{\theta} \mathcal{L}(\theta^{(l-1)})$$

- Proximal operator: Optimize $g(\theta)$

$$\begin{aligned} \theta^{(l)} &= \operatorname{argmin}_z \left(g(z) + \frac{1}{2} \|z - \hat{\theta}^{(l)}\|_2^2 \right) \\ &= \operatorname{argmin}_z \|z - \hat{\theta}^{(l)}\|_2 \text{ s.t. } \|z\|_0 \leq n \end{aligned}$$

Projection step for
iterative pruning

PGD is a special case of
proximal update with
indicator function


Special case: Revisit ℓ_1 regularization

$$\min_{\theta} \mathcal{L}(\theta) + \lambda \sum_i |\theta_i|$$

- Previously we considered direct SGD optimization of the regularized objective
- We can use the proximal update to better understand how ℓ_1 works
- Proximal operator of ℓ_1 : (known as soft thresholding)

$$\text{prox}_{\lambda \ell_1}(\theta_i) = \operatorname{argmin}_z \left(\lambda |z| + \frac{1}{2} \|z - \theta_i\|_2^2 \right)$$

$$= \begin{cases} \theta_i - \lambda, & \theta_i > \lambda \\ 0, & |\theta_i| \leq \lambda \\ \theta_i + \lambda, & \theta_i < -\lambda \end{cases}$$



Recall: Pros & Cons of ℓ_1

$$\text{prox}_{\lambda\ell_1}(\theta_i) = \begin{cases} \theta_i - \lambda, & \theta_i > \lambda \\ 0, & |\theta_i| \leq \lambda \\ \theta_i + \lambda, & \theta_i < -\lambda \end{cases}$$


- Pros:
 - Simple, effectively leading to sparsity
 - Pruning effect controlled by λ
- Cons:
 - **Biased**: the absolute value of large element always shrinks by λ , leading to loss of variance
- Improvement: change how large elements are handled

Improvements from ℓ_1

- There are a lot of ways to (partially) solve the “bias” problem of ℓ_1 , here we show some famous choices used in compressed sensing research

TABLE I: Regularization penalties and the corresponding proximity operator ($\lambda > 0$ is a thresholding parameter).

Penalty name	Penalty formulation	Proximity operator
(iii) ℓ_q -norm [6], [7]	$P_\lambda(x) = \lambda x ^q, 0 < q < 1$	$\text{prox}_{P_\lambda}(t) = \begin{cases} 0, & t < \tau \\ \{0, \text{sign}(t)\beta\}, & t = \tau \\ \text{sign}(t)y, & t > \tau \end{cases}$ <p>where $\beta = [2\lambda(1-q)]^{1/(2-q)}, \tau = \beta + \lambda q \beta^{q-1},$ $h(y) = \lambda q y^{q-1} + y - t = 0$ and $y \in [\beta, t]$</p>
(iv) q -shrinkage [10]	N/A ($q < 1$)	$\text{prox}_{P_\lambda}(t) = \text{sign}(t) \max\{ t - \lambda^{2-q} t^{q-1}, 0\}$
(v) SCAD [11]	$P_\lambda(x) = \begin{cases} \lambda x , & x < \lambda \\ \frac{2a\lambda x - x^2 - \lambda^2}{2(a-1)}, & \lambda \leq x < a\lambda \\ (a+1)\lambda^2/2, & x \geq a\lambda \end{cases}$	$\text{prox}_{P_\lambda}(t) = \begin{cases} \text{sign}(t) \max\{ t - \lambda, 0\}, & t \leq 2\lambda \\ \frac{(a-1)t - \text{sign}(t)a\lambda}{a-2}, & 2\lambda < t \leq a\lambda \\ t, & t > a\lambda \end{cases}$
(vi) MCP [12]	$P_{\lambda,\gamma}(x) = \lambda \int_0^{ x } \max(1 - t/(\gamma\lambda), 0) dt,$ <p>where $\gamma > 0$</p>	$\text{prox}_{P_{\lambda,\gamma}}(t) = \begin{cases} 0, & t \leq \lambda \\ \frac{\text{sign}(t)(t - \lambda)}{1 - 1/\gamma}, & \lambda < t \leq \gamma\lambda \\ t, & t > \gamma\lambda \end{cases}$
(vii) Firm thresholding [13]	$P_{\lambda,\mu}(x) = \begin{cases} \lambda[x - x^2/(2\mu)], & t \leq \mu \\ \lambda\mu/2, & t \geq \mu \end{cases}$ <p>where $\mu > \lambda$</p>	$\text{prox}_{P_{\lambda,\mu}}(t) = \begin{cases} 0, & t \leq \lambda \\ \frac{\text{sign}(t)(t - \lambda)\mu}{\mu - \lambda}, & \lambda \leq t \leq \mu \\ t, & t \geq \mu \end{cases}$



Further modification: Trimmed ℓ_1

- All previous variants of the ℓ_1 norm, including ℓ_1 itself, require a fixed strength/threshold as the hyperparameter, where a larger threshold leads to higher sparsity
- Recall iterative pruning, the pruning threshold is determined by the sparsity level we want to achieve
- We can combine the ideas by adding ℓ_1 penalty to a **certain percentage of the smallest weights**, while leaving larger weights unpenalized
- This is named “**Trimmed ℓ_1** ”, which allows sparsity exploration with a fixed upper bound

Formulation of the Trimmed ℓ_1

$$\begin{aligned} & \underset{\boldsymbol{\theta} \in \Omega, \boldsymbol{w} \in [0,1]^p}{\text{minimize}} && \mathcal{L}(\boldsymbol{\theta}; \mathcal{D}) + \lambda_n \sum_{j=1}^p w_j |\theta_j| \\ & \text{s. t.} && \mathbf{1}^\top \boldsymbol{w} \geq p - h. \end{aligned} \tag{1}$$

- p : dimension of parameter θ
- h : the number of nonzero parameters we want to keep
- “Gate variable” w : learnable ℓ_1 strength on each element (with constraint)
- The combination of the minimization objective and the constraint on the w enforces that the ℓ_1 penalty will be added on the smallest $p - h$ parameters at convergence

Optimization of Trimmed ℓ_1

$$\begin{aligned} & \underset{\boldsymbol{\theta} \in \Omega, \boldsymbol{w} \in [0,1]^p}{\text{minimize}} && \mathcal{L}(\boldsymbol{\theta}; \mathcal{D}) + \lambda_n \sum_{j=1}^p w_j |\theta_j| \\ & \text{s. t.} && \mathbf{1}^\top \boldsymbol{w} \geq p - h. \end{aligned}$$

Objective of \boldsymbol{w}

(1)

- Constrained optimization \rightarrow PGD
- PGD step for \boldsymbol{w}

$$\boldsymbol{w}^{k+1} \leftarrow \underset{\text{Projection}}{\text{proj}_{\mathcal{S}}} [\underset{\text{Gradient descent}}{\boldsymbol{w}^k - \tau \boldsymbol{r}(\boldsymbol{\theta}^k)}]$$

\mathcal{S} encodes the constraints $0 \leq w_i \leq 1, \mathbf{1}^\top \boldsymbol{w} = p - h$.

Yun, Jihun, et al. "Trimming the ℓ_1 Regularizer: Statistical Analysis, Optimization, and Applications to Deep Learning." *International Conference on Machine Learning*. 2019.

Optimization of Trimmed ℓ_1

$$\begin{aligned} & \underset{\boldsymbol{\theta} \in \Omega, \mathbf{w} \in [0,1]^p}{\text{minimize}} && \mathcal{L}(\boldsymbol{\theta}; \mathcal{D}) + \lambda_n \sum_{j=1}^p w_j |\theta_j| && \text{Objective of } \theta \\ & \text{s. t.} && \mathbf{1}^\top \mathbf{w} \geq p - h. && (1) \end{aligned}$$

- Regularized optimization -> Proximal gradient

- Proximal operator for the regularizer:

$$\text{prox}_{\eta\lambda\mathcal{R}(\cdot, \mathbf{w}^{k+1})}(\mathbf{z}) := \arg \min_{\boldsymbol{\theta}} \frac{1}{2\eta\lambda} \|\boldsymbol{\theta} - \mathbf{z}\|^2 + \sum_{j=1}^p w_j^{k+1} |\theta_j|$$

- Have a form of ℓ_1 , solution is the soft thresholding operator

- Proximal update step for θ

$$\boldsymbol{\theta}^{k+1} \leftarrow \text{prox}_{\eta\lambda\mathcal{R}(\cdot, \mathbf{w}^{k+1})}[\boldsymbol{\theta}^k - \eta \nabla \mathcal{L}(\boldsymbol{\theta}^k)]$$

Proximal

(soft thresholding)

Gradient
descent

Yun, Jihun, et al. "Trimming the ℓ_1 Regularizer: Statistical Analysis, Optimization, and Applications to Deep Learning." *International Conference on Machine Learning*. 2019.

Optimization of Trimmed ℓ_1

$$\begin{aligned} & \underset{\boldsymbol{\theta} \in \Omega, \boldsymbol{w} \in [0,1]^p}{\text{minimize}} \quad \mathcal{L}(\boldsymbol{\theta}; \mathcal{D}) + \lambda_n \sum_{j=1}^p w_j |\theta_j| \\ & \text{s. t.} \quad \mathbf{1}^\top \boldsymbol{w} \geq p - h. \end{aligned} \quad (1)$$

- Two operators:

- Projection for w : $\text{proj}_{\mathcal{S}}(\boldsymbol{z}) := \arg \min_{\boldsymbol{w} \in \mathcal{S}} \frac{1}{2} \|\boldsymbol{z} - \boldsymbol{w}\|^2$

- Proximal for θ : $\text{prox}_{\eta\lambda\mathcal{R}(\cdot, \boldsymbol{w}^{k+1})}(\boldsymbol{z}) := \arg \min_{\boldsymbol{\theta}} \frac{1}{2\eta\lambda} \|\boldsymbol{\theta} - \boldsymbol{z}\|^2 + \sum_{j=1}^p w_j^{k+1} |\theta_j|$

- Putting everything together

Algorithm 1 Block Coordinate Descent for (1)

Input: λ , η , and τ .

Initialize: $\boldsymbol{\theta}^0$, \boldsymbol{w}^0 , and $k = 0$.

while not converged **do**

$\boldsymbol{w}^{k+1} \leftarrow \text{proj}_{\mathcal{S}}[\boldsymbol{w}^k - \tau \boldsymbol{r}(\boldsymbol{\theta}^k)]$

$\boldsymbol{\theta}^{k+1} \leftarrow \text{prox}_{\eta\lambda\mathcal{R}(\cdot, \boldsymbol{w}^{k+1})}[\boldsymbol{\theta}^k - \eta \nabla \mathcal{L}(\boldsymbol{\theta}^k)]$

$k \leftarrow k + 1$

end while

Output: $\boldsymbol{\theta}^k$, \boldsymbol{w}^k .

Yun, Jihun, et al. "Trimming the ℓ_1 Regularizer: Statistical Analysis, Optimization, and Applications to Deep Learning." *International Conference on Machine Learning*. 2019.

Summary of proximal operator

- Proximal operator makes regularized optimization smoother by introducing the proximity term
 - L0 -> projected gradient descent (hard thresholding)
 - L1 -> Soft thresholding
- Soft thresholding reveals the “bias” problem of L1, partially solved by SCAD & MCP etc.
- Trimmed L1 combines hard and soft thresholding for better sparsity inducing performance on DNN models

Moving beyond absolute values

- Previous projection/proximal methods use the absolute value of the weights as pruning/penalization criterion
- What if a small weight element is also useful?
 - Small changes to activation may cause major differences in the final output (i.e. adversarial example)
- Measuring the importance of the weight element
 - How loss changes when setting a weight element to zero

$$L(x, y, \Theta_{w \leftarrow 0}) = L(x, y, \Theta) - \frac{\partial L(x, y, \Theta)}{\partial w} (0 - w) + o(w^2) \quad \text{Taylor expansion}$$

$$|L(x, y, \Theta_{w \leftarrow 0}) - L(x, y, \Theta)| = \left| \frac{\partial L(x, y, \Theta)}{\partial w} w \right| = T(x, y, w) \quad \text{Ignore high-order term}$$

Saliency

Global Sparse Momentum (GSM) SGD

- Momentum SGD with weight decay for weights with high saliency
- **No gradient update** for low-saliency elements, only do weight decay.

$$B_{m,n}^{(k)} = \begin{cases} 1 & \text{if } T(x, y, W_{m,n}^{(k)}) \geq \text{the } Q\text{-th greatest value in } T(x, y, \Theta^{(k)}) , \\ 0 & \text{otherwise .} \end{cases}$$

$$\begin{array}{ll} \text{Momentum} & \mathbf{Z}^{(k+1)} \leftarrow \beta \mathbf{Z}^{(k)} + \eta \mathbf{W}^{(k)} + \overset{\text{Weight decay}}{B^{(k)}} \circ \overset{\text{Gradient}}{\frac{\partial L(x, y, \Theta)}{\partial \mathbf{W}^{(k)}}} , \\ \text{Weight update} & \mathbf{W}^{(k+1)} \leftarrow \mathbf{W}^{(k)} - \alpha \mathbf{Z}^{(k+1)} , \end{array}$$

- Global ranking across all layers
 - Saliency is more comparable across layers, than abs value
 - Enable more flexible sparsity allocation

Pruning results

Table 1: Pruning results on MNIST.

Model	Result	Base Top1	Pruned Top1	Origin / Remain Params	Compress Ratio	Non-zero Ratio
LeNet-300	Han et al. [21]	98.36	98.41	267K / 22K	12.1×	8.23%
LeNet-300	L-OBS [13]	98.24	98.18	267K / 18.6K	14.2×	7%
LeNet-300	Zhang et al. [56]	98.4	98.4	267K / 11.6K	23.0×	4.34%
LeNet-300	DNS [18]	97.72	98.01	267K / 4.8K	55.6×	1.79%
LeNet-300	GSM	98.19	98.18	267K / 4.4K	60.0×	1.66%
LeNet-5	Han et al. [21]	99.20	99.23	431K / 36K	11.9×	8.35%
LeNet-5	L-OBS [13]	98.73	98.73	431K / 3.0K	14.1×	7%
LeNet-5	Srinivas et al. [47]	99.20	99.19	431K / 22K	19.5×	5.10%
LeNet-5	Zhang et al. [56]	99.2	99.2	431K / 6.05K	71.2×	1.40%
LeNet-5	DNS [18]	99.09	99.09	431K / 4.0K	107.7×	0.92%
LeNet-5	GSM	99.21	99.22	431K / 3.4K	125.0×	0.80%
LeNet-5	GSM	99.21	99.06	431K / 1.4K	300.0×	0.33%

Table 2: Pruning results on CIFAR-10.

Model	Result	Base Top1	Pruned Top1	Origin / Remain Params	Compress Ratio	Non-zero Ratio
ResNet-56	GSM	94.05	94.10	852K / 127K	6.6×	15.0%
ResNet-56	GSM	94.05	93.80	852K / 85K	10.0×	10.0%
DenseNet-40	GSM	93.86	94.07	1002K / 150K	6.6×	15.0%
DenseNet-40	GSM	93.86	94.02	1002K / 125K	8.0×	12.5%
DenseNet-40	GSM	93.86	93.90	1002K / 100K	10.0×	10.0%

Direct parameter selection

- All previously mentioned methods are aiming for making more parameters closer to zero

- Can we directly select the parameters to be pruned?

- This is a process for parameter selection, or gating (masking)

$$\theta_j = \tilde{\theta}_j z_j, z_j \in \{0,1\}, \tilde{\theta}_j \neq 0, ||\theta||_0 = \sum z_j$$

- To make it trainable with gradient-based optimization, we need a **differentiable, nonlinear selection (gating) function**

- Stochastic gate with hard concrete distribution
 - Continuous sparsification

Stochastic binary gates

$$\theta_j = \tilde{\theta}_j z_j, z_j \in \{0,1\}, \tilde{\theta}_j \neq 0, ||\theta||_0 = \sum z_j$$

- Here we consider the stochastic case, where the state of the gate is z_j determined by a Bernoulli random variable

$$q(z_j | \pi_j) = \text{Bern}(\pi_j)$$

- The ℓ_0 regularization can be converted to

$$\tilde{\theta}^*, \pi^* = \operatorname{argmin}_{\tilde{\theta}, \pi} \mathcal{L}(\tilde{\theta} \cdot z) + \lambda \sum \pi_j$$

The Bernoulli distribution is still discrete, cannot pass gradient from z to π directly. Further smoothing is required for optimization, details in the reading material

Continuous Sparsification

$$\theta_j = \tilde{\theta}_j z_j, z_j \in \{0,1\}, \tilde{\theta}_j \neq 0, ||\theta||_0 = \sum z_j$$

- Stochastic gate method requires multiple steps of smoothing and estimation, leading to large variance and poor scalability to larger models
- Rethink the problem we are facing
 - Training requires the gate to be continuous and differentiable
 - Sparsity requires gate to be binary for the **final model**
- No need to have binary gate throughout training process, can keep it continuous, while gradually moving towards binary as training progresses

Continuous Sparsification

- A function converting continuous input towards binary output -> Sigmoid with temperature

$$\sigma(\beta s) = \frac{1}{1 + e^{-\beta s}}, \beta \in [1, \infty)$$

- $\sigma(\beta s)$ is continuously differentiable, yet go towards binary with $\beta \rightarrow \infty$, satisfying our need for the mask
- Using $\sigma(\beta s)$ as mask:
 - Original $L = L(\theta) + \lambda ||\theta||_0 = L(\tilde{\theta}_j z_j) + \lambda \sum z_j, z_j \in \{0, 1\}$
 - Converted $L_\beta = L(\tilde{\theta}_j \sigma(\beta s_j)) + \lambda \sum \sigma(\beta s_j), s_j \in \mathbb{R}$
- Increasing β from 1 to ∞ during training, leading to sparse model in the end

Summary of ℓ_0 regularization

- ℓ_0 regularization is an effective method for model pruning, as it directly measures the sparsity, without shrinking larger values, and can result in exact zero parameters
- The hard-concrete distribution provides a particle way of approximating and optimizing the ℓ_0 regularization, but suffers from high variance in estimating the gradient, and is hard to scale to larger problems
- Continuous sparsification makes the minimization of ℓ_0 regularization more stable

Moving beyond pruning

- Pruning aims for removing part of the model
- Often leads to loss in model capacity -> loss in accuracy, leading to efficiency-accuracy tradeoff
- Can we get the best of both worlds?
 - Models are designed in this way for a reason, need the capacity to learn the entire dataset
 - But may not need that large capacity to learn **easy data**



(a) Red wine

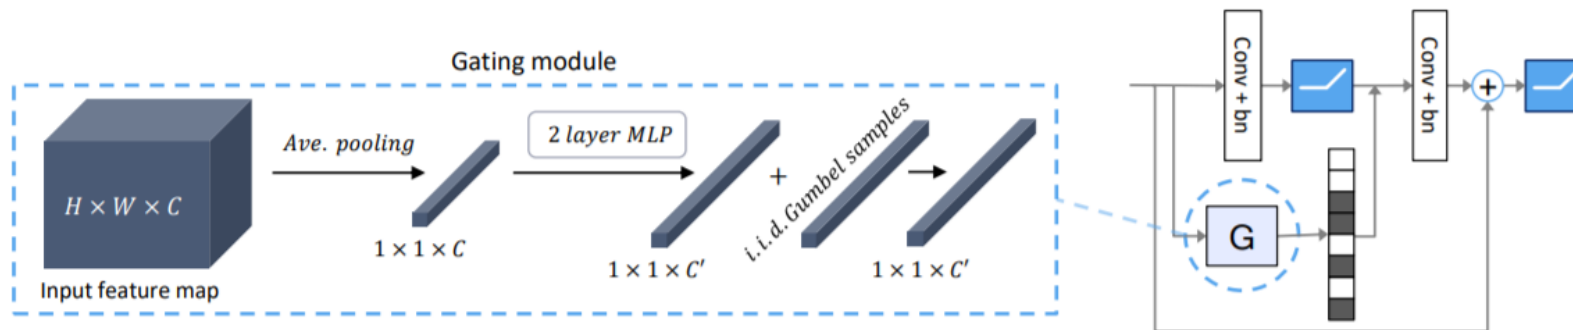
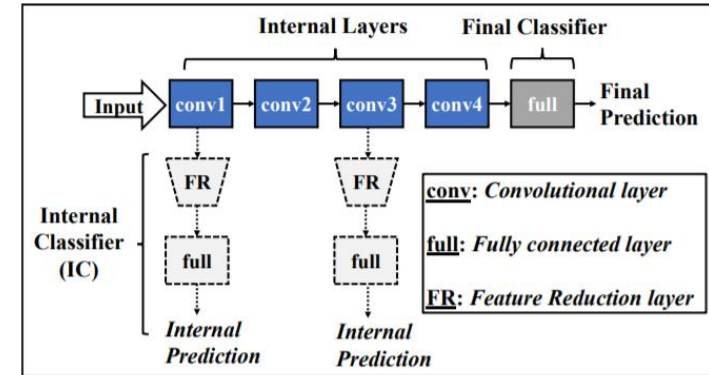


(b) Volcano

Huang, G., Chen, D., Li, T., Wu, F., van der Maaten, L., & Weinberger, K. Q. (2017). Multi-scale dense networks for resource efficient image classification. *arXiv preprint arXiv:1703.09844*.

Processing easy data with less effort

- Make the model input-adaptive
 - Multi-exit: make classification with intermediate layer feature if a certain condition is met
- Dynamic execution: dynamically changing (gating) model architecture for inputs with different features

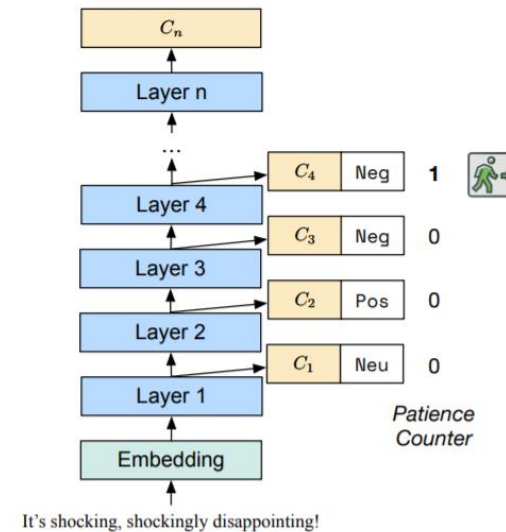


Early exit model

- Model training objective
 - Weighted sum of multi-exits' losses
 - Weighing earlier exit more will encourage earlier exit, but may hurt overall accuracy
- Early exit criteria
 - Entropy threshold
 - Output consistency

$$L_{\text{branchynet}}(\hat{\mathbf{y}}, \mathbf{y}; \theta) = \sum_{n=1}^N w_n L(\hat{\mathbf{y}}_{\text{exit}_n}, \mathbf{y}; \theta)$$

$$\text{entropy}(\mathbf{y}) = \sum_{c \in \mathcal{C}} y_c \log y_c$$




Teerapittayanon, S., McDanel, B., & Kung, H. T. (2016, December). Branchynet: Fast inference via early exiting from deep neural networks. In *2016 23rd International Conference on Pattern Recognition (ICPR)* (pp. 2464-2469). IEEE.


Zhou, W., Xu, C., Ge, T., McAuley, J., Xu, K., & Wei, F. (2020). Bert loses patience: Fast and robust inference with early exit. *arXiv preprint arXiv:2006.04152*.

Effectiveness of early exiting

- Overcoming overfitting and overthinking



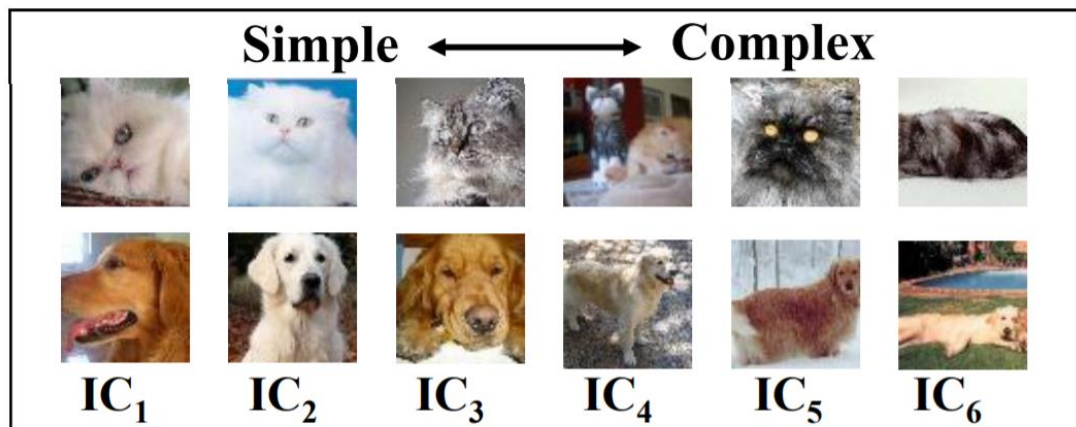
(a) Overfitting in training



(b) Overthinking in inference

Zhou, W., Xu, C., Ge, T., McAuley, J., Xu, K., & Wei, F. (2020). Bert loses patience: Fast and robust inference with early exit. *arXiv preprint arXiv:2006.04152*.

- Identify easy data at early stage



Kaya, Y., Hong, S., & Dumitras, T. (2019, May). Shallow-deep networks: Understanding and mitigating network overthinking. In *International Conference on Machine Learning* (pp. 3301-3310). PMLR.

Dynamic model architecture

- Recall in CNN lectures: CNN utilize multiple channels to capture different features in each layer
- The model needs to capture all possible features to fit a dataset well
- Each input may only utilize part of the feature
- Can let the model dynamically decide which channel to use based on the input

Typical gating module

- **Gating module** decides which channel to choose
- For each layer, the input information is captured from the output feature of previous layer
- Gating module perform a cheap computation based on previous layer's output
 - Global average pooling + FC
 - L0 reg on the gate to encourage efficiency

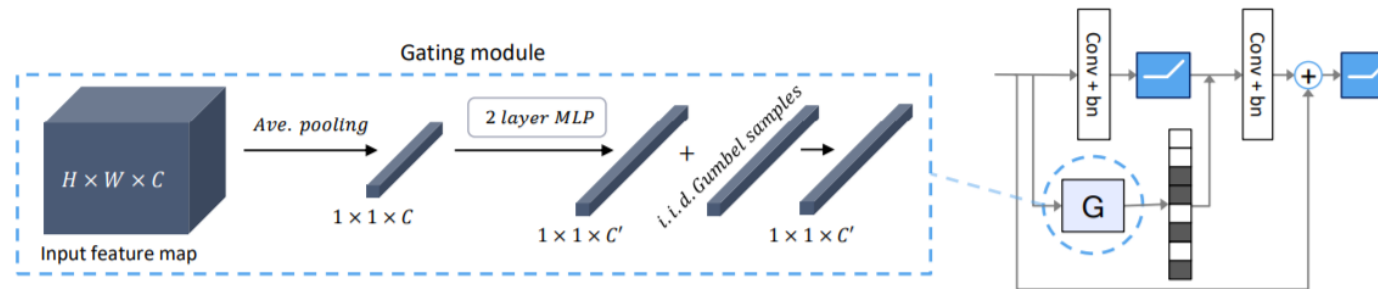




Figure 2: Illustration of our channel gated ResNet block and the gating module.

What is learnt in the gate?

- Gate activation pattern: few gates are totally useless, while most of them are used occasionally



- Similar classes activates similar channels

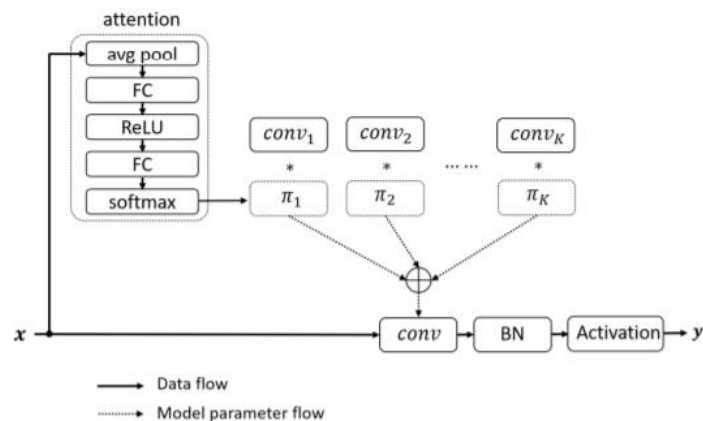


Variant of dynamic model



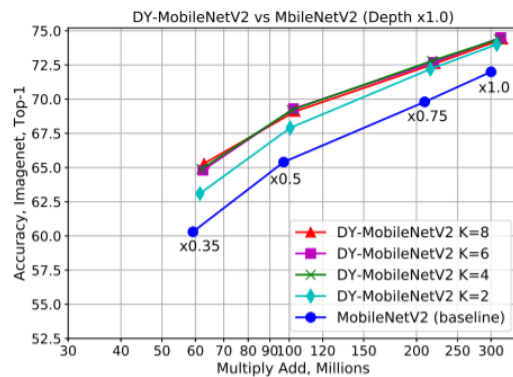
- Dynamic models make inference cost smaller under the same architecture
- Similarly, it will enable the use of a larger model with similar inference cost -> **Overparameterization**
- Overparameterized model use more parameters in the training process, while dynamically combining them back to the original model size at inference
 - Can learn more diverse features
 - Can converge to lower loss in complex tasks

Case study: Dynamic convolution

- Have multiple trainable kernels in each layer
- Use a weighted sum of all kernels for forward pass, weights generate dynamically based on previous layer's feature



Chen, Y., Dai, X., Liu, M., Chen, D., Yuan, L., & Liu, Z. (2020). Dynamic convolution: Attention over convolution kernels. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition* (pp. 11030-11039).



In this lecture, we learned:

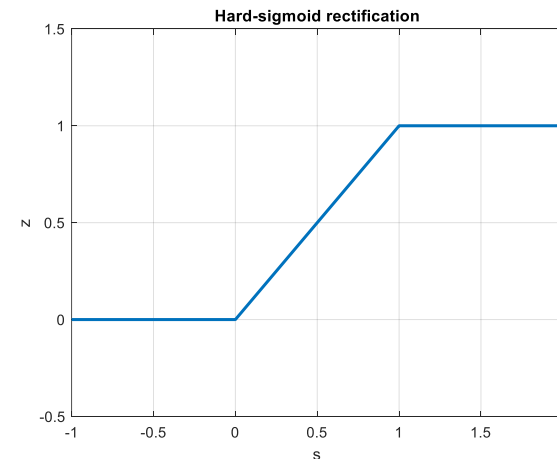
- Proximal optimization methods
 - Soft thresholding and beyond
 - Trimmed L1
- Alternative pruning criteria
 - Saliency-based pruning
- Continuous smoothing for ℓ_0 optimization
- Dynamic NN model
- See cited papers for details on each technique

Smoothing the discrete gate (reading)

$$\tilde{\theta}^*, \pi^* = \operatorname{argmin}_{\tilde{\theta}, \pi} \mathcal{L}(\tilde{\theta} \cdot z) + \lambda \sum \pi_j$$

- The Bernoulli distribution is still discrete, cannot pass gradient from z to π directly
- Consider a random variable s defined by $s \sim q(s|\phi)$, we can smooth z by letting $z = \min(1, \max(0, s))$
 - This is named “[hard-sigmoid rectification](#)”
- Recall that $\pi_j = p_r(z \neq 0)$, with $Q(\cdot)$ being the CDF of s we have:

$$\pi_j = 1 - Q(s_j \leq 0 | \phi)$$



Hard concrete distribution

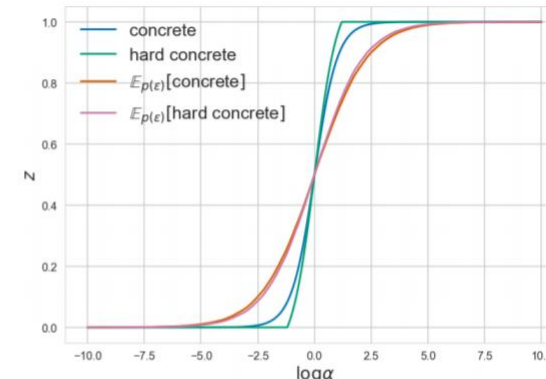

$$\tilde{\theta}^*, \pi^* = \operatorname{argmin}_{\tilde{\theta}, \pi} \mathcal{L}(\tilde{\theta} \cdot z) + \lambda \sum \pi_j$$

- Now z and π_j are related to s with some differentiable function, only thing left is to define $q(s|\phi)$

- Let s follow a **concrete distribution** with $\phi = \log \alpha$

$$u \sim \mathcal{U}(0, 1), \quad s = \operatorname{Sigmoid}((\log u - \log(1 - u) + \log \alpha)/\beta), \quad \bar{s} = s(\zeta - \gamma) + \gamma$$

- $z = \min(1, \max(0, s))$ will be following a **hard-concrete distribution**



Hard concrete ℓ_0 regularization

- The ℓ_0 norm can be relaxed with the hard-concrete distribution as

$$\mathcal{L}_C = \sum_{j=1}^{|\theta|} (1 - Q_{\tilde{s}_j}(0|\phi)) = \sum_{j=1}^{|\theta|} \text{Sigmoid}(\log \alpha_j - \beta \log \frac{-\gamma}{\zeta}).$$

$$\hat{\mathbf{z}} = \min(\mathbf{1}, \max(\mathbf{0}, \text{Sigmoid}(\log \boldsymbol{\alpha})(\zeta - \gamma) + \gamma))$$

$$\tilde{\boldsymbol{\theta}}^*, \alpha^* = \operatorname{argmin}_{\tilde{\boldsymbol{\theta}}, \alpha} \mathcal{L}(\tilde{\boldsymbol{\theta}} \cdot \hat{\mathbf{z}}) + \lambda \mathcal{L}_C$$

- Here α is a trainable variable, same size as $\tilde{\boldsymbol{\theta}}$
- The combined loss is differentiable
- $0 < \beta < 1$ is a hyperparameter for the distribution
 - $\beta \rightarrow 0$, \mathbf{z} will have a high probability to be 0 or 1, closer approximation to original Bernoulli
 - $\beta \rightarrow 1$, more density between 0 and 1, a flatter density function leading to easier optimization